

OS Overview / Process and Thread

국민대학교 임베디드 연구실
경 주 현

오늘 강의 계획

- OS 기본 이론 및 강의 내용 설명
- 리눅스 기본 이론
 - 성능: 멀티코어, 리눅스 Lock and scalability
 - 파워: 리눅스 PM, Timer 기반 DVFS
- 리눅스 스케줄러
 - 리눅스 스케줄러 class, 트레이스(실습)
 - CFS, CFS Parameters, Load Balancer
 - Group Scheduling, Bandwidth Control, PELT
 - EAS(Energy Aware Scheduling)
 - Android PM, Deadline scheduler

Goals

**Understand Linux Scheduler/DVFS in detail
by designing and implementing**

What do apps want from a O/S?

What do applications want from a O/S?

- **Abstract the hardware for convenience and portability**
- **Multiplex the hardware among multiple applications**
- **Isolate applications to contain bugs**
- **Allow sharing among applications**

What services does a O/S typically provide?

What services does a O/S typically provide?

- **Processes**
- **Memory**
- **File contents**
- **Directories and file names**
- **Security**
- **Many others: users, IPC, network, time, terminals**

What is key requirements for kernels?

Key requirements for kernels

- **Isolation**
- **Multiplexing**
- **Interaction**

Focus on isolation

Constraining requirement.

What is isolation?

- **Enforced separation to contain effects of failures**
- **The process is the usual unit of isolation**

What is isolation?

- Prevent process X from wrecking or spying on process Y
- Prevent a process from wrecking the operating system itself
- Prevent a malice as well as bugs

In order to isolate,

In order to isolate

- **Apps must use OS interface**
- **Apps cannot directly interact with hardware**
- **Apps cannot harm operating system**
- **Apps cannot directly affect each other**



Processor

Processors provide mechanisms

- **Hardware provides user mode and kernel mode**
- **Some instructions can only be executed in kernel mode**
- **Device access, processor configuration, isolation mechanisms**

Hardware isolation mechanisms

- **Operating systems runs in kernel mode**
 - kernel is a big program
 - services: processes, file system, net
 - low-level: devices, virtual memory
 - all of kernel runs with full hardware privilege
 - convenient
- **Applications run in user mode**
 - Kernel sets up per-process isolated address space
 - **System calls** switch between user and kernel mode



O/S

How to isolate process memory?

How to isolate process memory?

- **Address Space**
 - Give memory - code, variables, heap, stack
- **prevent it from accessing other memory**
 - kernel or other processes
- **"paging hardware" in the memory management unit (MMU)**
- **MMU translates (or "maps") every address issued by program**
- **VA -> PA**

How to isolate CPU?

How to isolate CPU?

- **H/W provides a periodic "clock interrupt"**
 - Forcefully suspends current process
 - Jumps into kernel
 - Kernel can switch to a different process
- **Preventing long computations, or buggy infinite loops**
- **Kernel must save/restore process state (registers)**
 - Context switch

Process

- **An abstract virtual machine**
- **As if it had its own CPU and memory**
- **Not accidentally affected by other processes.**
- **Motivated by isolation**

Thread

- An abstraction that contains enough state of a running program
- It can be stopped and resumed
- Multiple threads **share** an address space

Overview of switch between two threads

User -> kernel transition (system call or timer)

Kernel -> kernel switch

Kernel -> user transition

Conclusion

**Apps requirements -> O/S -> Isolation -> scheduler
-> process/thread**

Reference

- <https://pdos.csail.mit.edu/6.828/2016/schedule.html>
- <http://web.mit.edu/6.033>
- <http://www.rdrop.com/~paulmck/>
- "Is Parallel Programming Hard, And If So, What Can You Do About It?"
- Davidlohr Bueso. 2014. Scalability techniques for practical synchronization primitives. *Commun. ACM* 58

<http://queue.acm.org/detail.cfm?id=2698990>

- "CPUFreq and The Scheduler Revolution in CPU Power Management", Rafael J. Wysocki
- <https://sites.google.com/site/embedwiki/oses/linux/pm/pm-qos>
- <https://intl.aliyun.com/forum/read-916>
- User-level threads : co-routines

<http://www.gamedevforever.com/291>

https://www.youtube.com/watch?v=YYtzQ355_Co

- Scheduler Activations
 - <https://cgi.cse.unsw.edu.au/~cs3231/12s1/lectures/SchedulerActivations.pdf>
- [https://en.wikipedia.org/wiki/FIFO_\(computing_and_electronics\)](https://en.wikipedia.org/wiki/FIFO_(computing_and_electronics))
- <http://jake.dothome.co.kr/>
- <http://www.linuxjournal.com/magazine/completely-fair-scheduler?page=0.0>
- https://www2.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/6_CPU_Scheduling.html
- "Energy Aware Scheduling", Byungchul Park, LG Electronic
- "Update on big.LITTLE scheduling experiments", ARM
- "EAS Update" 2015 september ARM
- "EAS Overview and Integration Guide", ARM TR
- "Drowsy Power Management", Matthew Lentz, SOSP 2015
- <https://www.slideshare.net/nanik/learning-aosp-android-hardware-abstraction-layer-hal>
- <https://www.youtube.com/watch?v=oTGQXqD3CNI>
- <https://www.youtube.com/watch?v=P80NcKUKpuo>
- <https://lwn.net/Articles/398470/>
- "SCHED_DEADLINE: It's Alive!", ARM, 2017