

# Linux Power

**Basic, Linux PM, DVFS, OPP, Runtime PM, PM QoS**


국민대학교 임베디드 연구실  
경주현

# Outline




- **Processor Power Management**
- **Linux Power Management**

# How to reduce power?

# How to reduce power?



WIFI, Modem chip, LCD, Touch ....



- Power Rail Check
- Total Current Check

# Background

- **Energy**
  - joule
- **Power**
  - watt
- **Processor Power Management**
  - P-state, C-state



# **Processor Power Management**



# Silicon Power Management Features



















- **P – States are a sub-state of the C0 state**
  - They offer reduced power consumption while the processor is executing code.
- **C-States are a sub-state of the S0 state**
  - They offer reduced power consumption while the system is fully on.
- **S-States**
  - Suspend to idle, Standby, Suspend to ram, suspend to disk

# Processor Power Management

- **CPU C-State (Idle Power State)**

- C0 - Full On State
- C1 - Auto Halt State
- C2, C4 (C3, C4, C5), C6

- **CPU P-State (Performance State)**

|                  | Active state  | Sleep states  |   |   |   |
|------------------|---|---|---|---|---|
|                  | <u>C0</u>   | <u>C1/C1E</u>   | <u>C3</u>   | <u>C6</u>   | <u>C7</u>   |
|                  | Operating   | Halt  | Sleep   | Deep Sleep  |   |
| Core clock       |  | off   | off   | off   | off   |
| PLL              |  |  | off   | off   | off   |
| Core caches      |  |  | flushed   | flushed   | flushed   |
| Shared cache     |  |  |  |  | flushed   |
| Wakeup time*     | active  |  |  |  |  |
| Core Idle power* |  |  |  | ~ 0   | < C6  |

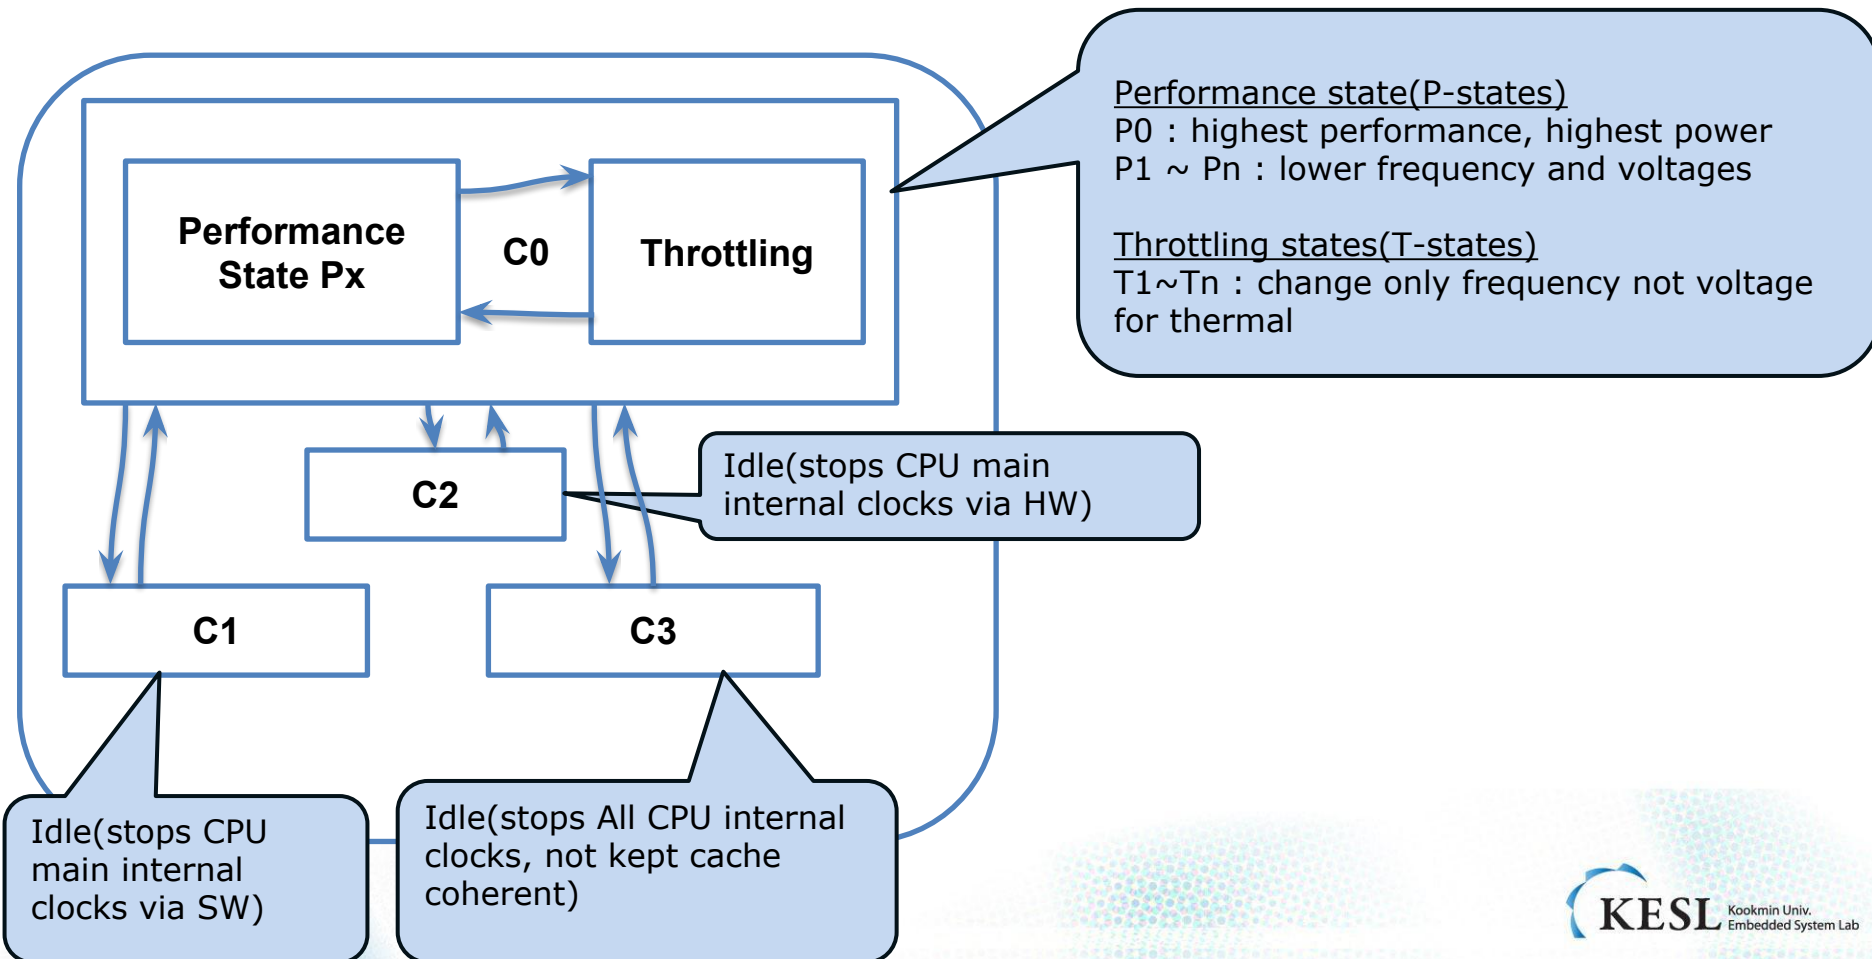
\* Rough approximation

<http://download.intel.com/embedded/technology/PowerManagementforEmbeddedApps.pdf>



# Silicon Power Management Features

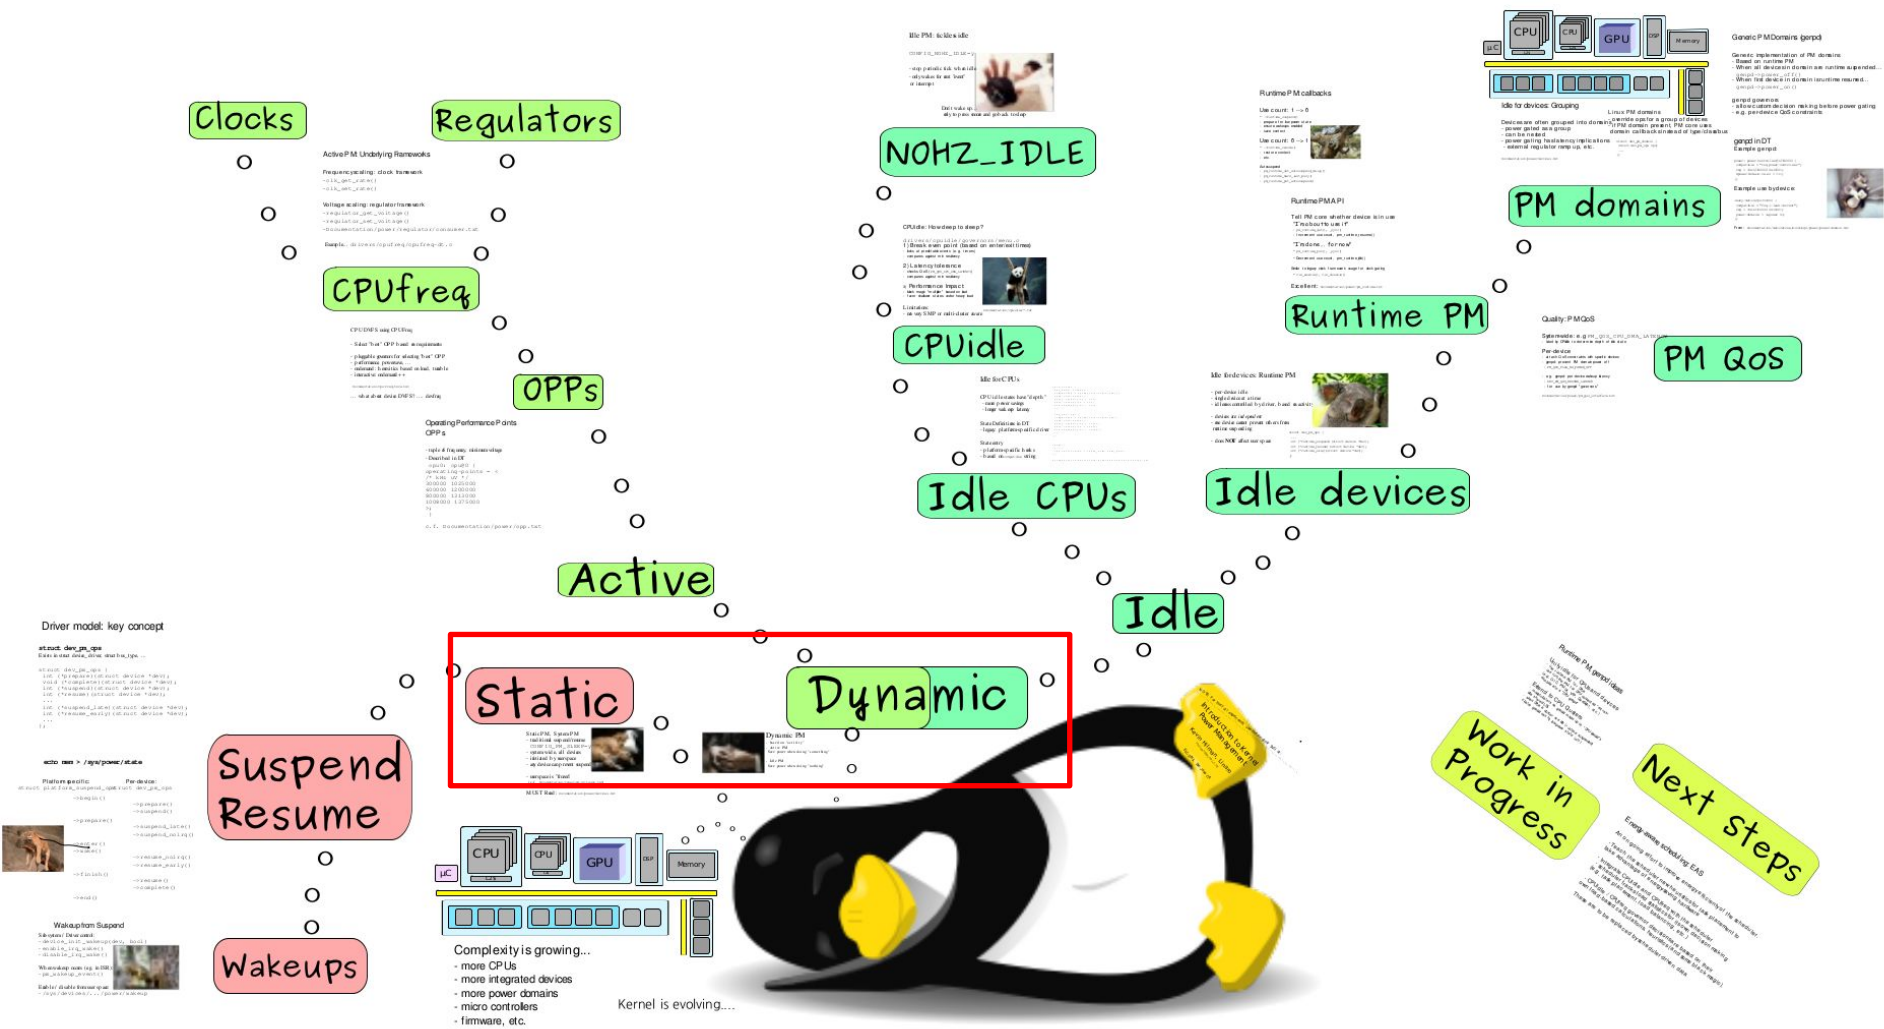
- Advanced Configuration and Power Interface (ACPI)





# Linux Power Management

# Tree view - Dynamic and Static PM



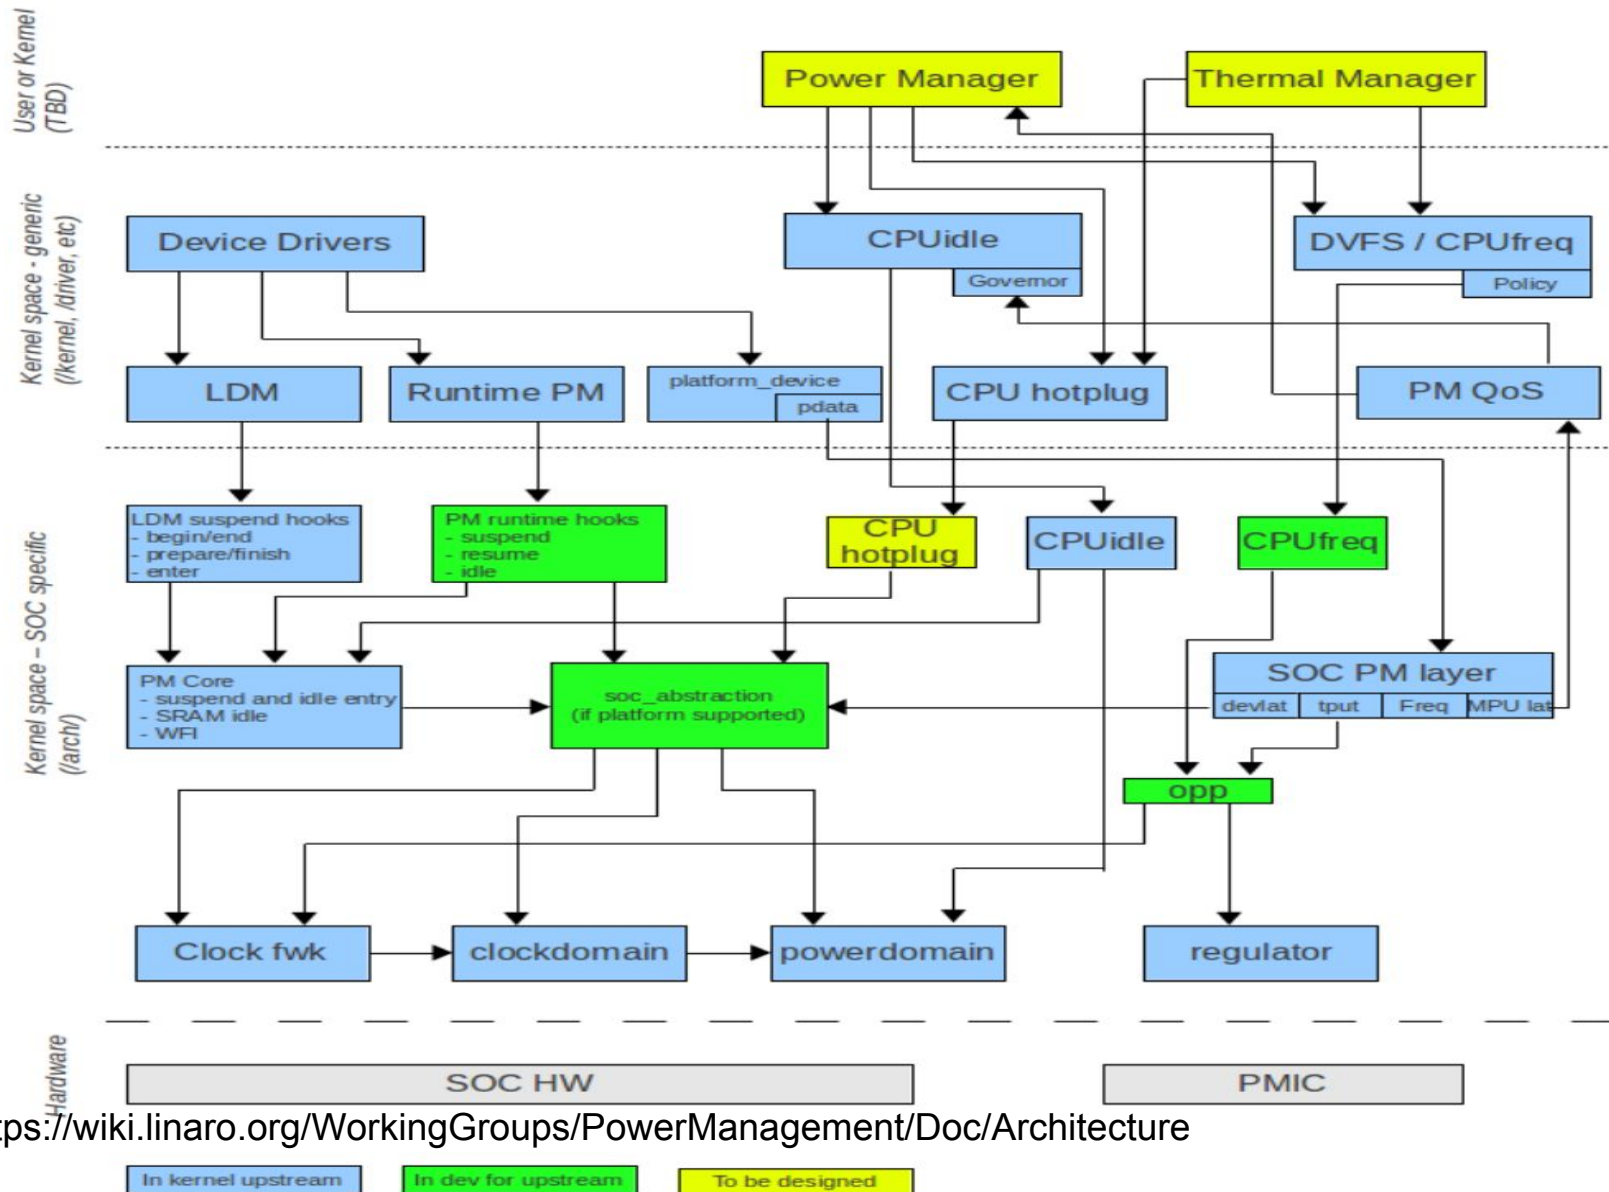
Reference from "Introduction to Kernel Power Management" Kevin Hilman, Linaro

# Linux Driver Model

- Example

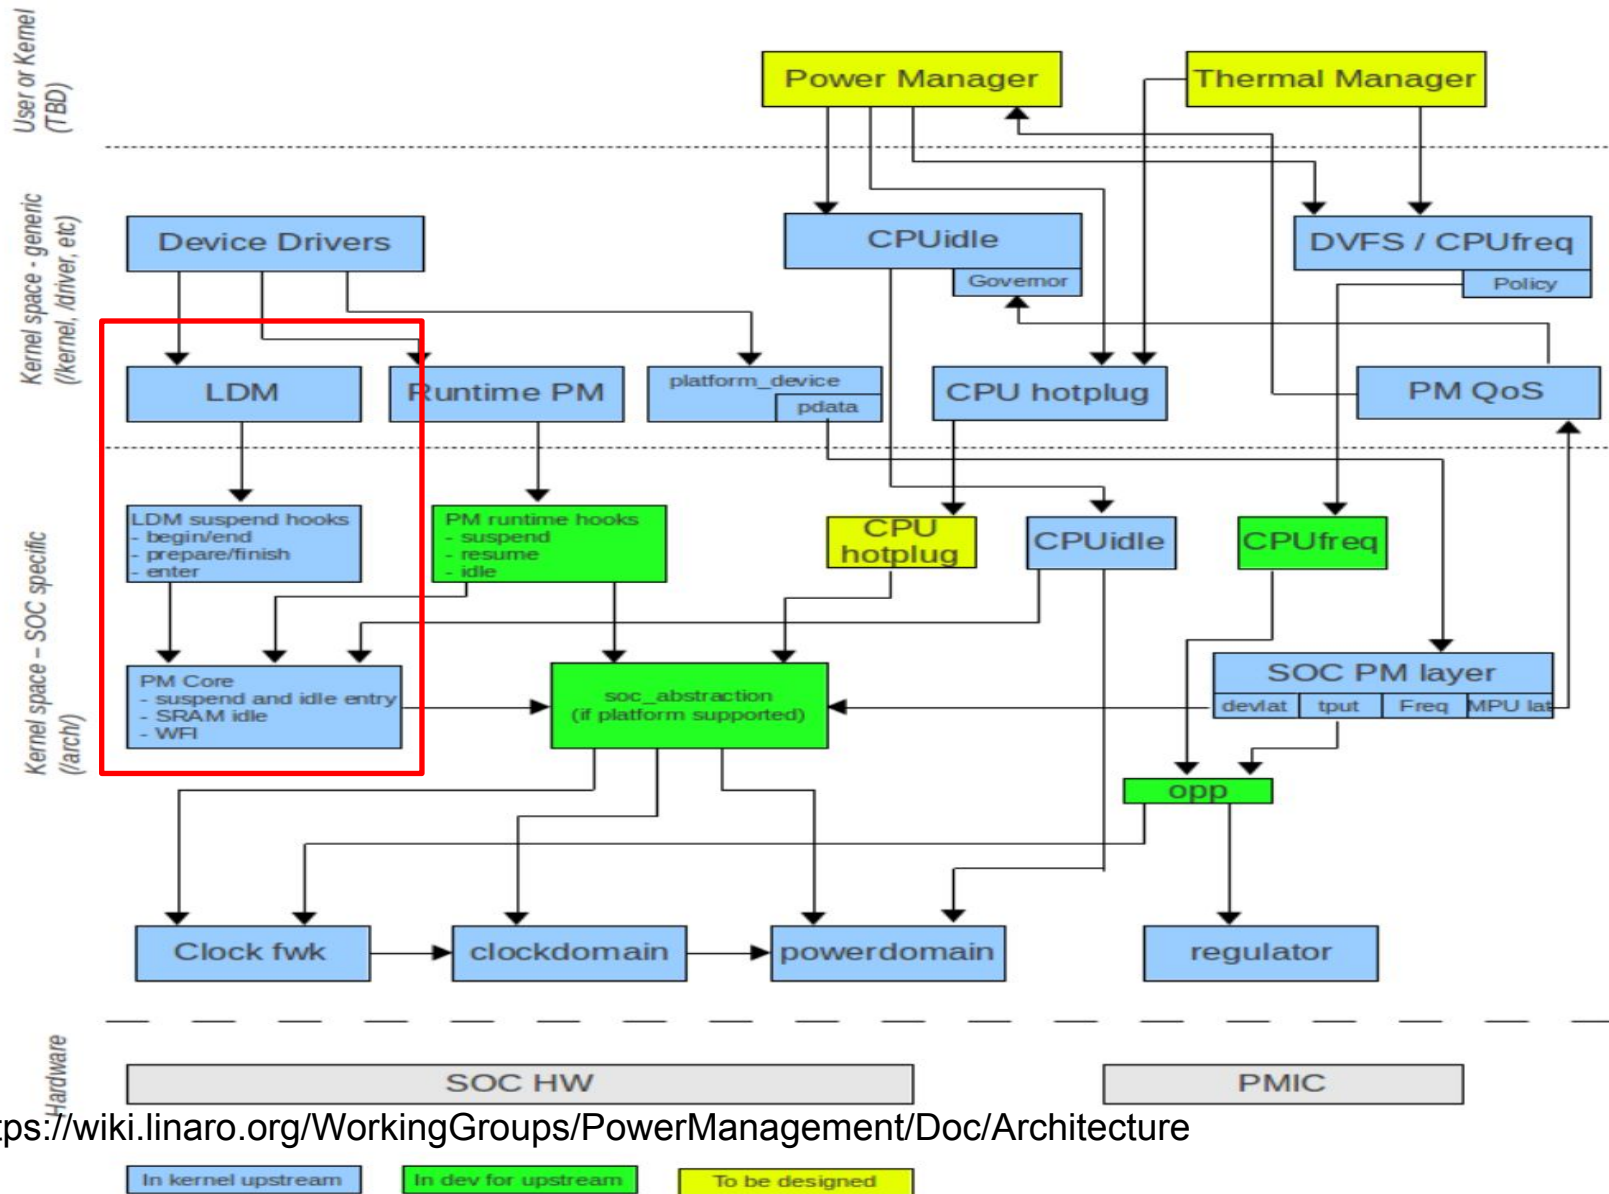
```
static struct device_driver eeepro100_driver = {  
    .name      = "eeepro100",  
    .bus       = &pci_bus_type,  
  
    .probe     = eeepro100_probe,  
    .remove    = eeepro100_remove,  
    .suspend   = eeepro100_suspend,  
    .resume    = eeepro100_resume,  
};
```

# Linux Power Framework





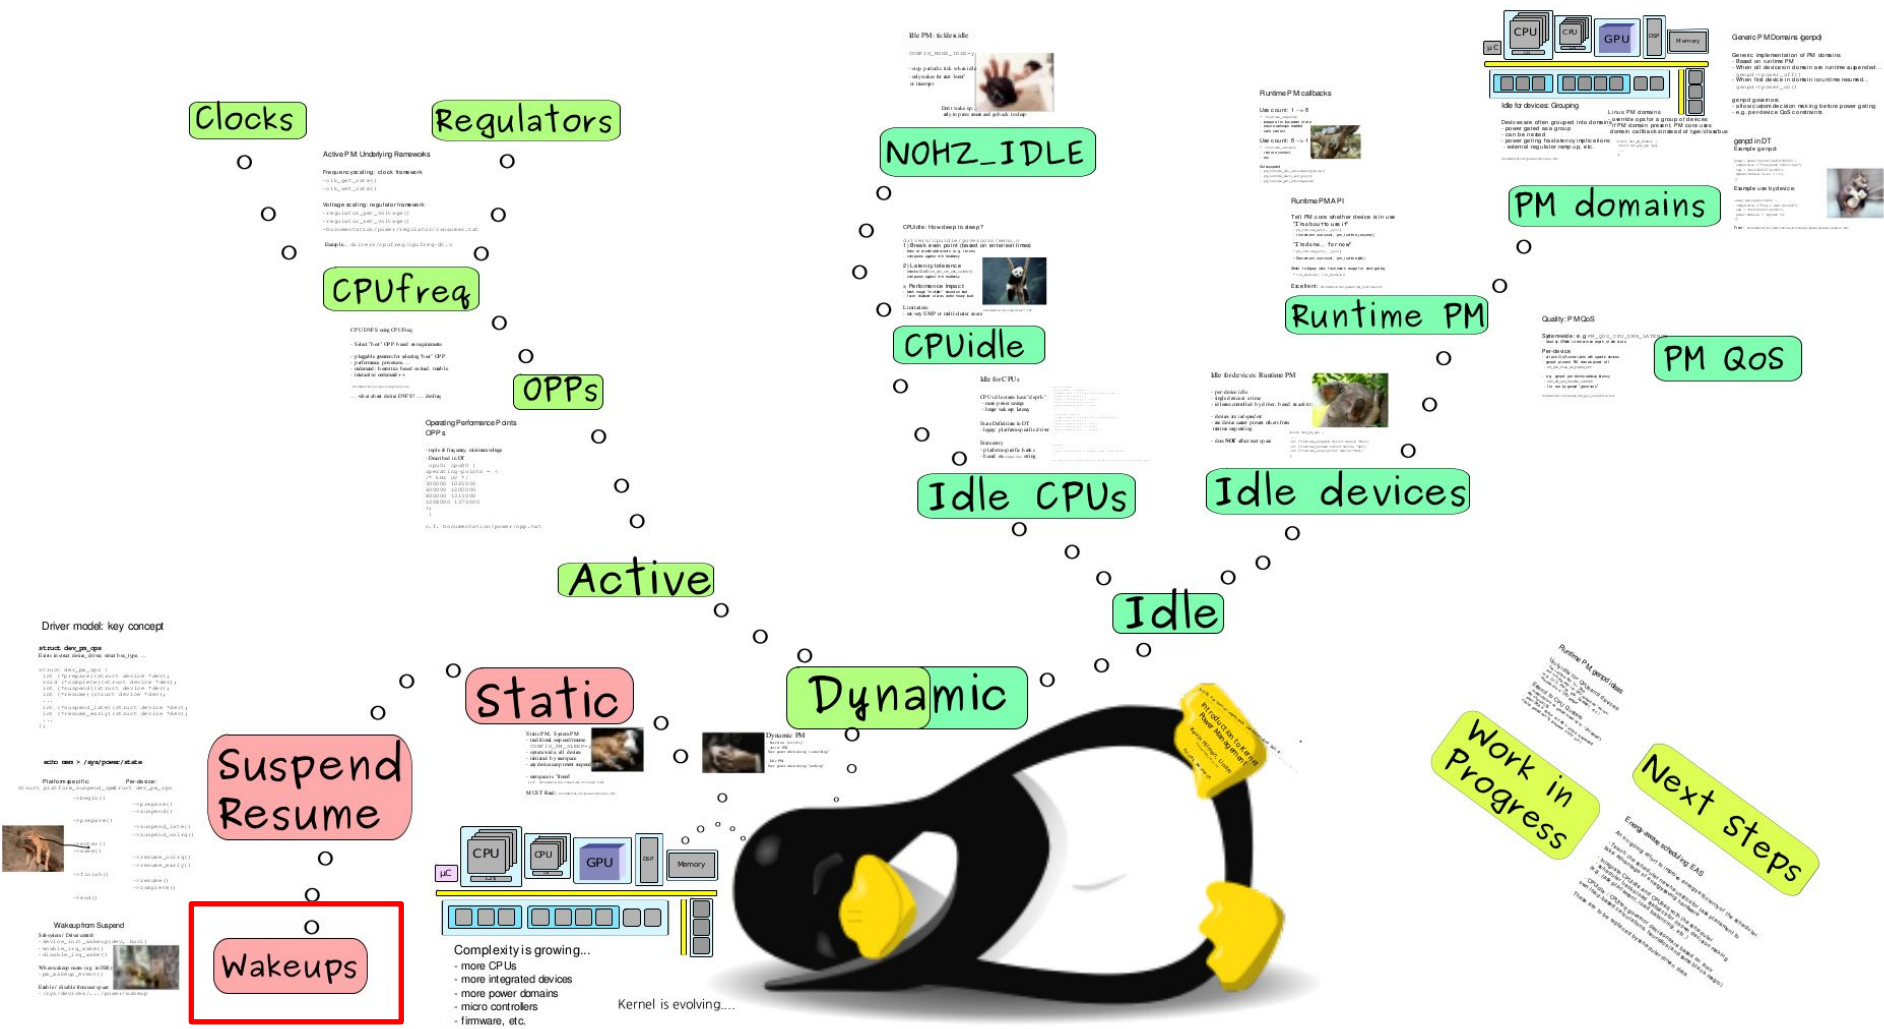
# Linux Driver Model



<https://wiki.linaro.org/WorkingGroups/PowerManagement/Doc/Architecture>

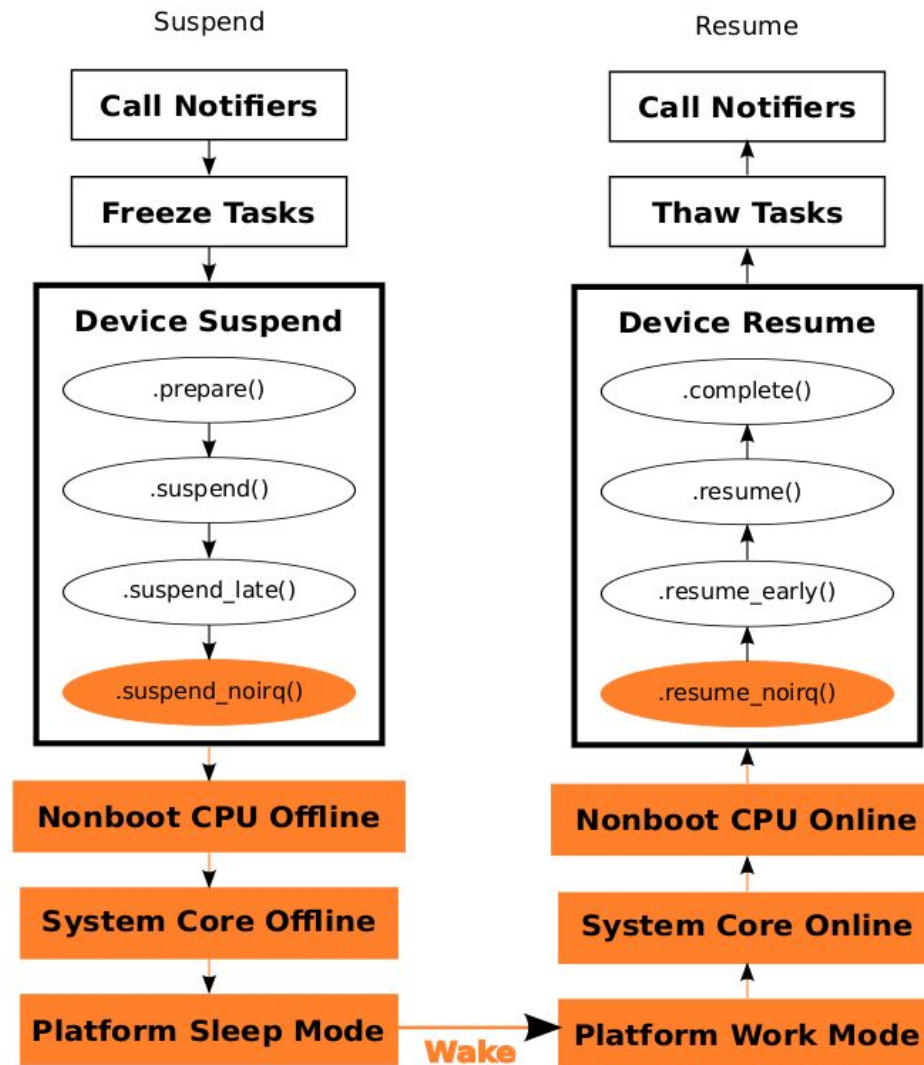


# Tree view - Dynamic and Static PM



Reference from "Introduction to Kernel Power Management" Kevin Hilman, Linaro

# General System Suspend/Resume Control Flow



# Problem

- **How to reduce system idle power?**

# Problem

- **How to reduce system idle power?**
- **What about CPU Hotplugging?**

# Problem

- How to reduce system idle power?
- What about CPU Hotplugging?
  - Expensive latency
- Solution?

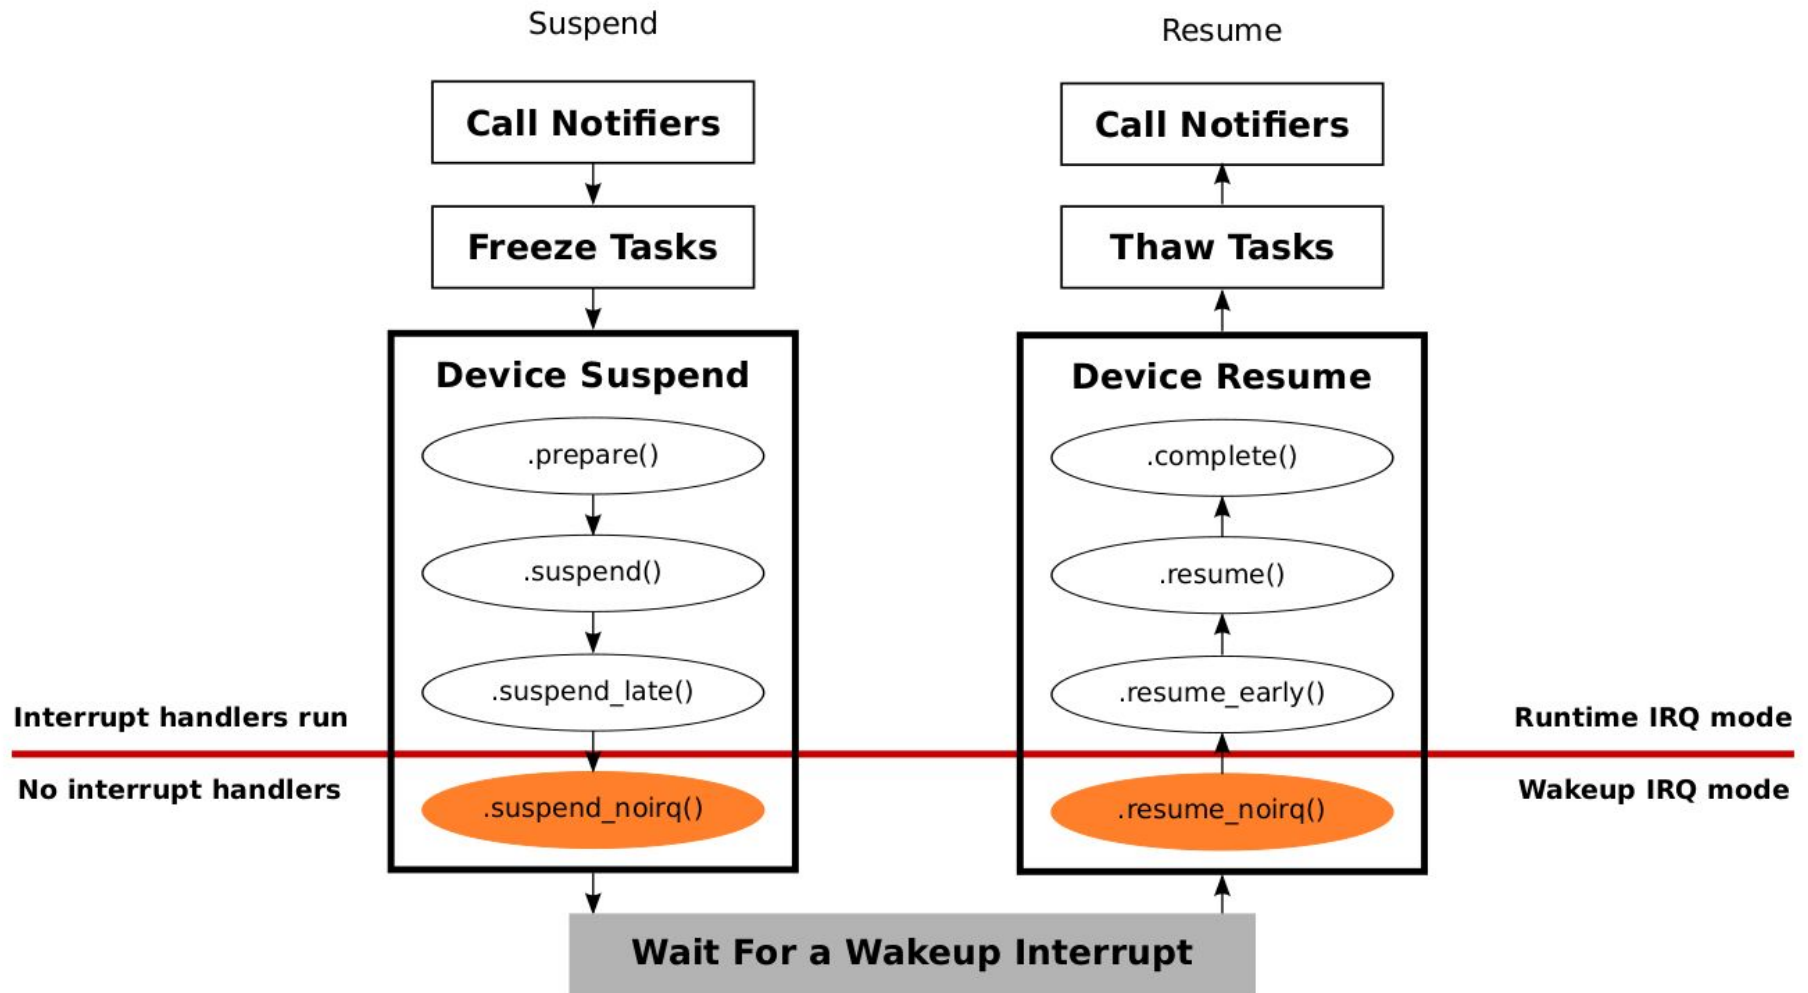
# Problem

- How to reduce system idle power?
- What about CPU Hotplugging?
  - Expensive latency
- Solution?


**S2I(Suspend to Idle)**



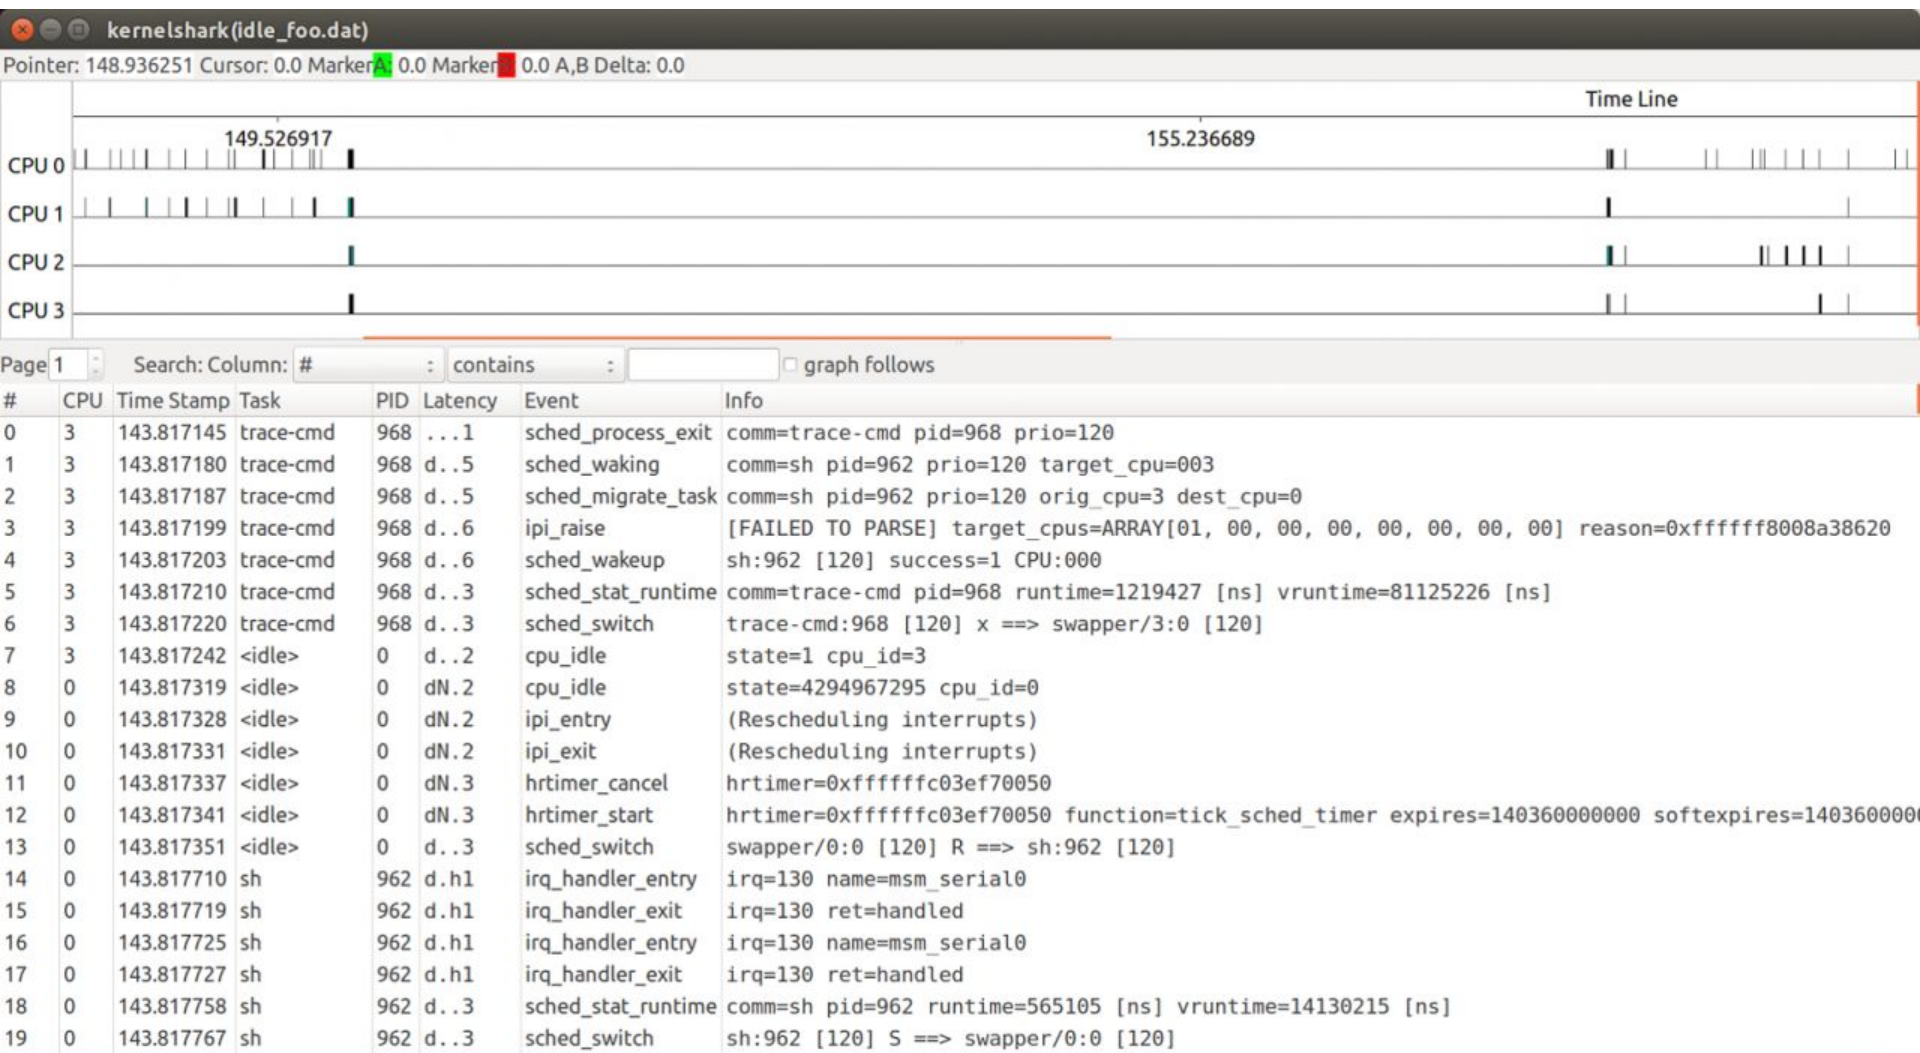
# Suspend-to-Idle



# Suspend-to-Idle

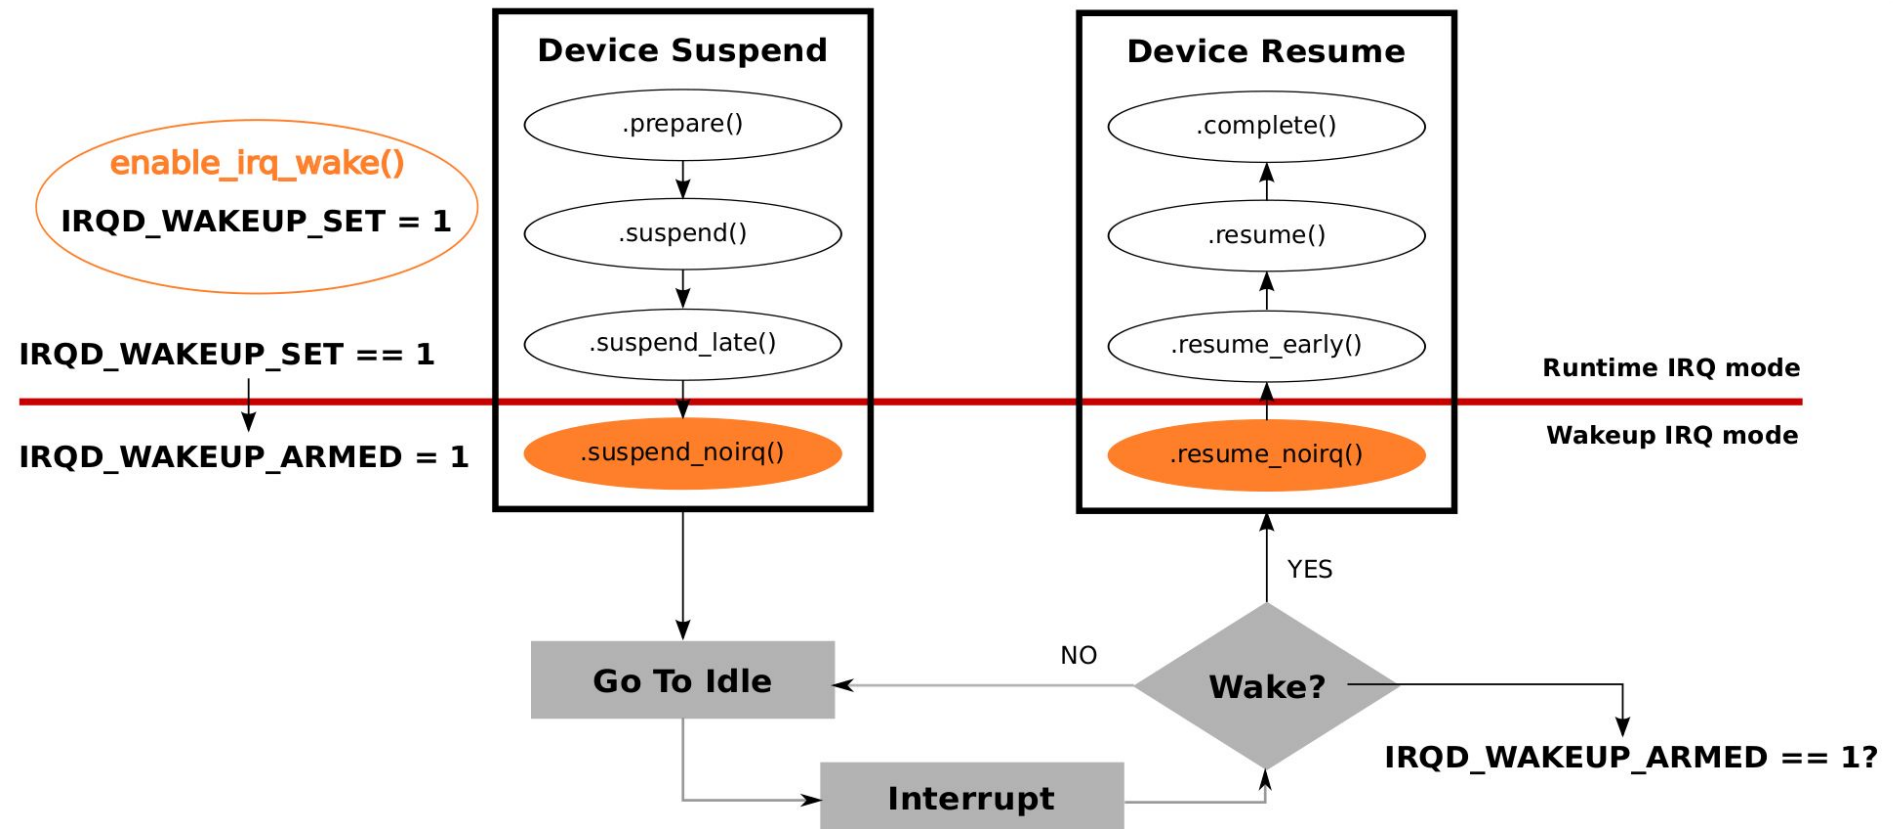


# Suspend-to-Idle



<https://www.linaro.org/blog/suspend-to-idle/>

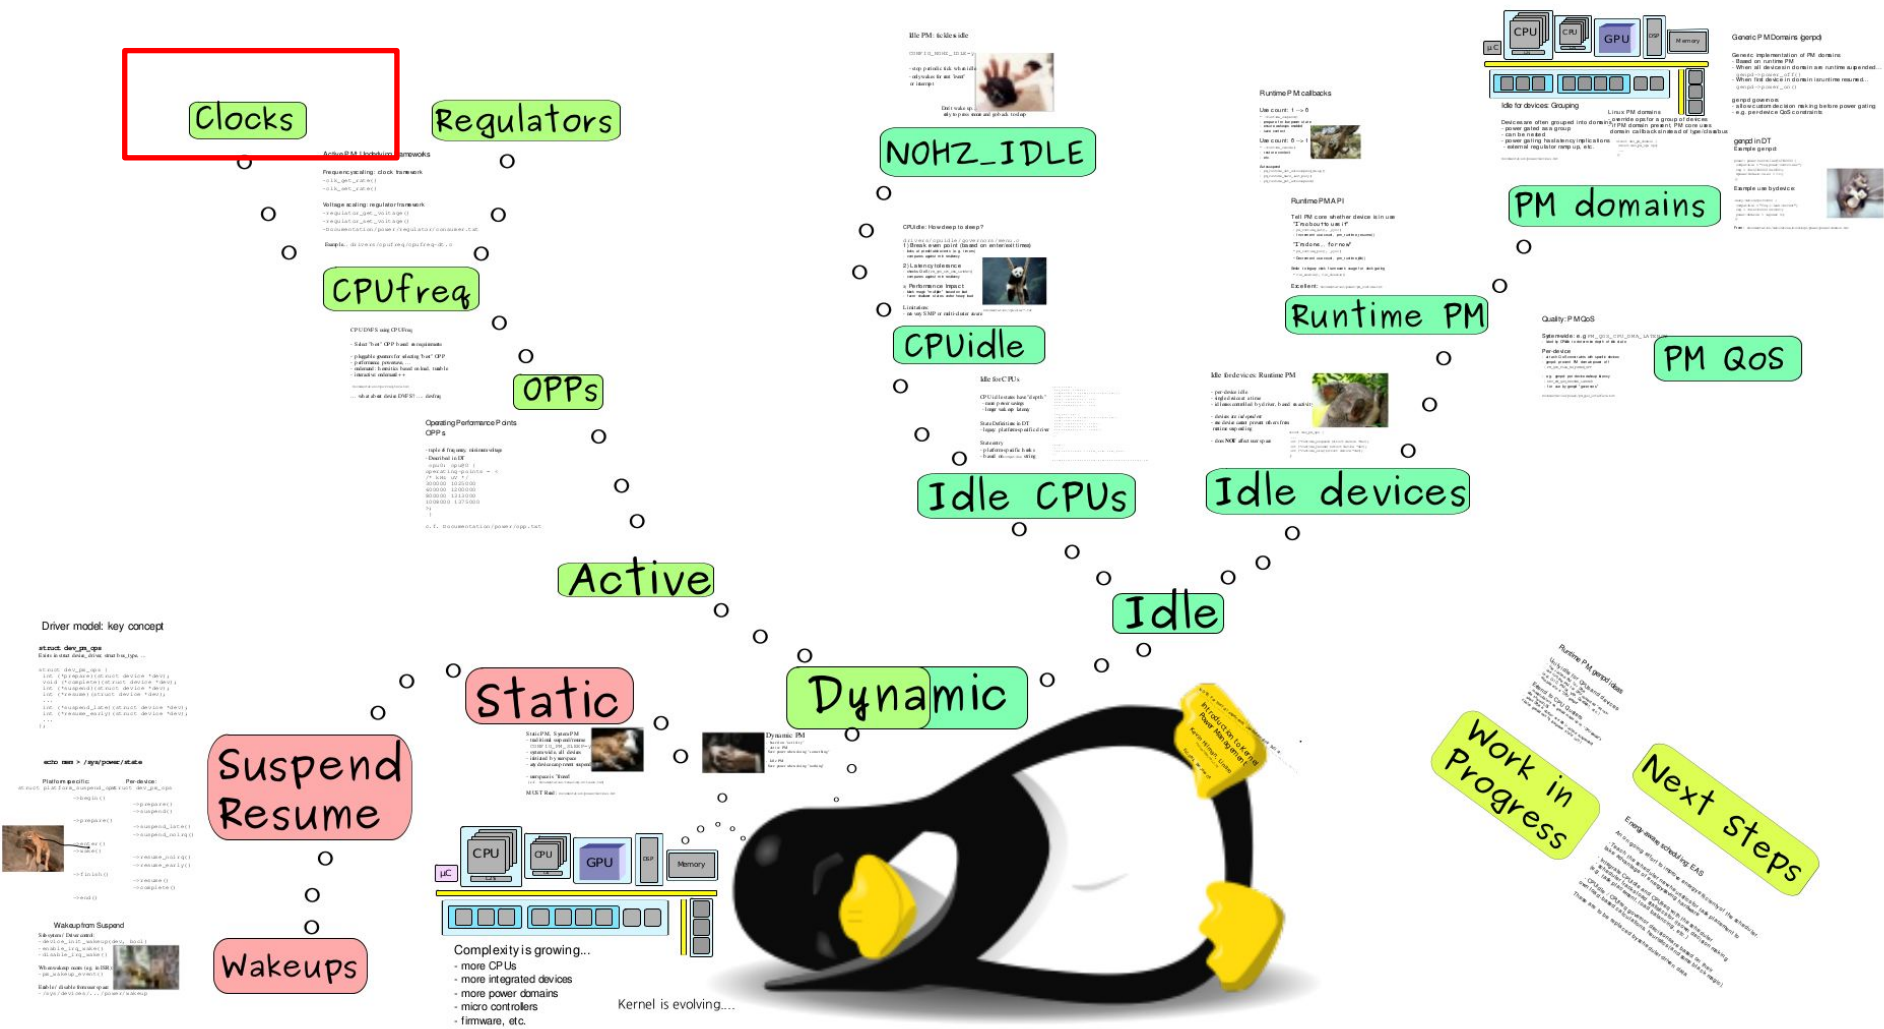
# Suspend-to-Idle Wakeup



“What Is Suspend-to-Idle and How To Make It Work”, Rafael J. Wysocki

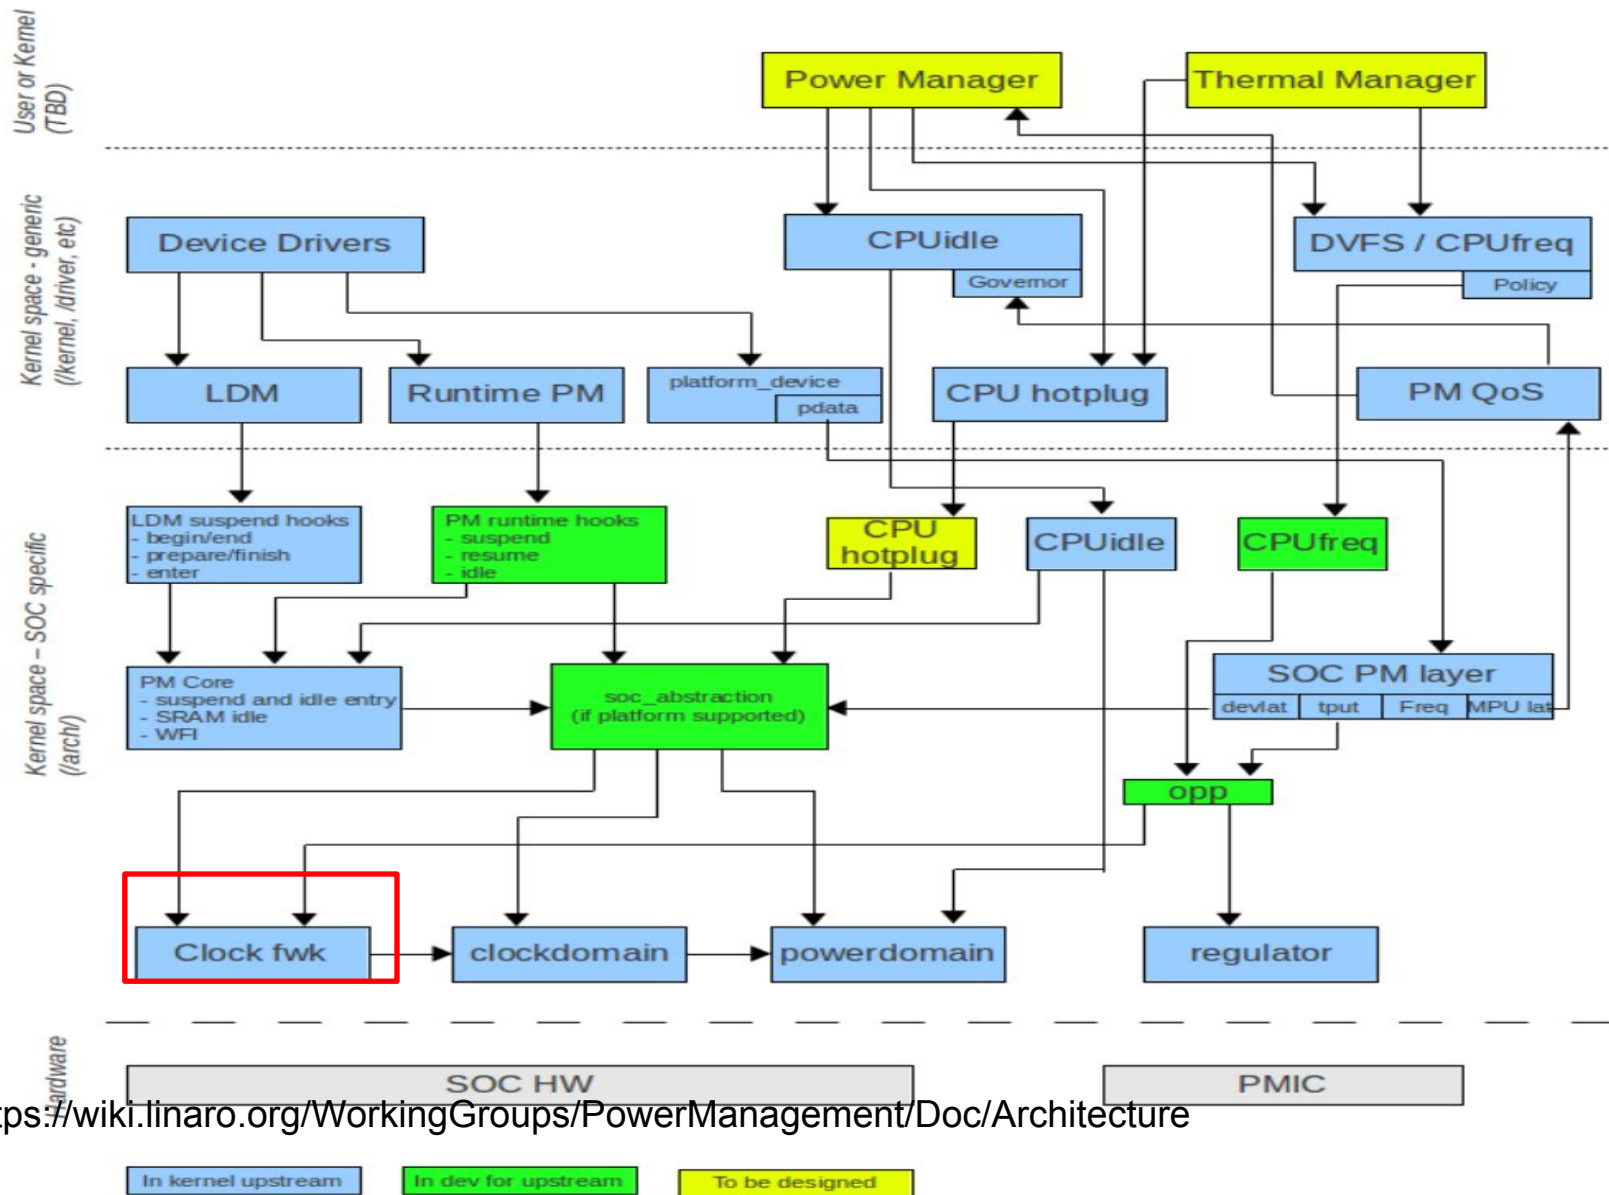


# Tree view - Dynamic and Static PM



Reference from "Introduction to Kernel Power Management" Kevin Hilman, Linaro

# Linux Power Framework

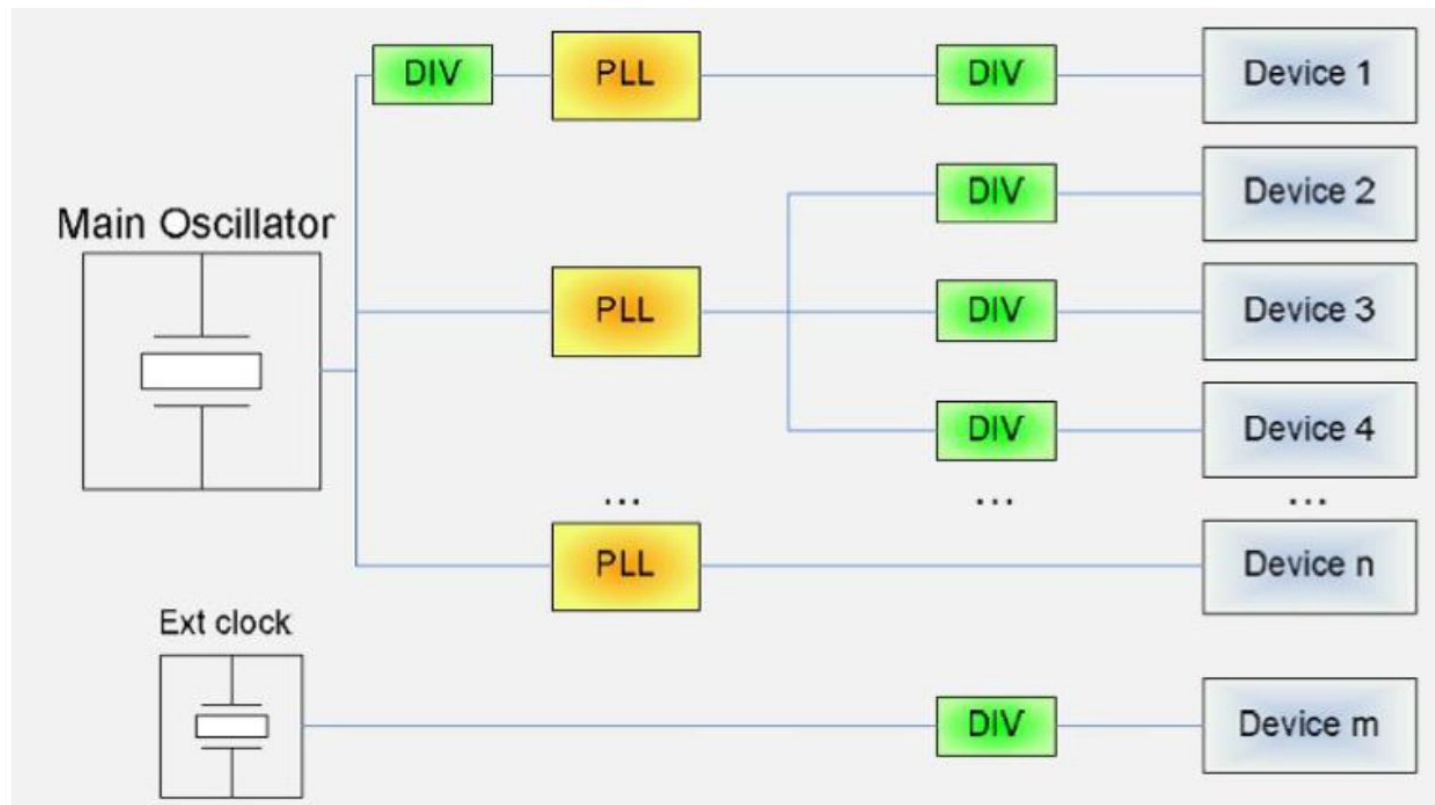


<https://wiki.linaro.org/WorkingGroups/PowerManagement/Doc/Architecture>



# Clock Framework

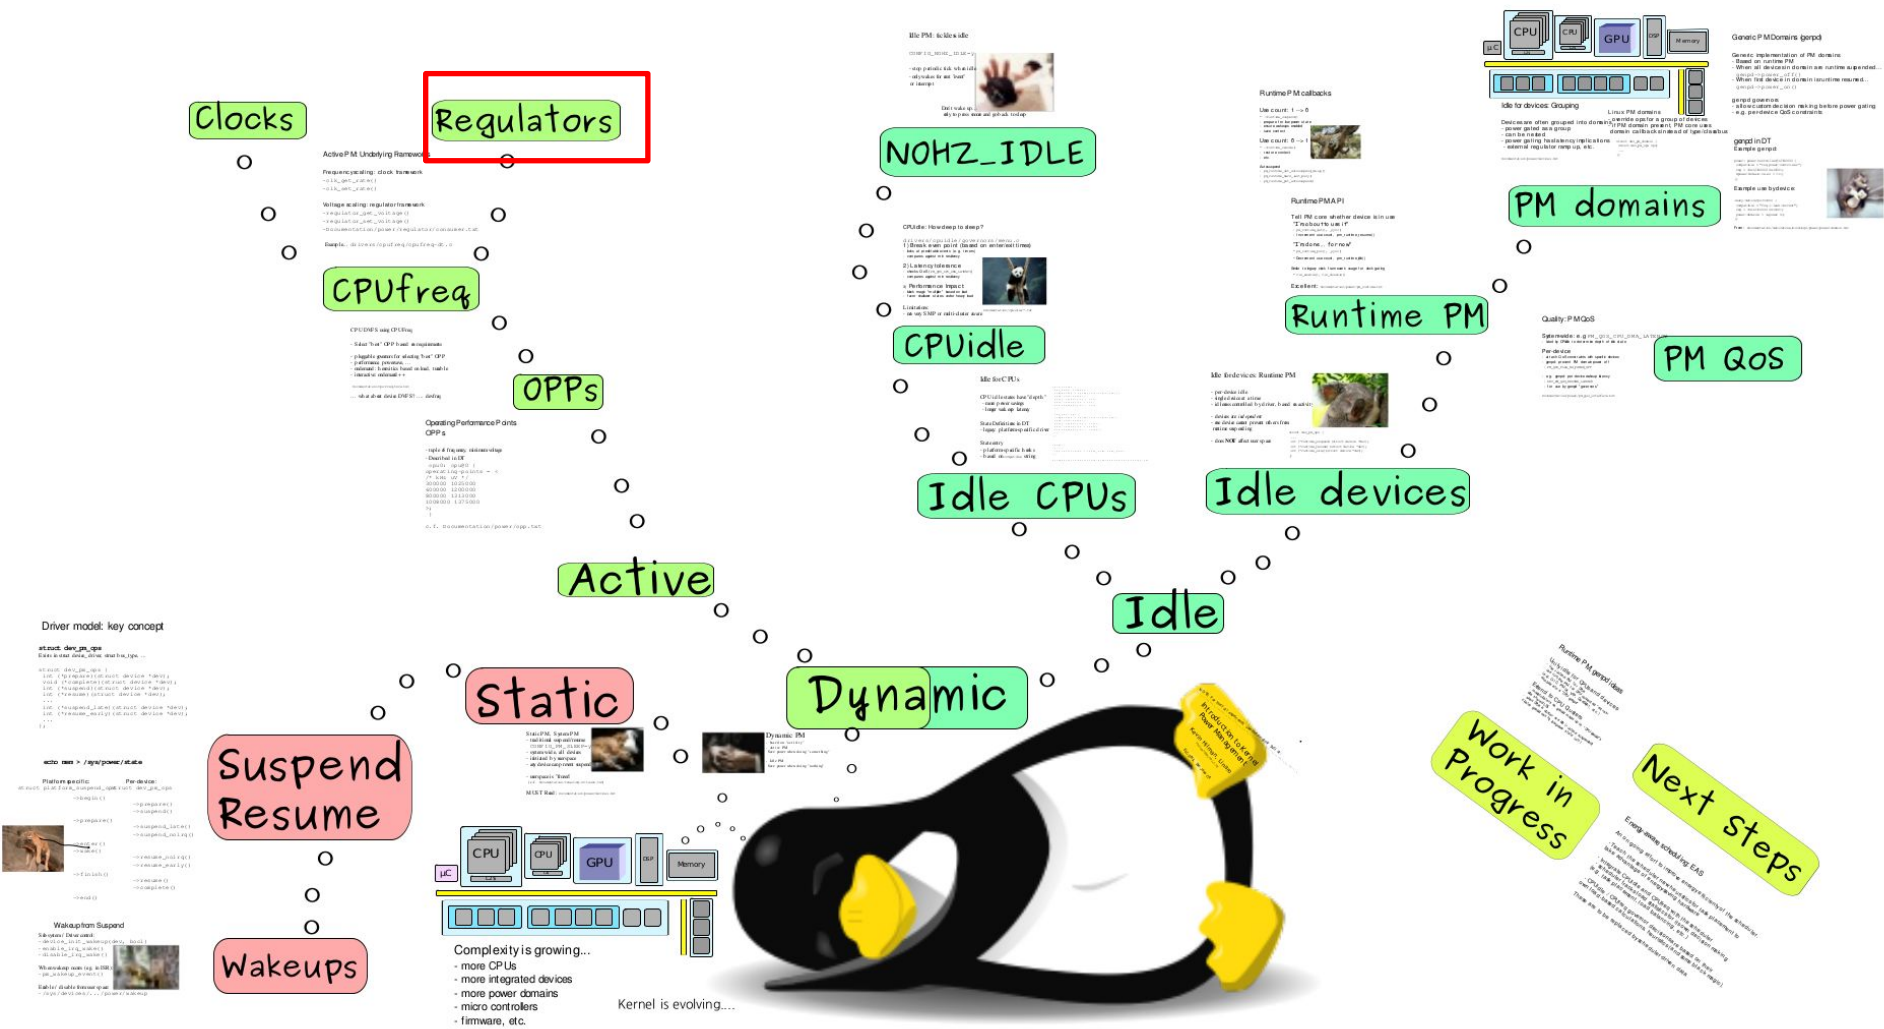
- Clock Domains



# Clock Framework – function list

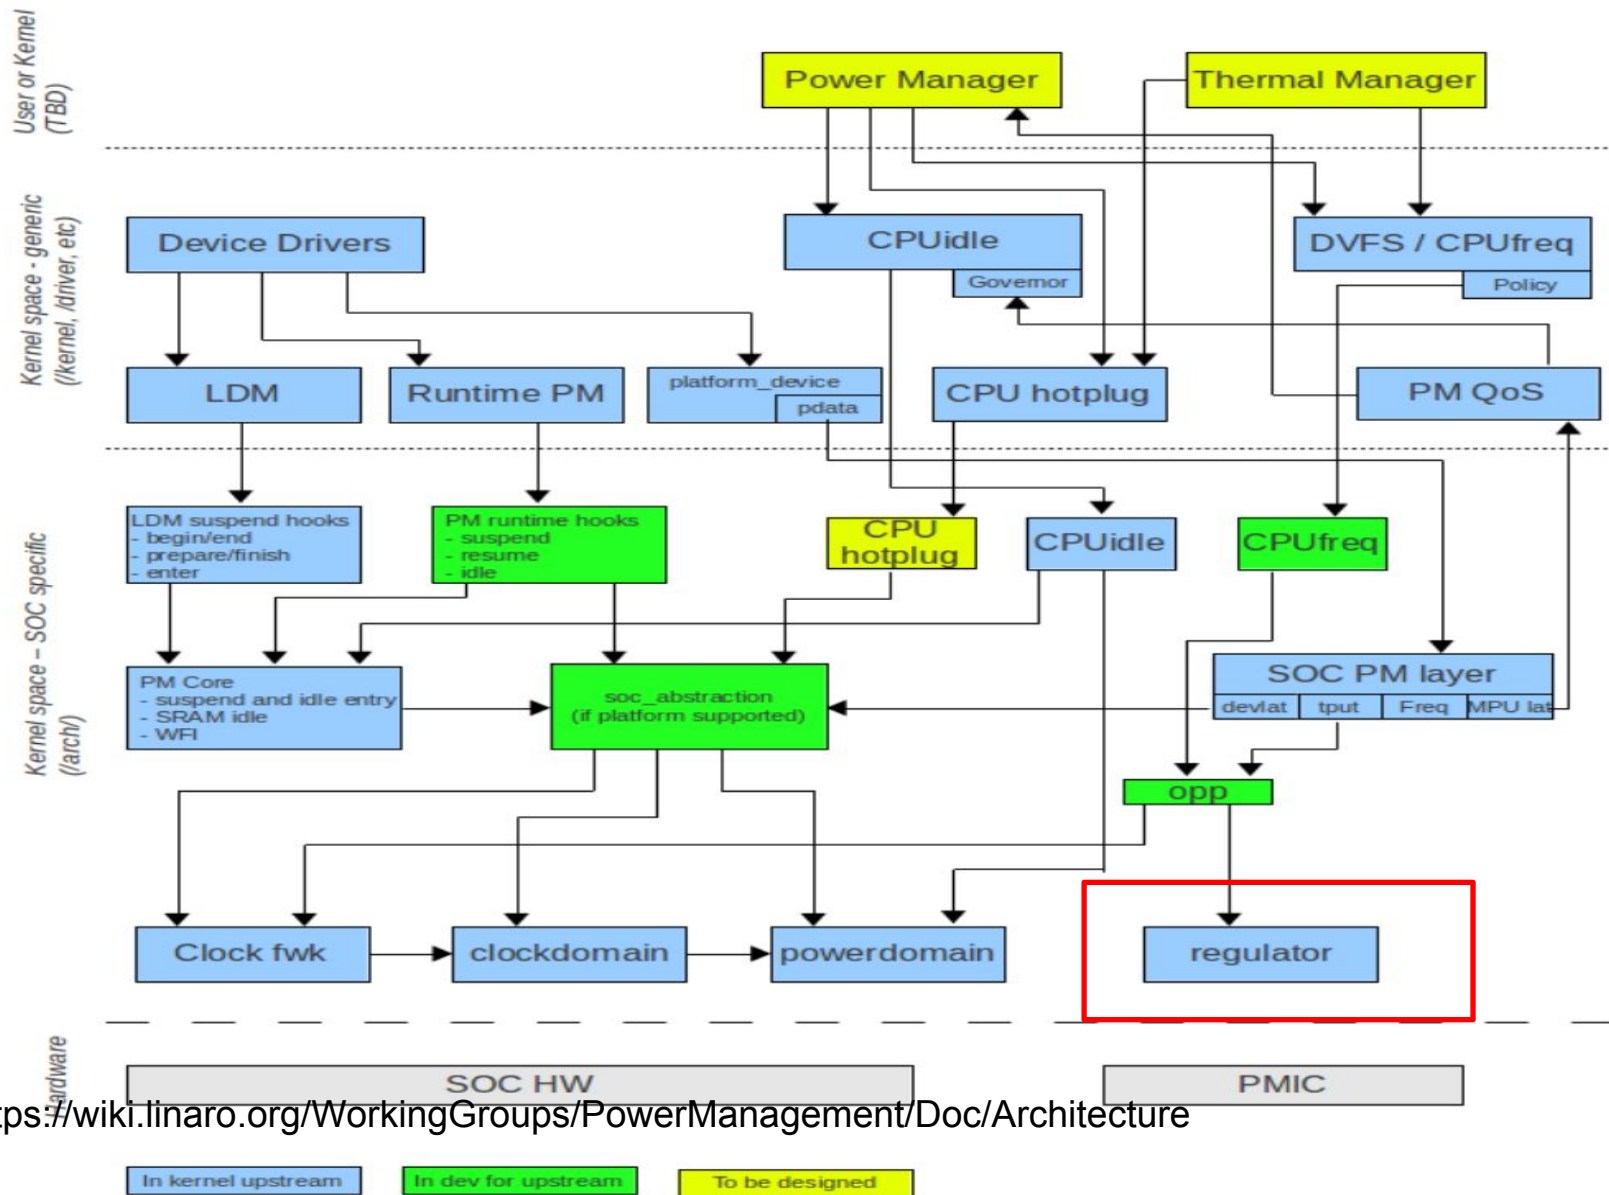
- **clk\_get**
  - Lookup and obtain a reference to a clock producer
- **clk\_enable**
  - Inform the system when the clock source should be running.
- **clk\_disable**
  - Inform the system when the clock source is no longer required.
- **clk\_put**
  - “free” the clock source

# Tree view - Dynamic and Static PM



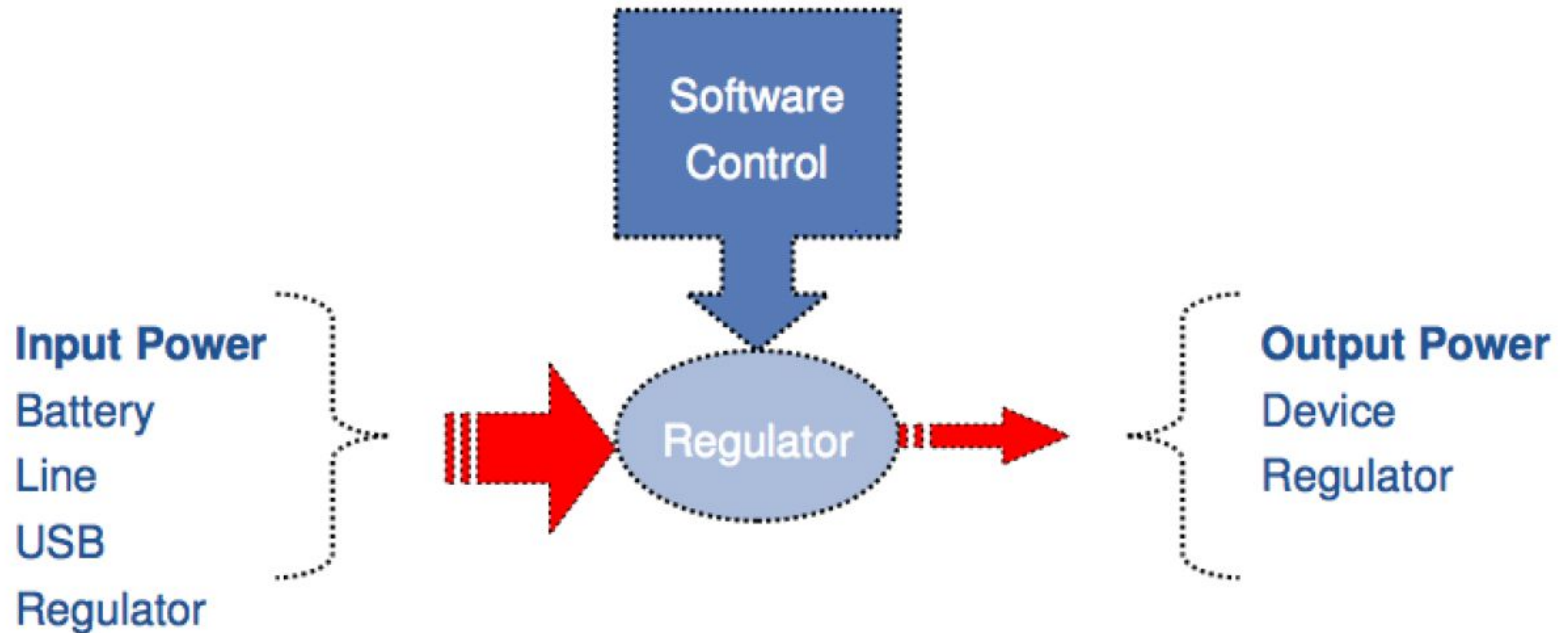
Reference from "Introduction to Kernel Power Management" Kevin Hilman, Linaro

# Linux Power Framework



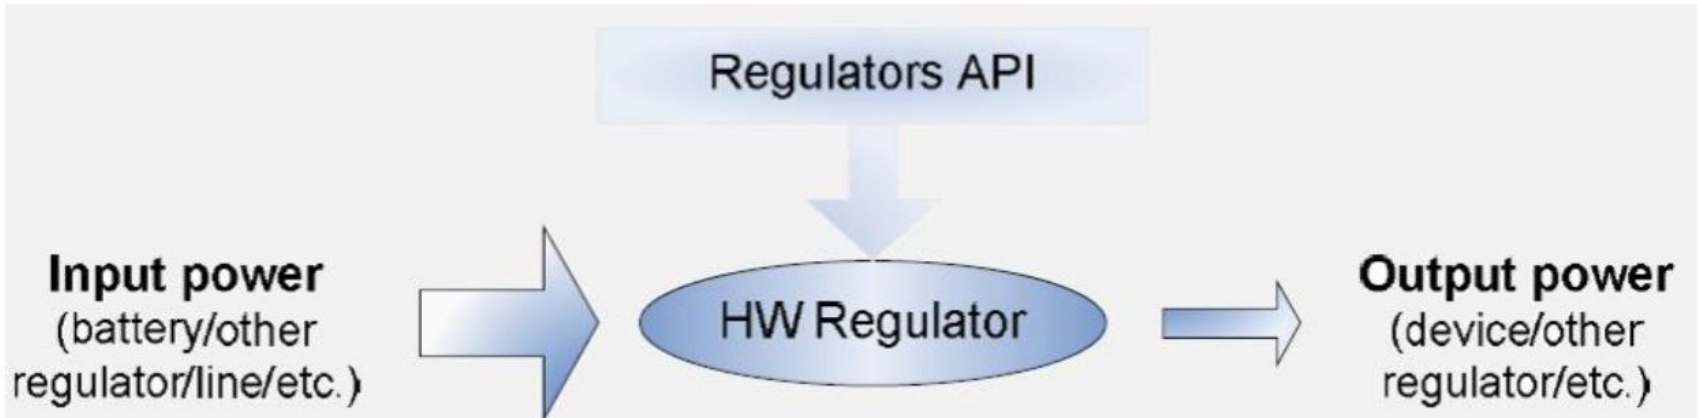
<https://wiki.linaro.org/WorkingGroups/PowerManagement/Doc/Architecture>

# Regulator Framework - 1





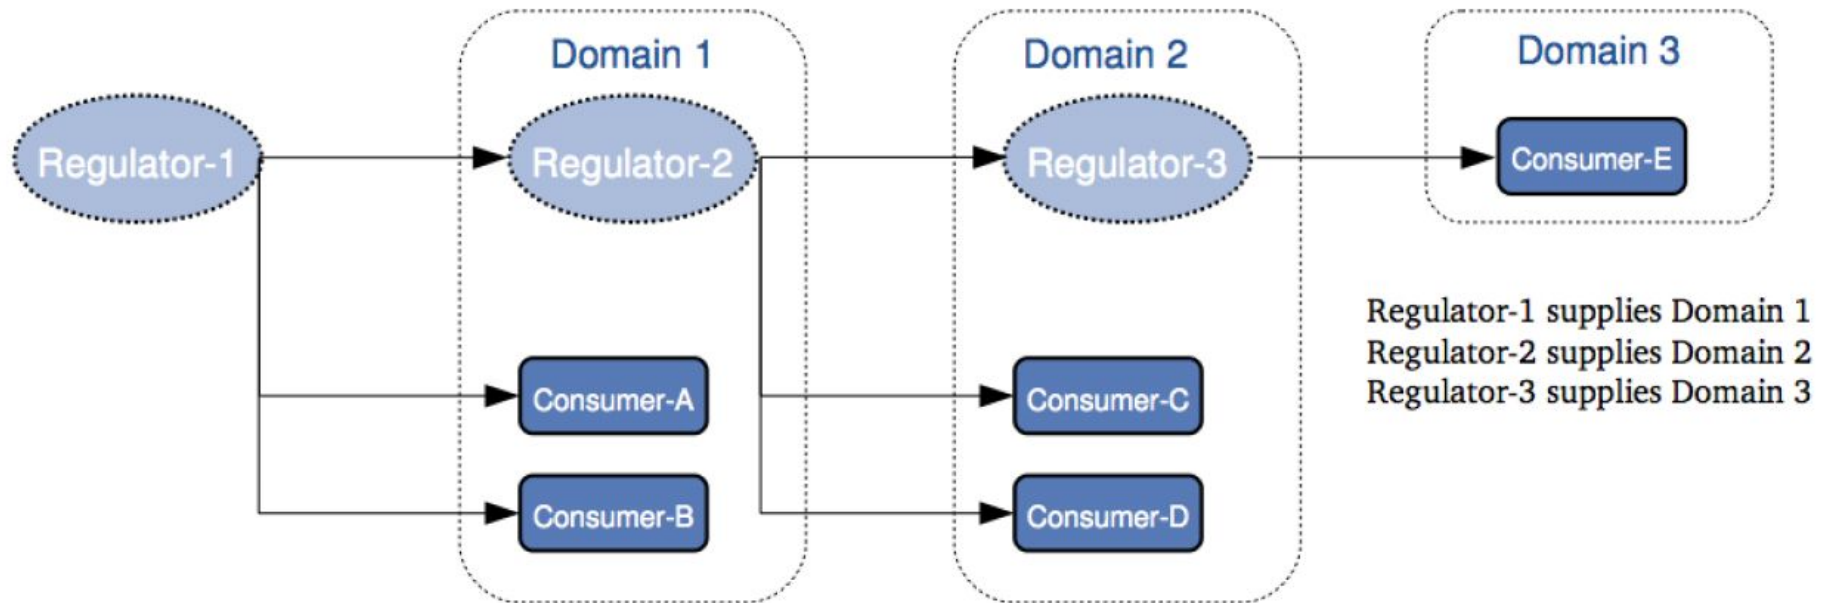
# Regulator Framework - 4



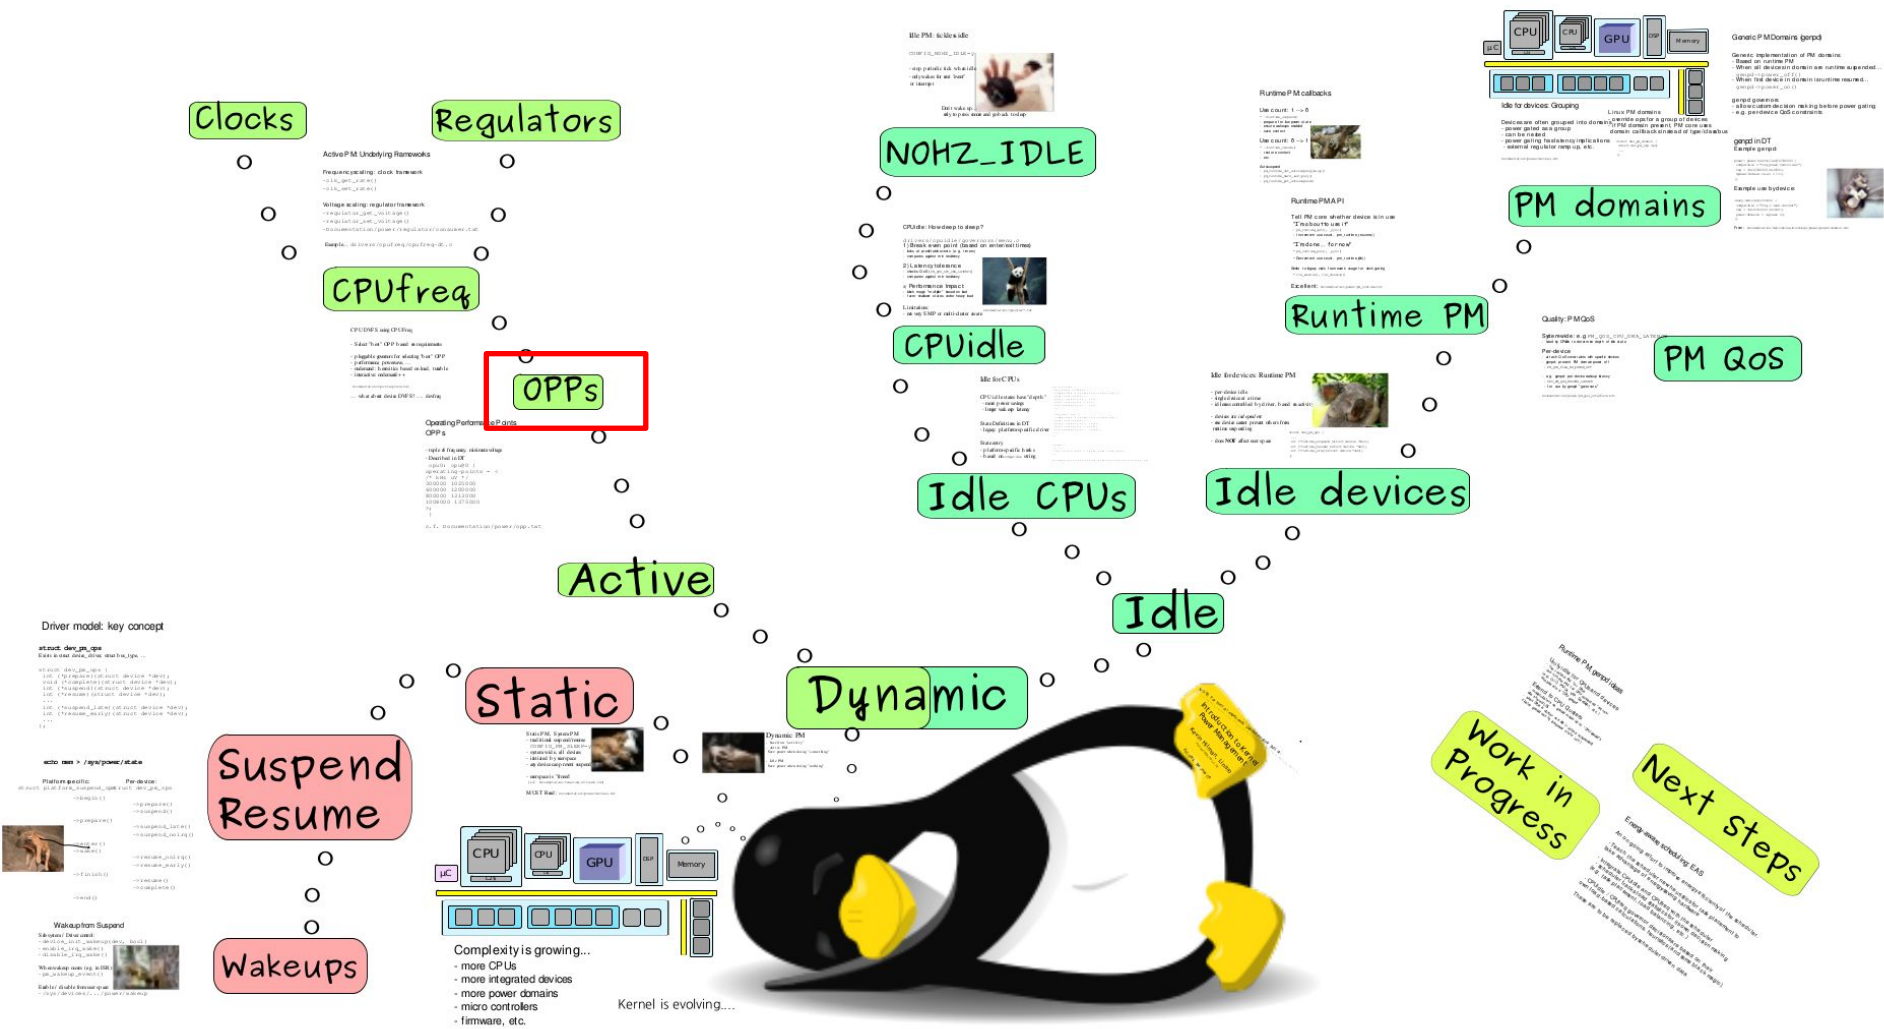
- Regulators have power constraints to protect hardware
- Sysfs interface (/sys/class/regulator/.../)
- Regulator independent abstraction (drivers/regulators/core.c)
- **regulator\_get** - lookup and obtain a reference to a regulator
- **regulator\_put** - "free" the regulator source
- **regulator\_enable** - enable regulator output
- **regulator\_disable** - disable regulator output



# Regulator Framework - 2

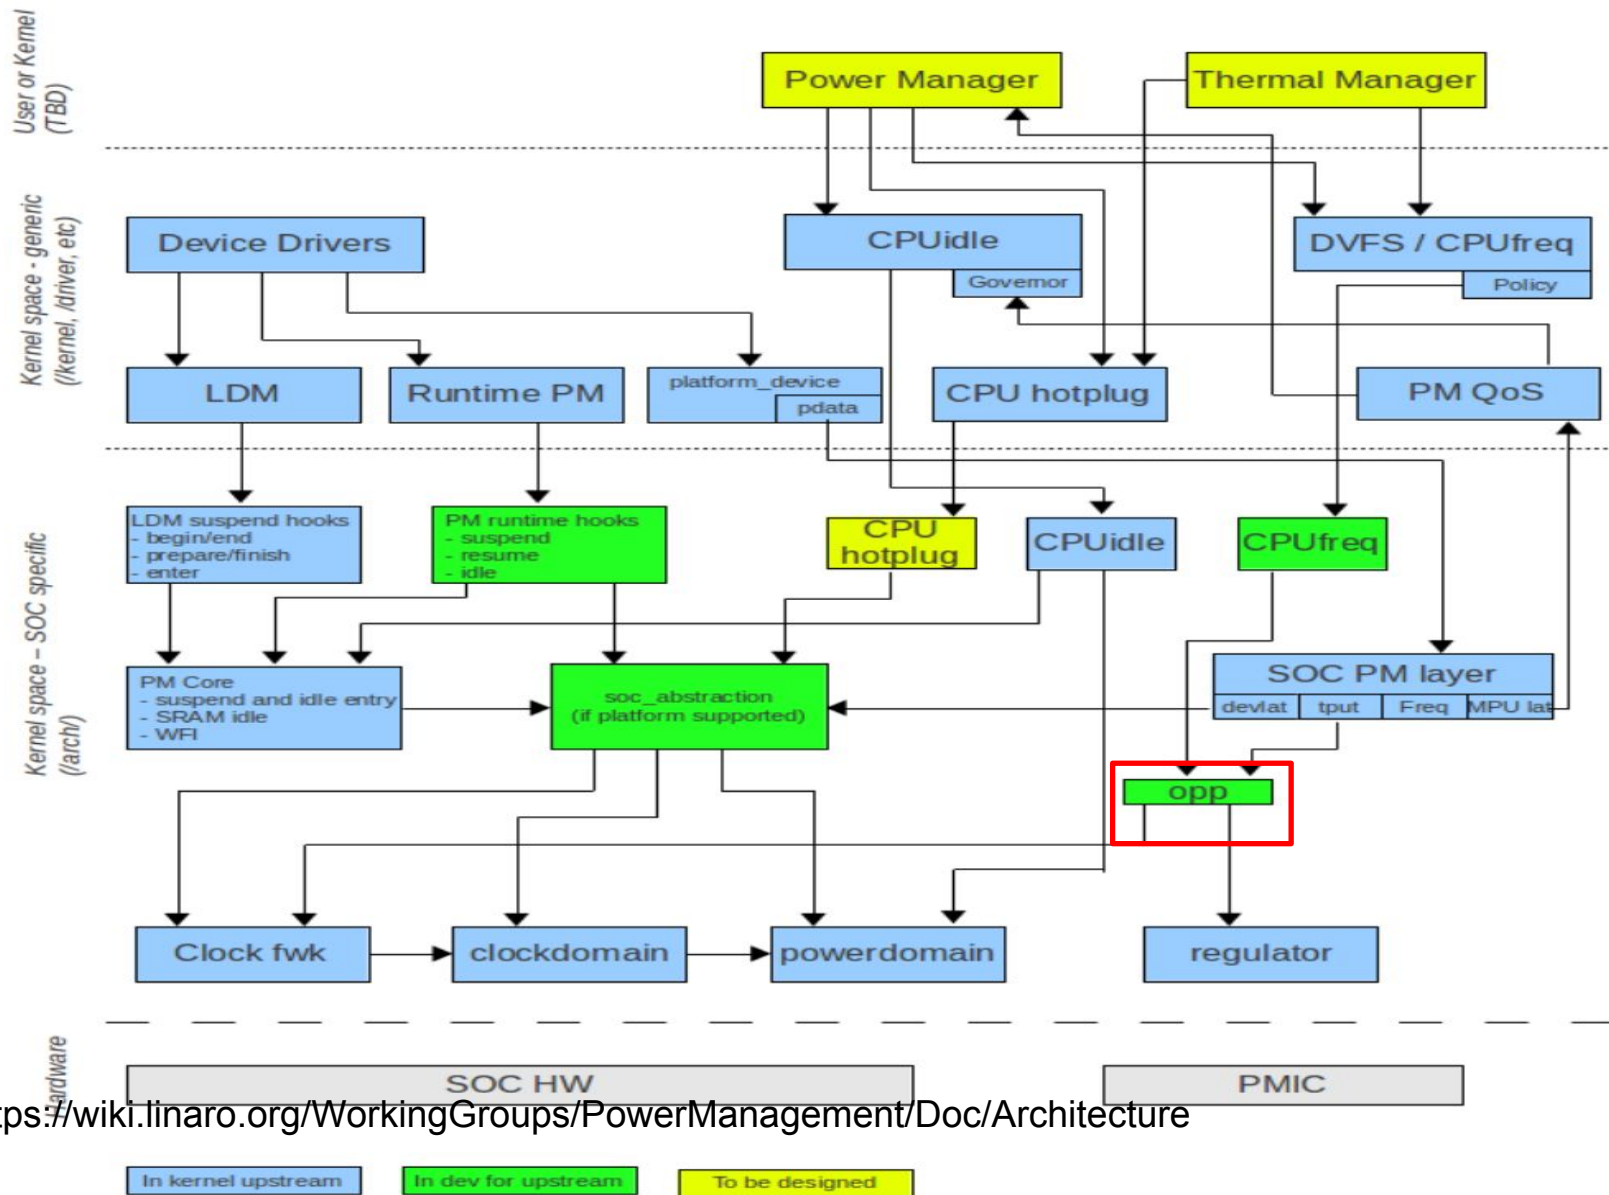


# Tree view - Dynamic and Static PM



Reference from "Introduction to Kernel Power Management" Kevin Hilman, Linaro

# Linux Power Framework

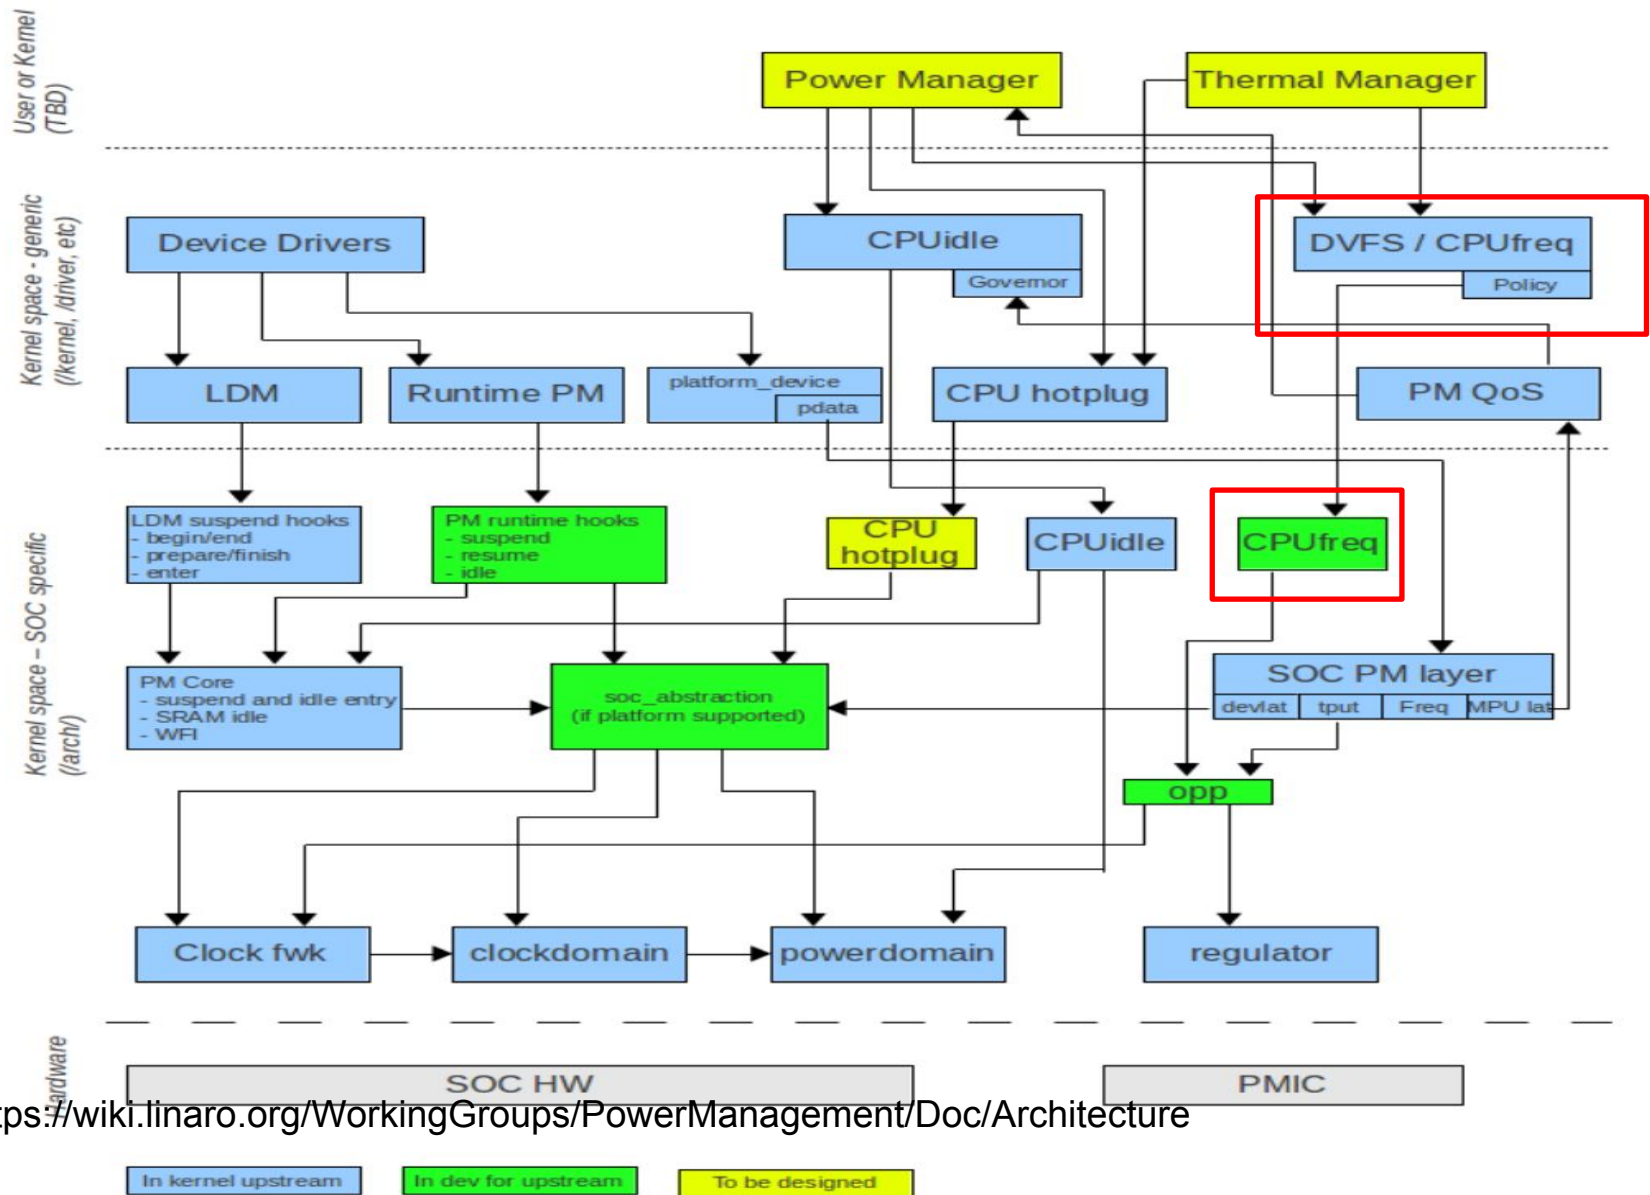


# OPP - Operating Performance Point

- Tuple of frequency, minimum voltage.
- Described in DT

```
cpu0: cpu@0 {  
    operating-points = <  
    /* kHz uV */  
    300000 1025000  
    600000 1200000  
    800000 1313000  
    1008000 1375000  
    >;  
}
```

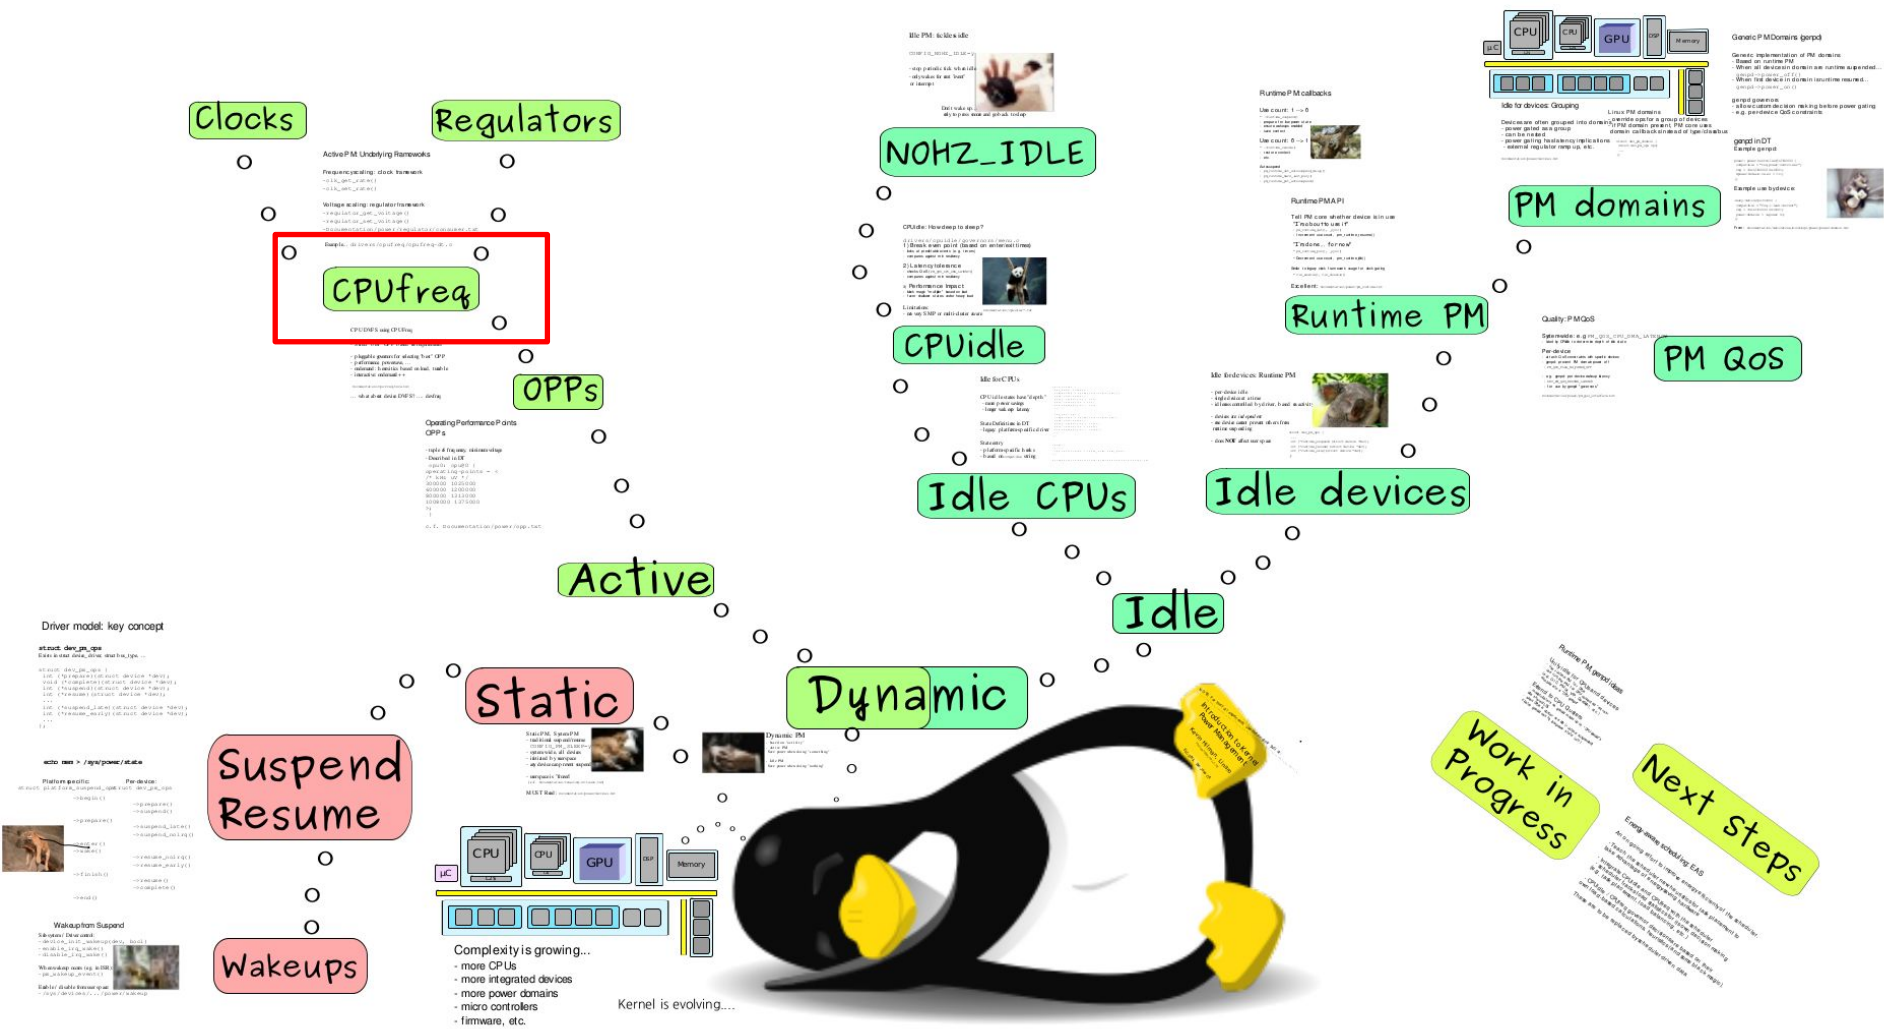
# Linux Power Framework



<https://wiki.linaro.org/WorkingGroups/PowerManagement/Doc/Architecture>



# Tree view - Dynamic and Static PM



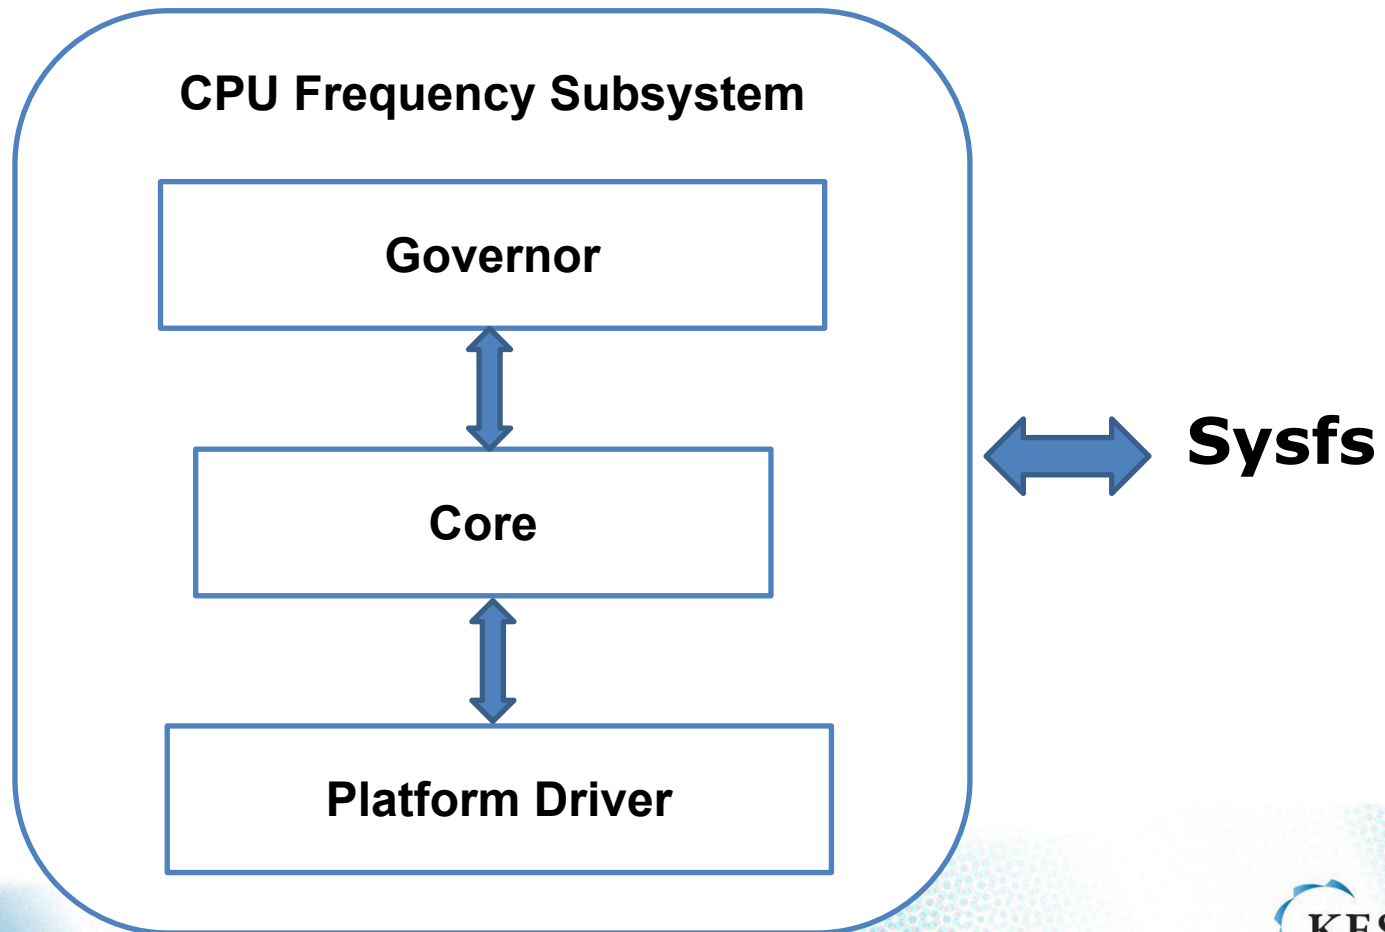
Reference from "Introduction to Kernel Power Management" Kevin Hilman, Linaro



# **CPU Frequency Governor 4.5 and Earlier version of Linux -**

# CPU Frequency Subsystem

- Linux CPU frequency subsystem manages CPU frequency scaling.



# CPU Frequency Subsystem


- **Comprise three parts**
  - Governor
    - Implemented CPU frequency scaling policy
    - Governors in Linux Kernel
      - Performance, Powersave, Userspace, Ondemand, Conservative, Interactive
  - Core
    - Provide generic interface for governor and platform driver to communicate with each other
  - Platform driver
    - SoC-chip specific cpu frequency scaling driver
      - Thermal throttle, DVFS, and so on

# Governors In Linux Kernel

- **Powersave governor**
  - Set frequency statically to lowest available frequency.
- **Performance governor**
  - Set frequency statically to highest available frequency.
- **Userspace governor**
  - Set frequency manually.
- **Ondemand governor**
  - Scale frequency based on CPU utilization.
- **Interactive governor**
  - More responsive to interactive workloads.




# Governors In Linux Kernel



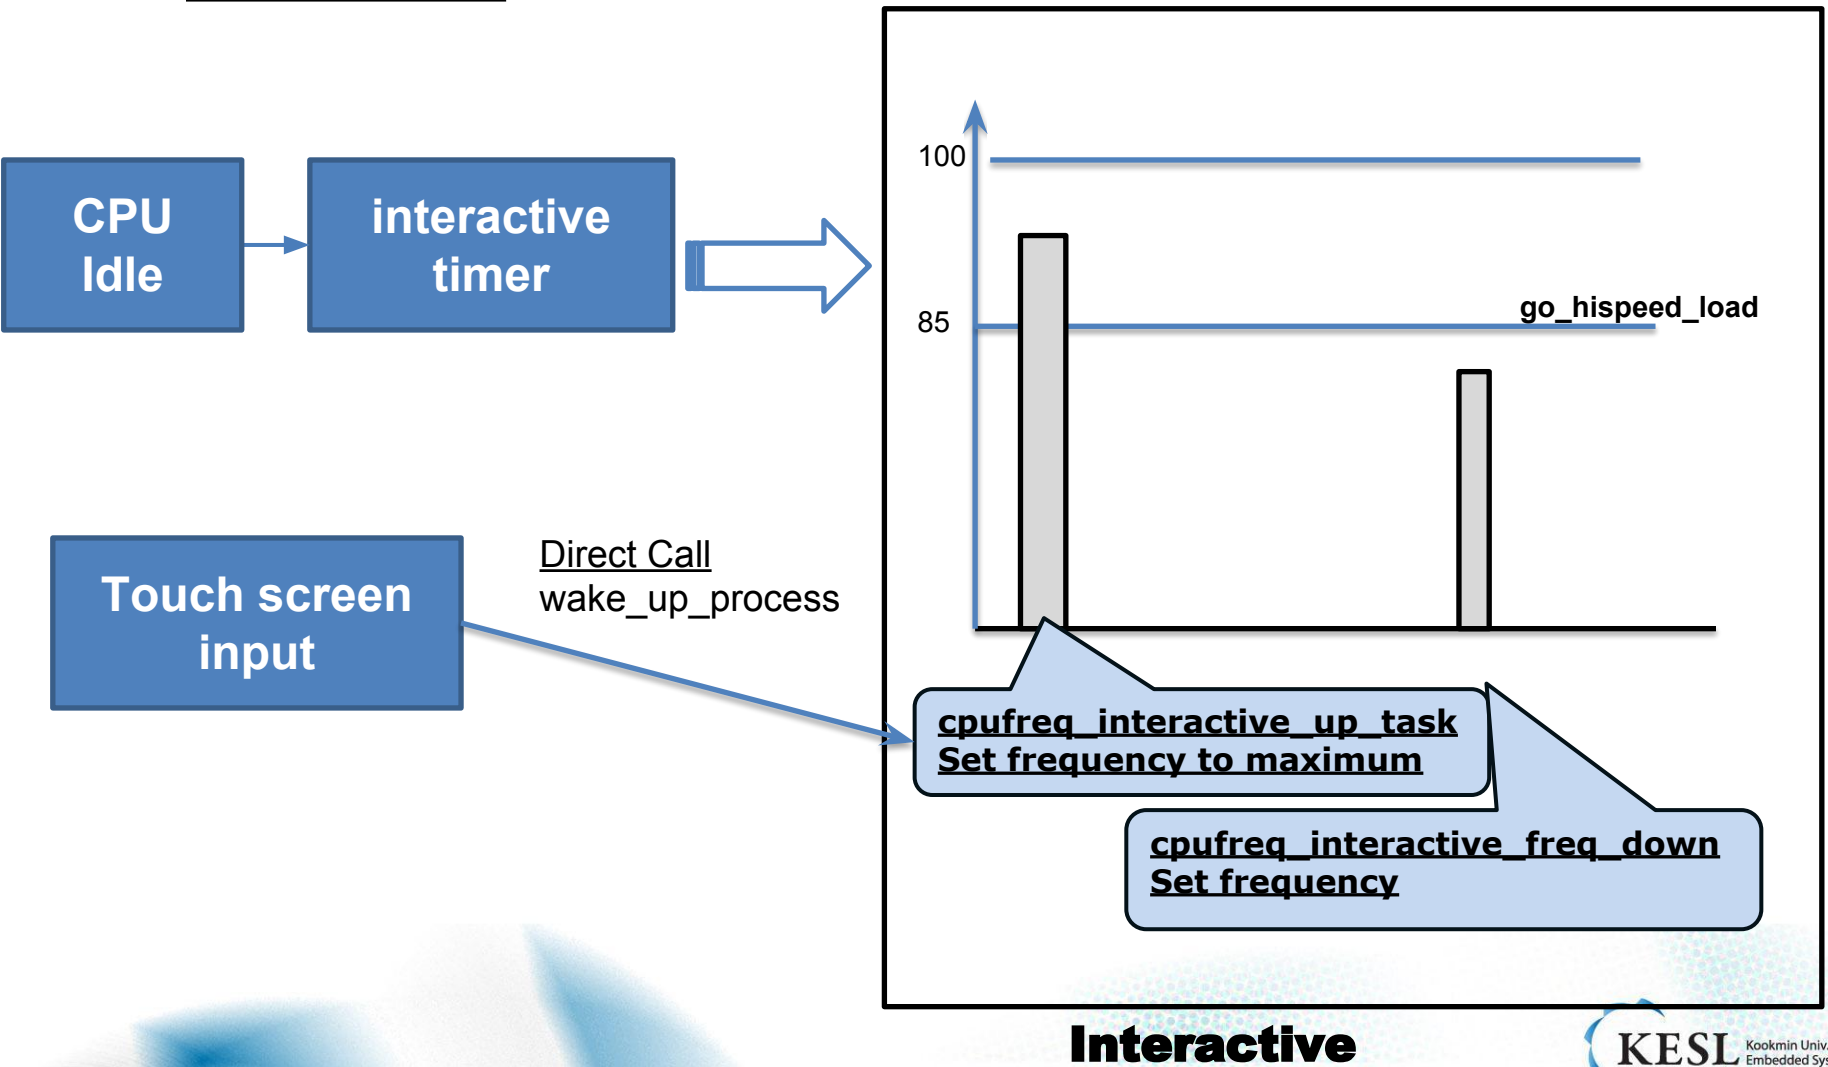
"CPUFreq and The Scheduler Revolution in CPU Power Management", Rafael J. Wysocki

# DVFS in Linux (cpufreq)

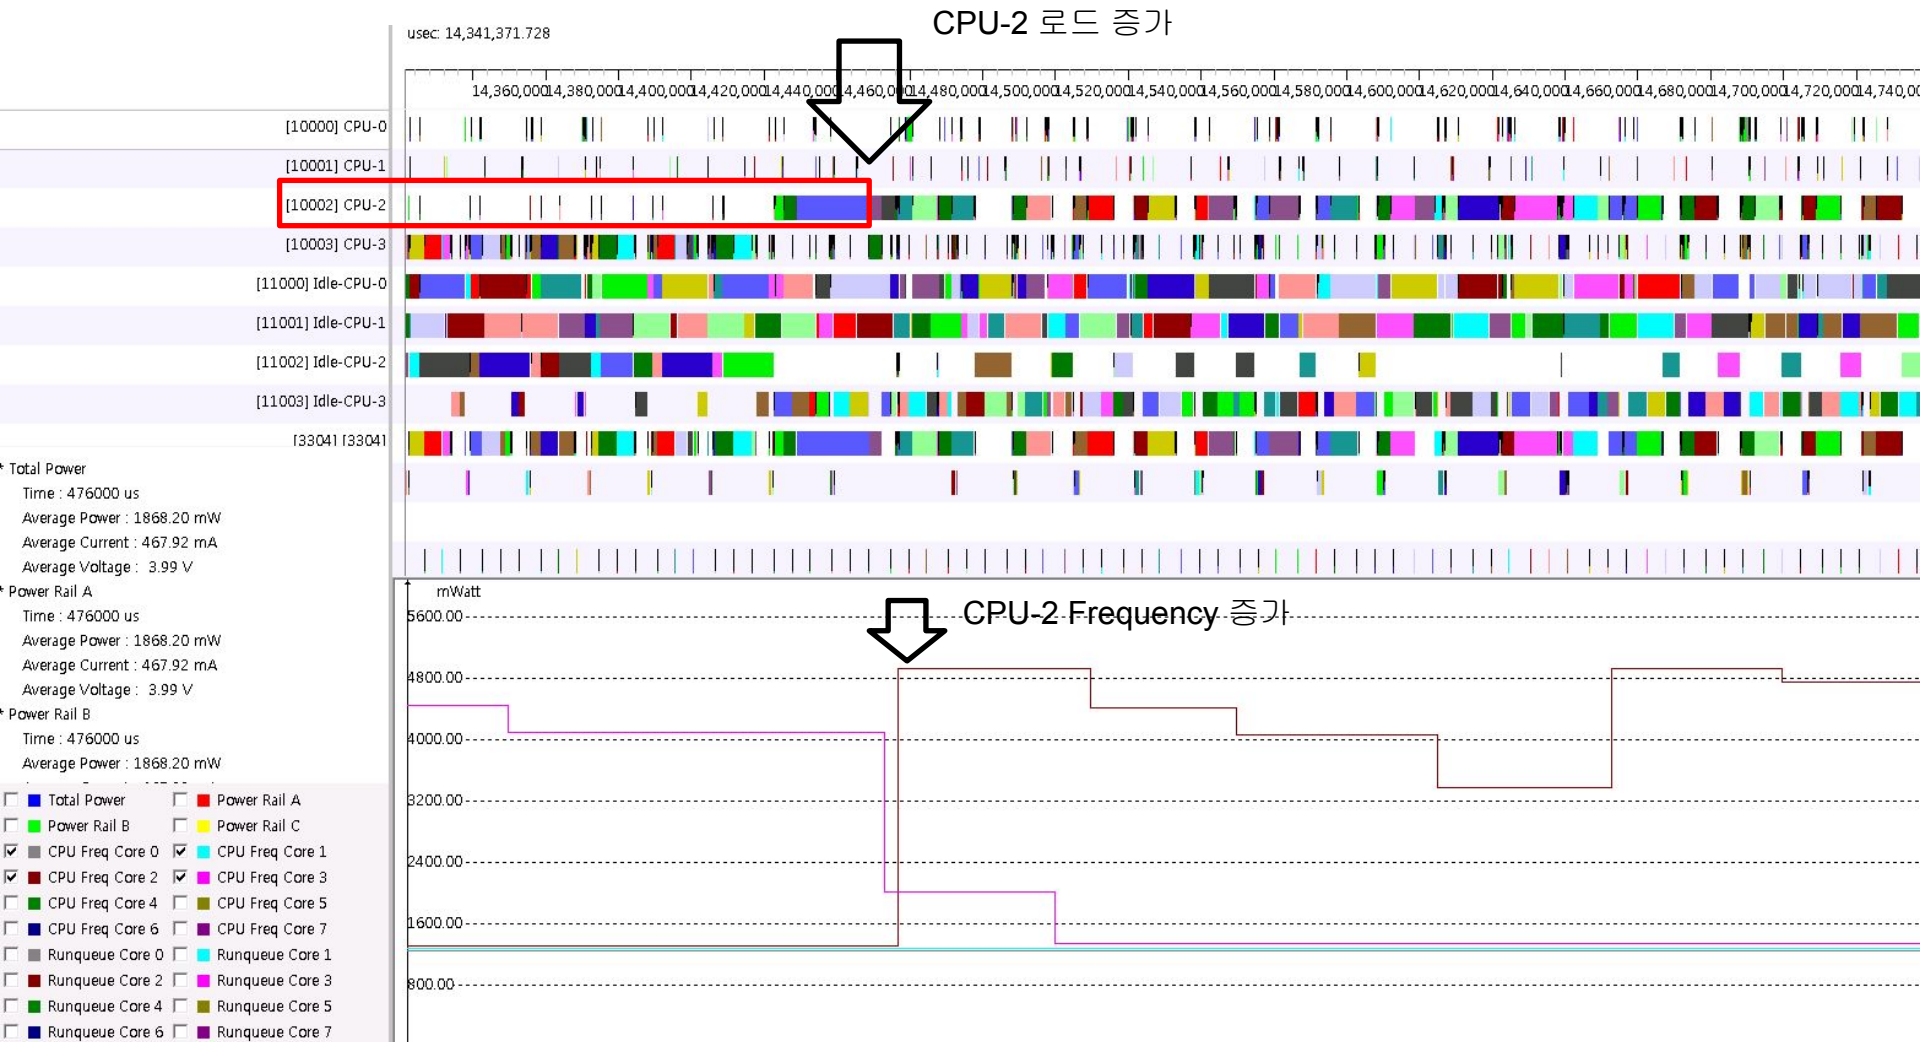


# Ex) Interactive Governor


- Interactive



# Interface



# Boost

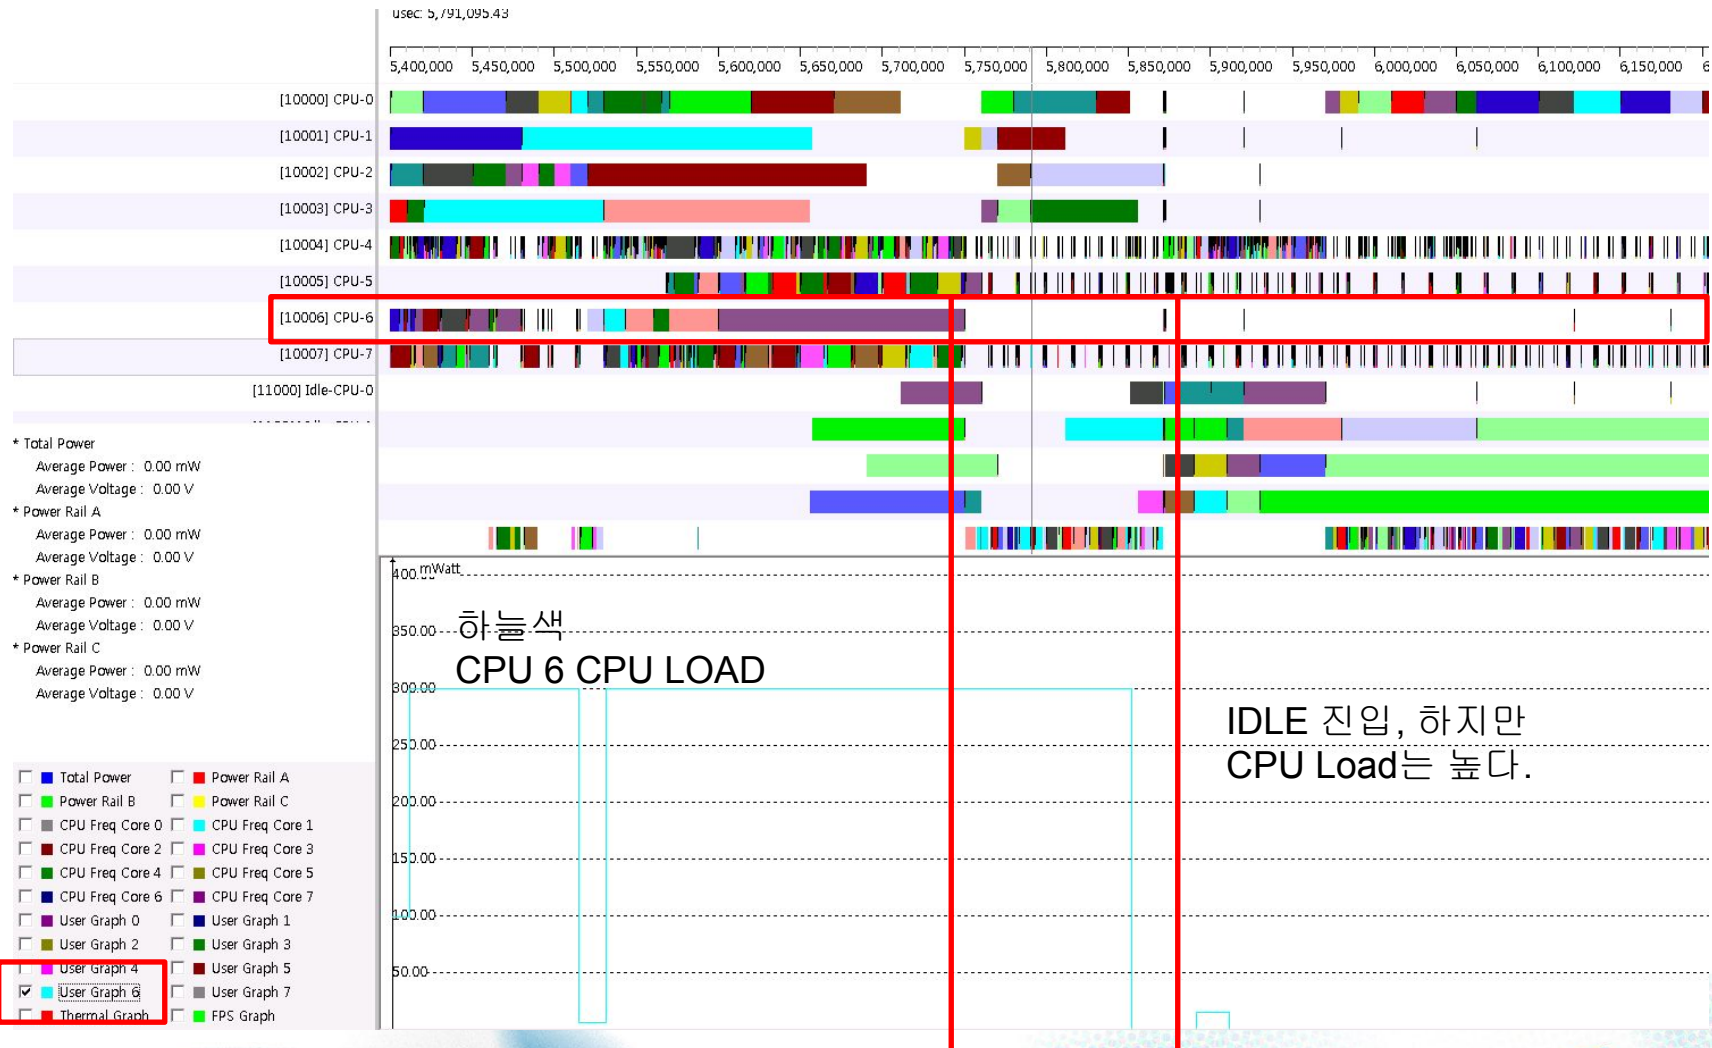




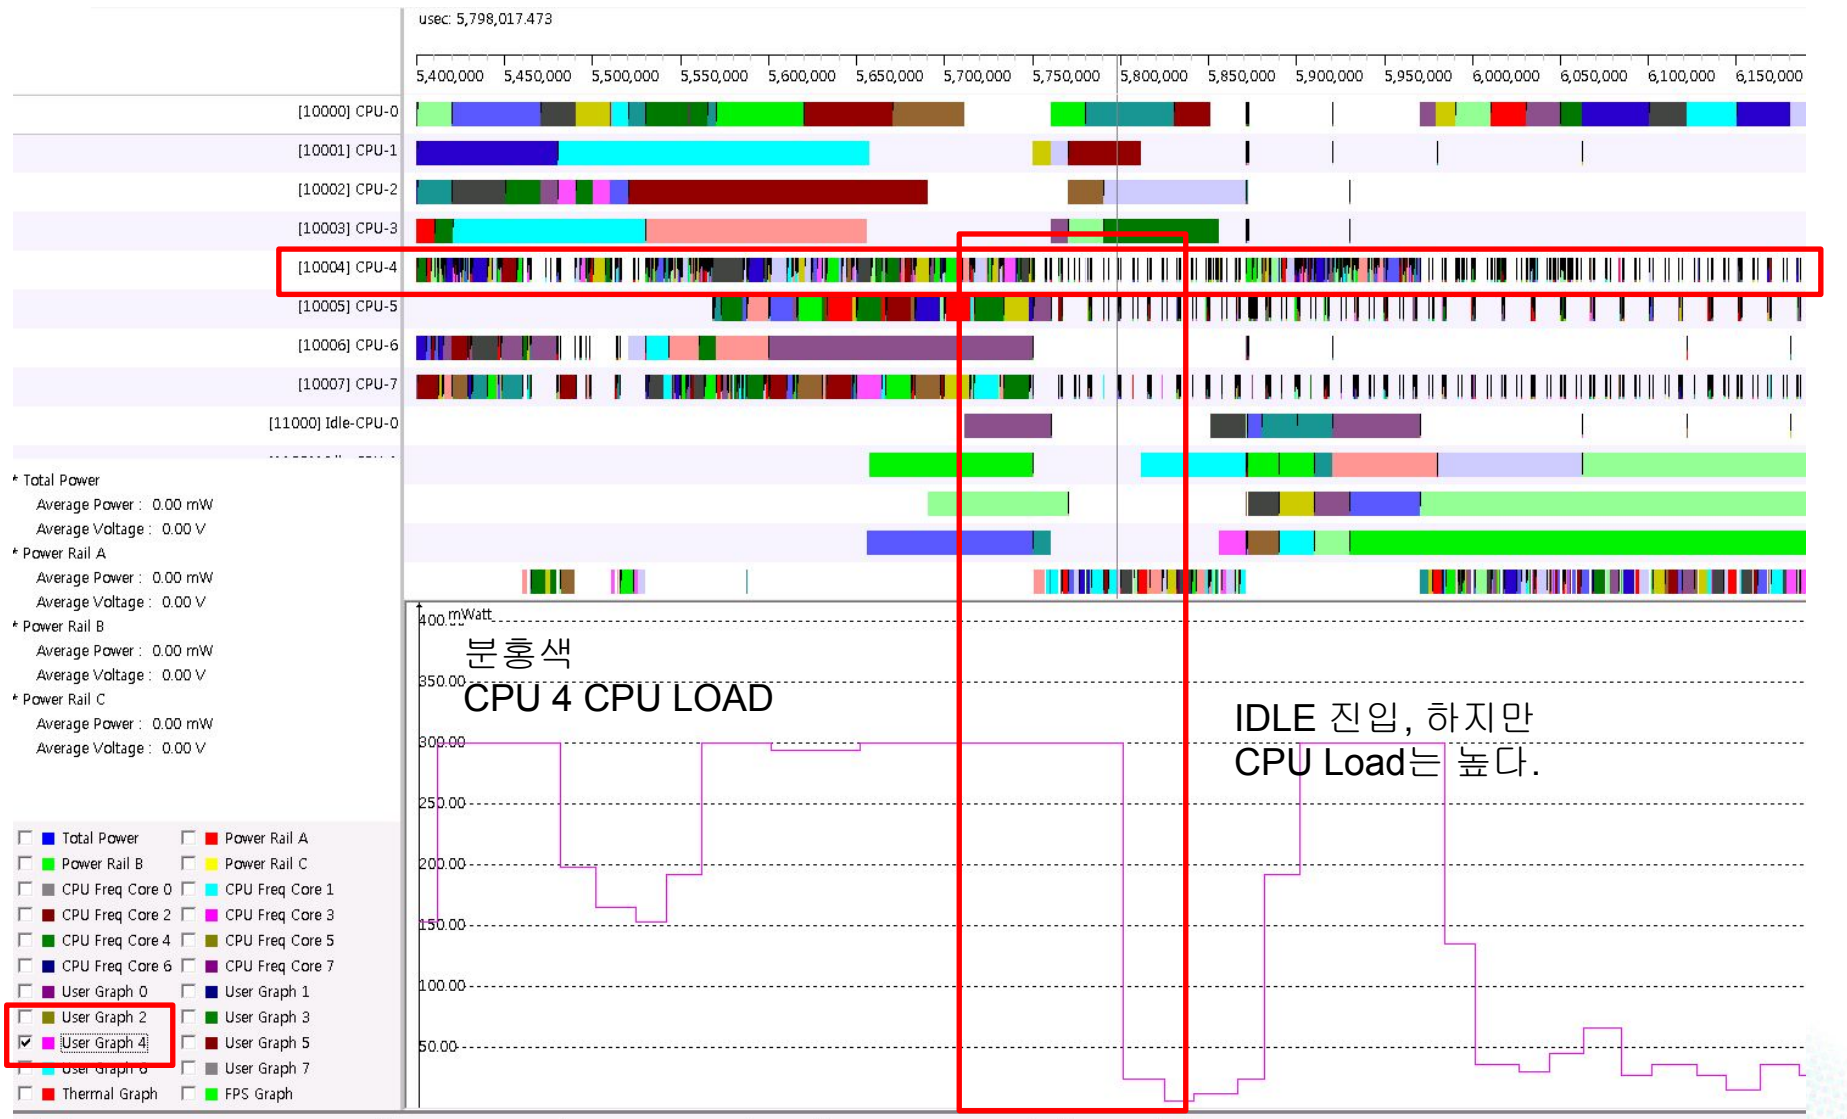
# Floor validate\_time



# Deferrable timer





# Deferrable timer



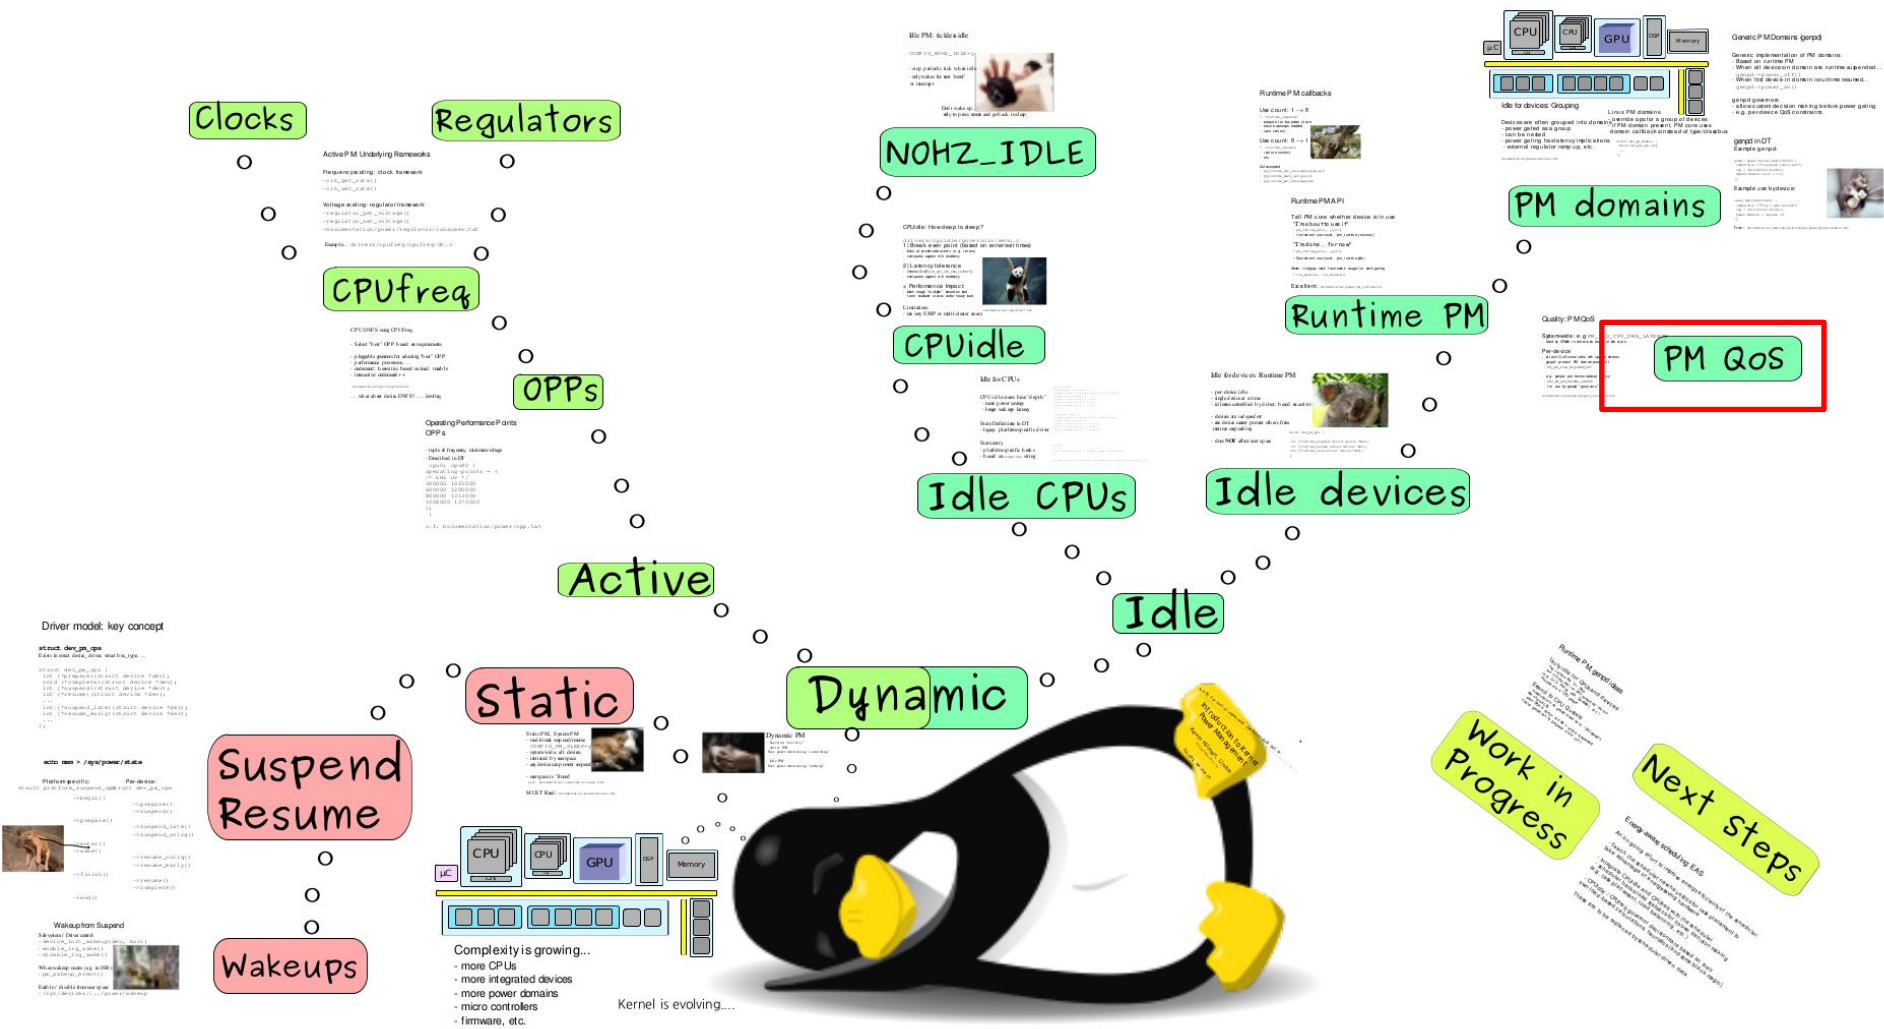
# DVFS in Linux (cpufreq)

- **Problem**
  - Sampling based governors are slow to respond and hard to tune.
- **Sampling too fast**
  - Freq changes for small utilization spikes.
- **Sampling too slow**
  - Average too low





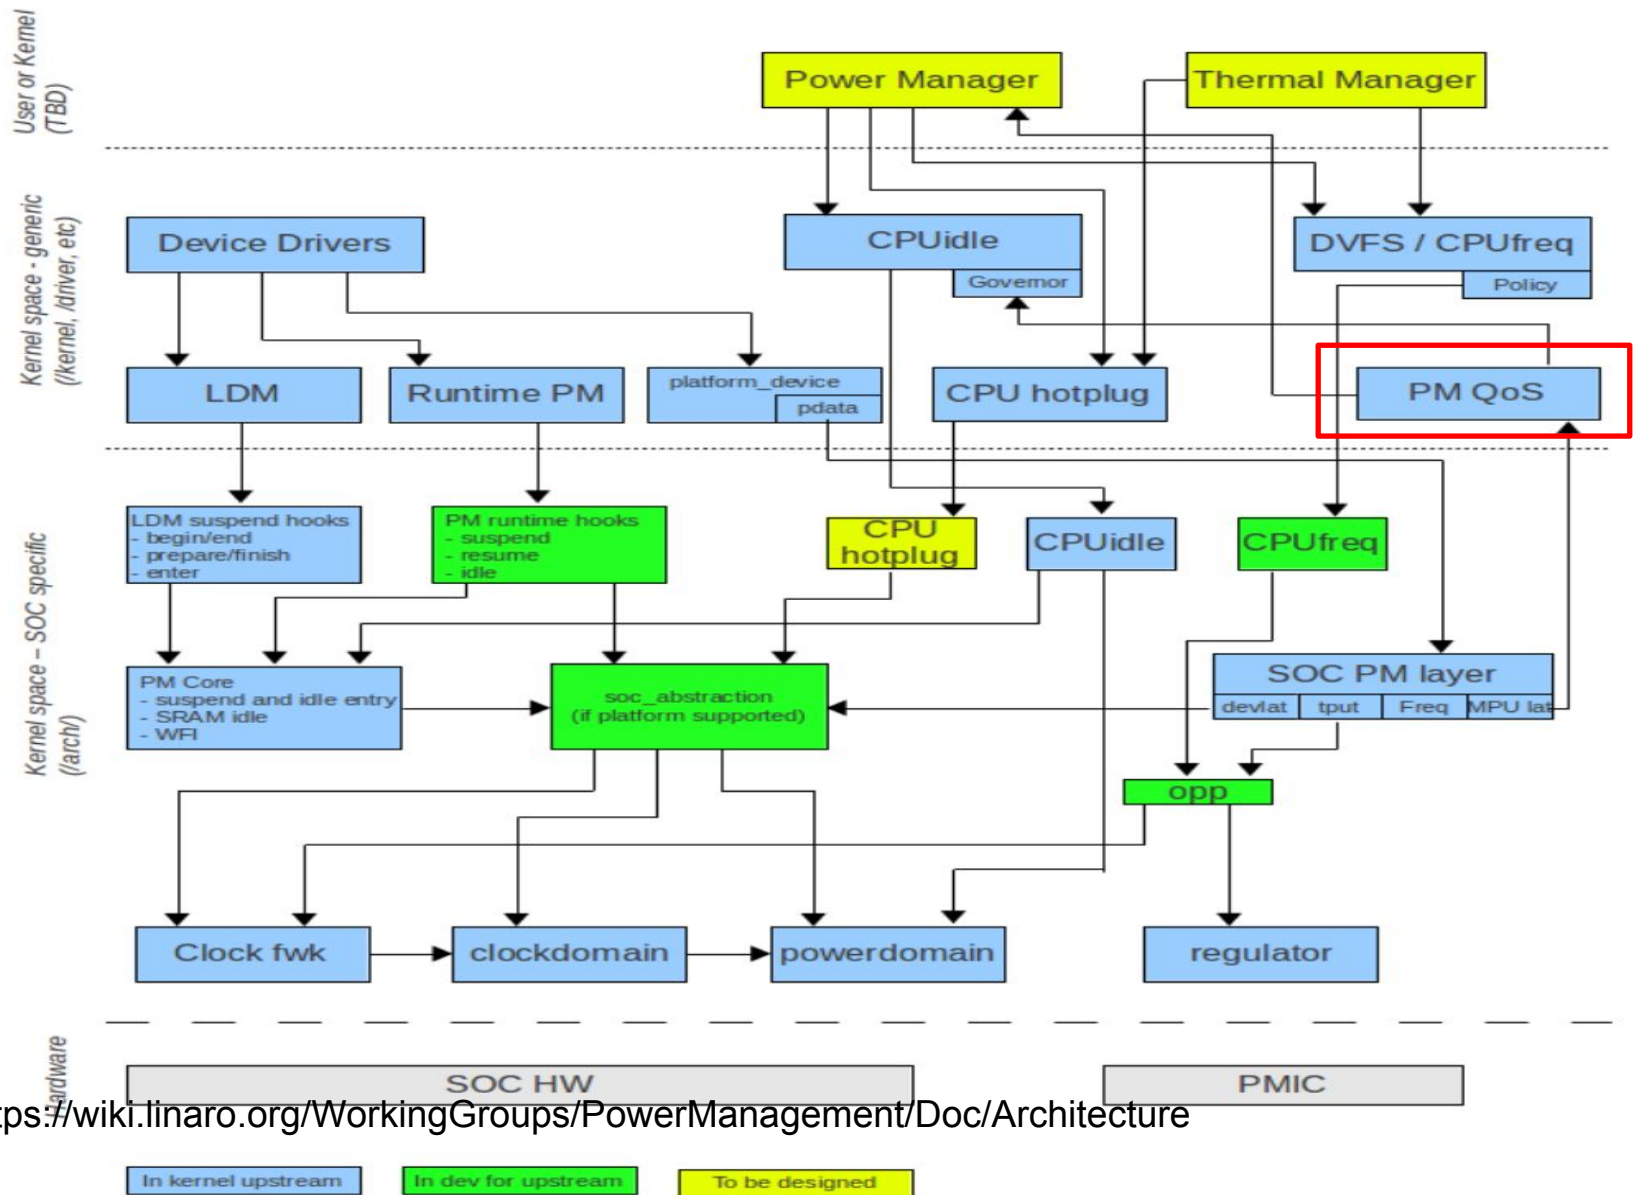
# Tree view - Dynamic and Static PM



Reference from "Introduction to Kernel Power Management" Kevin Hilman, Linaro



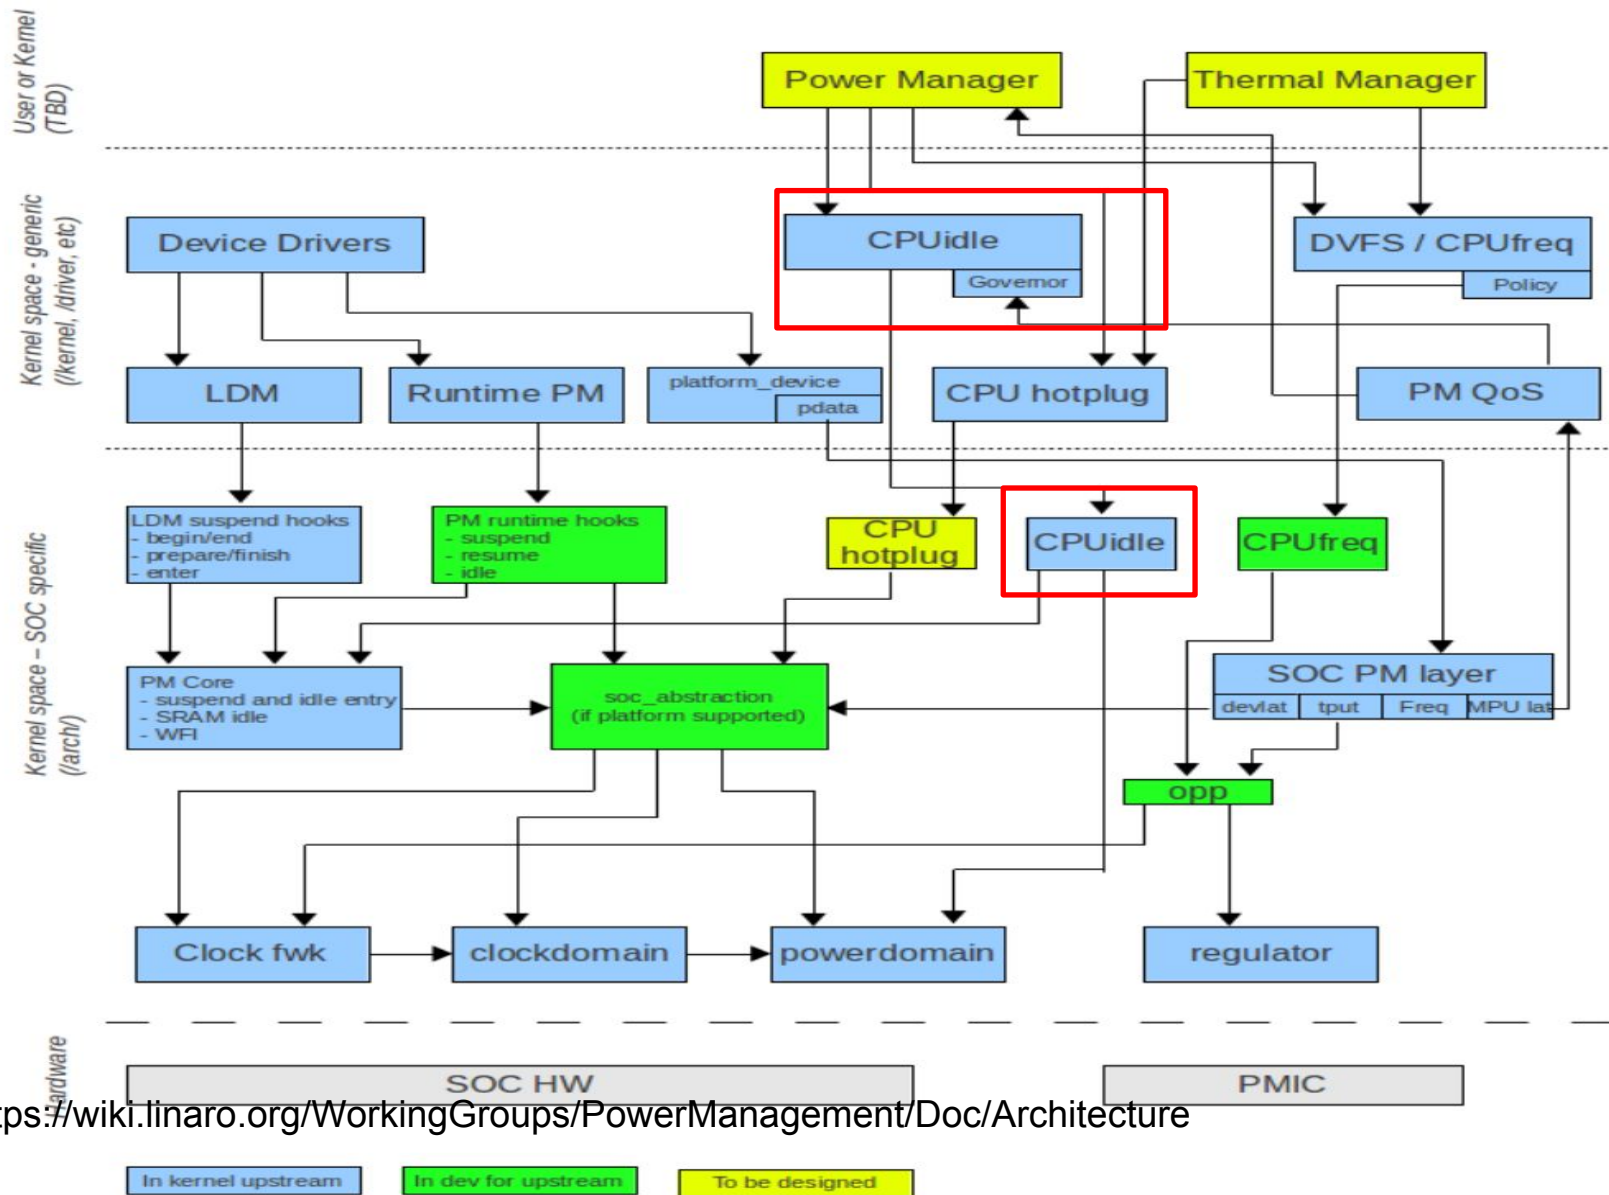
# Linux Power Framework



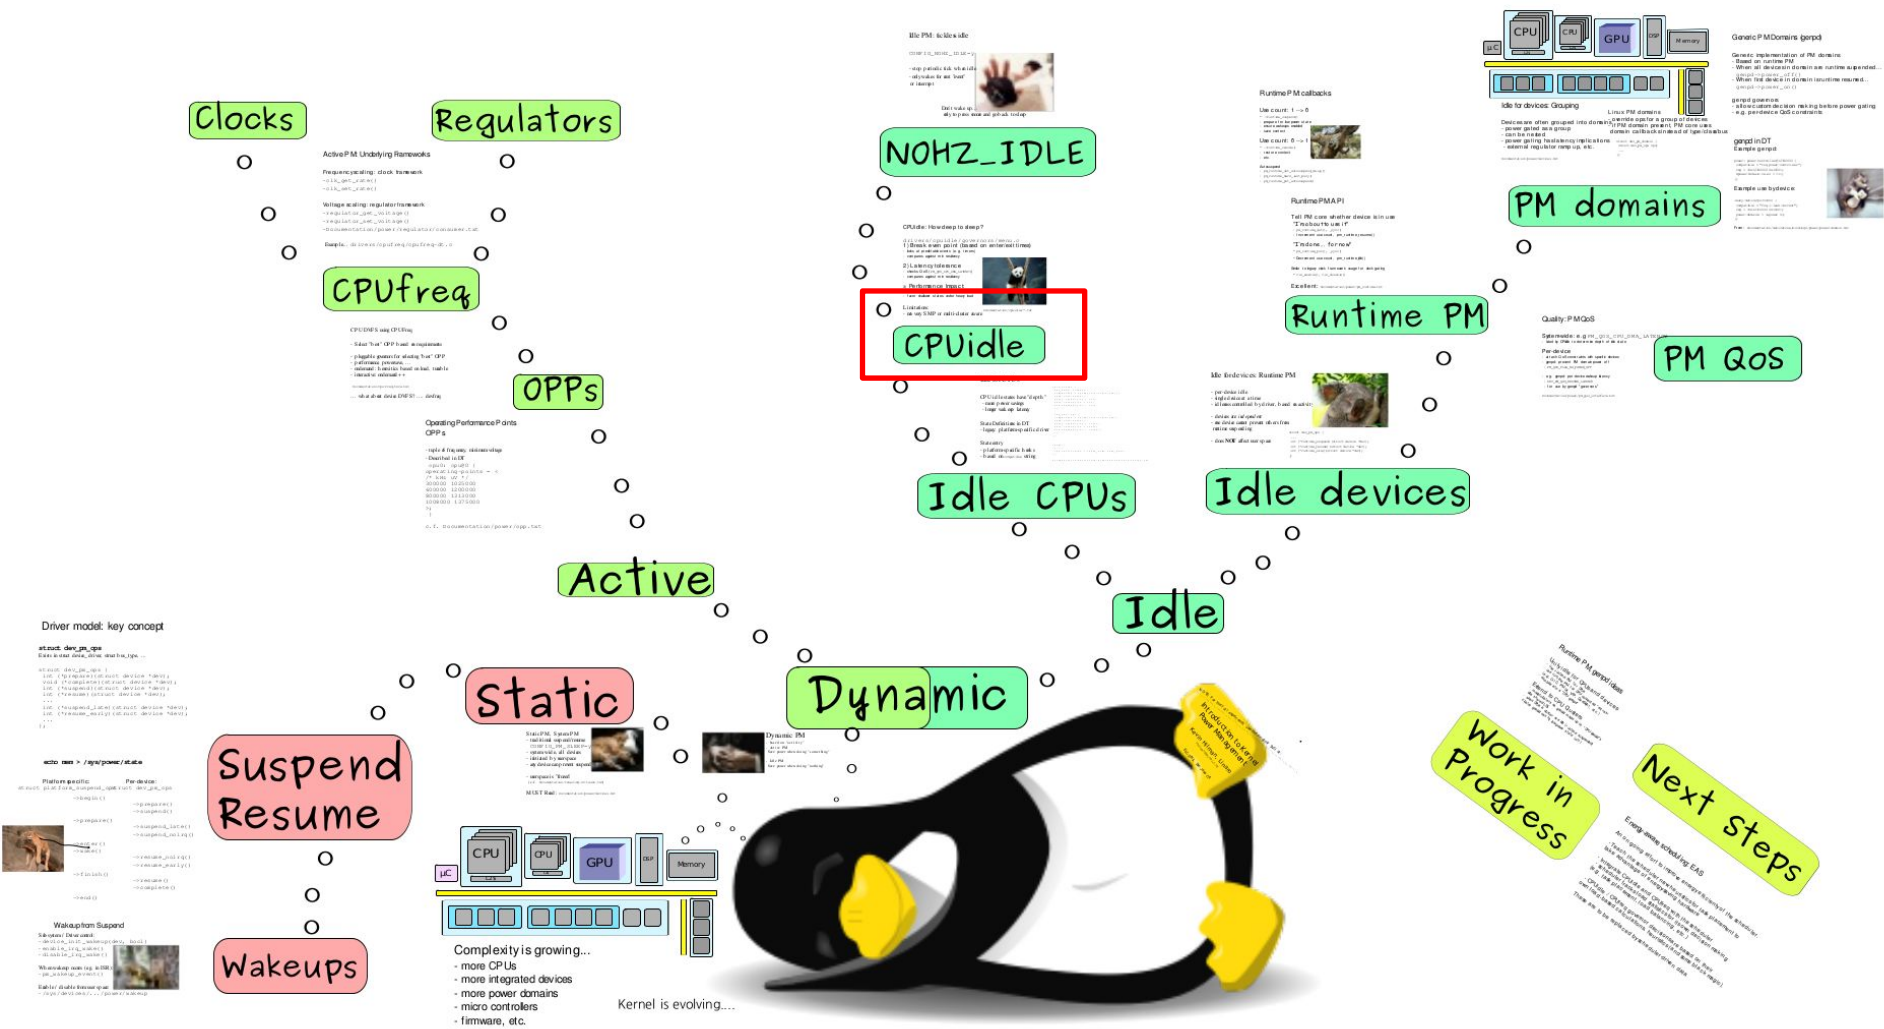
# PM\_QoS

- **Goal : to reduce energy**
- **A coordination mechanism**
  - User's performance needs
- **Examples**
  - PM QoS with cpuidle
  - Per-device PM QoS with runtime pm.

# Linux Power Framework



















# Tree view - Dynamic and Static PM



Reference from "Introduction to Kernel Power Management" Kevin Hilman, Linaro

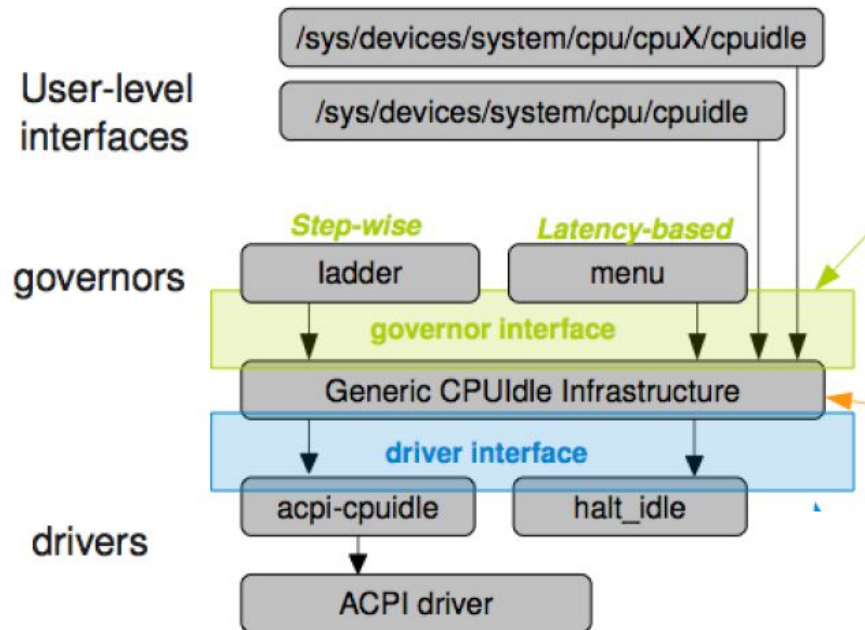
# Processor Power Management

|                  | Active state  | Sleep states   |   |  |  |
|------------------|---|--|---|--|--|
|                  | <u>C0</u>   | <u>C1/C1E</u>  | <u>C3</u>   | <u>C6</u>  | <u>C7</u>  |
|                  | Operating   | Halt   | Sleep   | Deep Sleep   |  |
| Core clock       |    | off  | off   | off  | off  |
| PLL              |    |     | off   | off  | off  |
| Core caches      |    |     | flushed   | flushed  | flushed  |
| Shared cache     |    |     |    |   | flushed  |
| Wakeup time*     | active  |    |   |  |  |
| Core Idle power* |  |  |  | $\sim 0$   | $< C6$   |

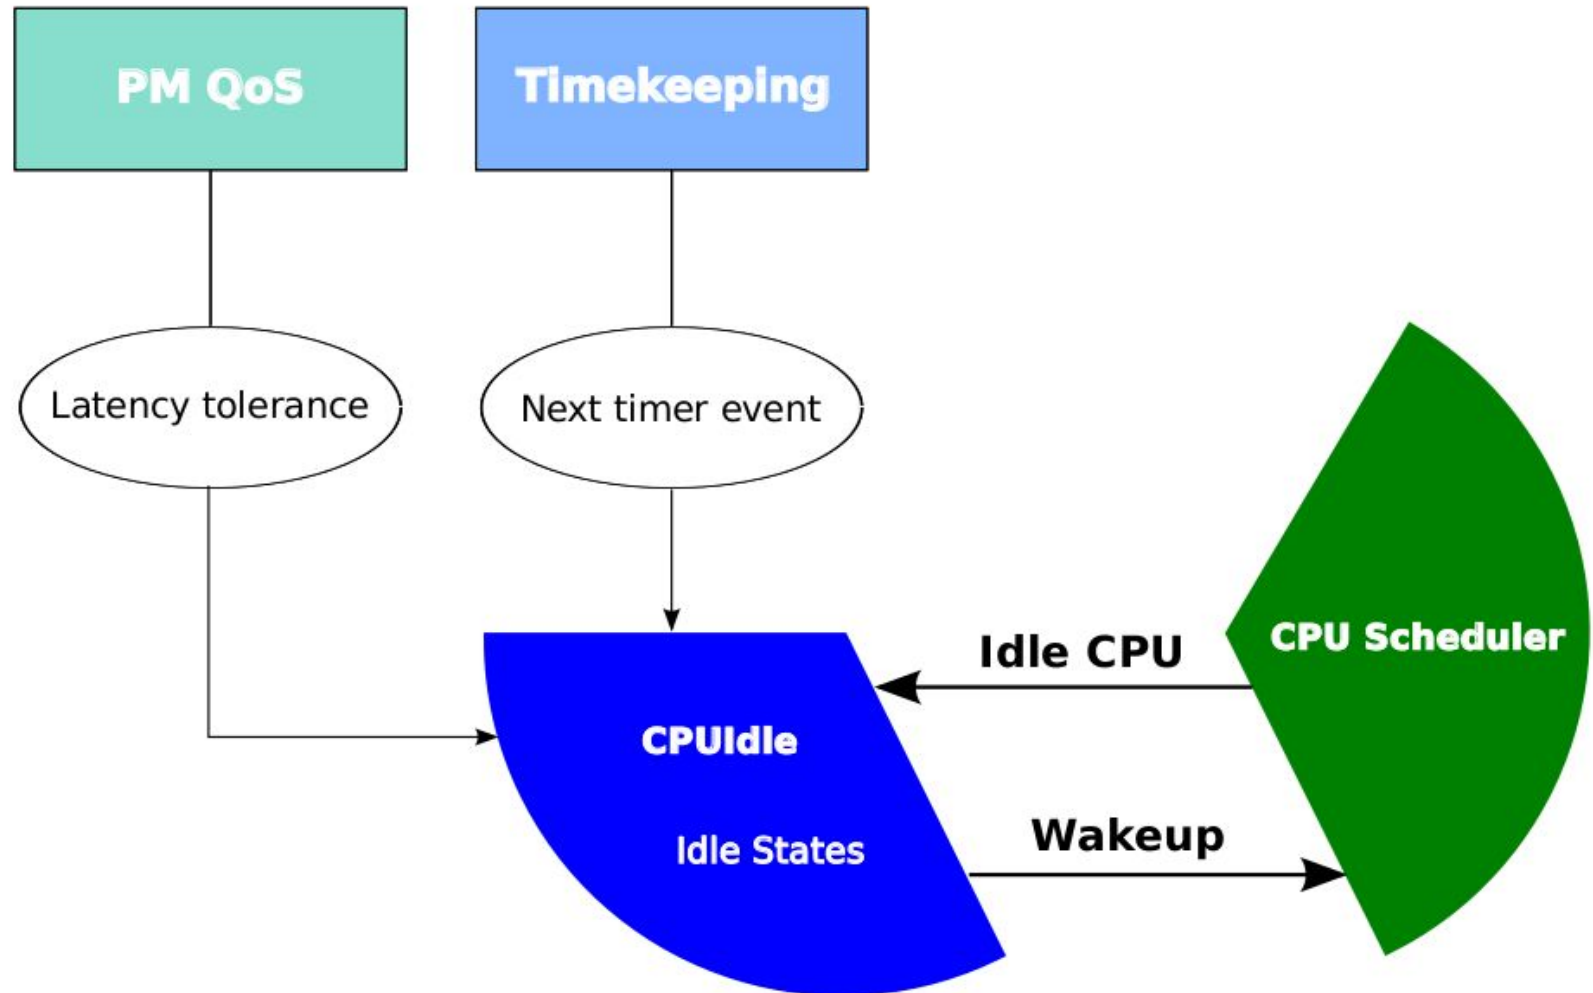
\* Rough approximation



# CPU idle - C-state management



# CPU idle



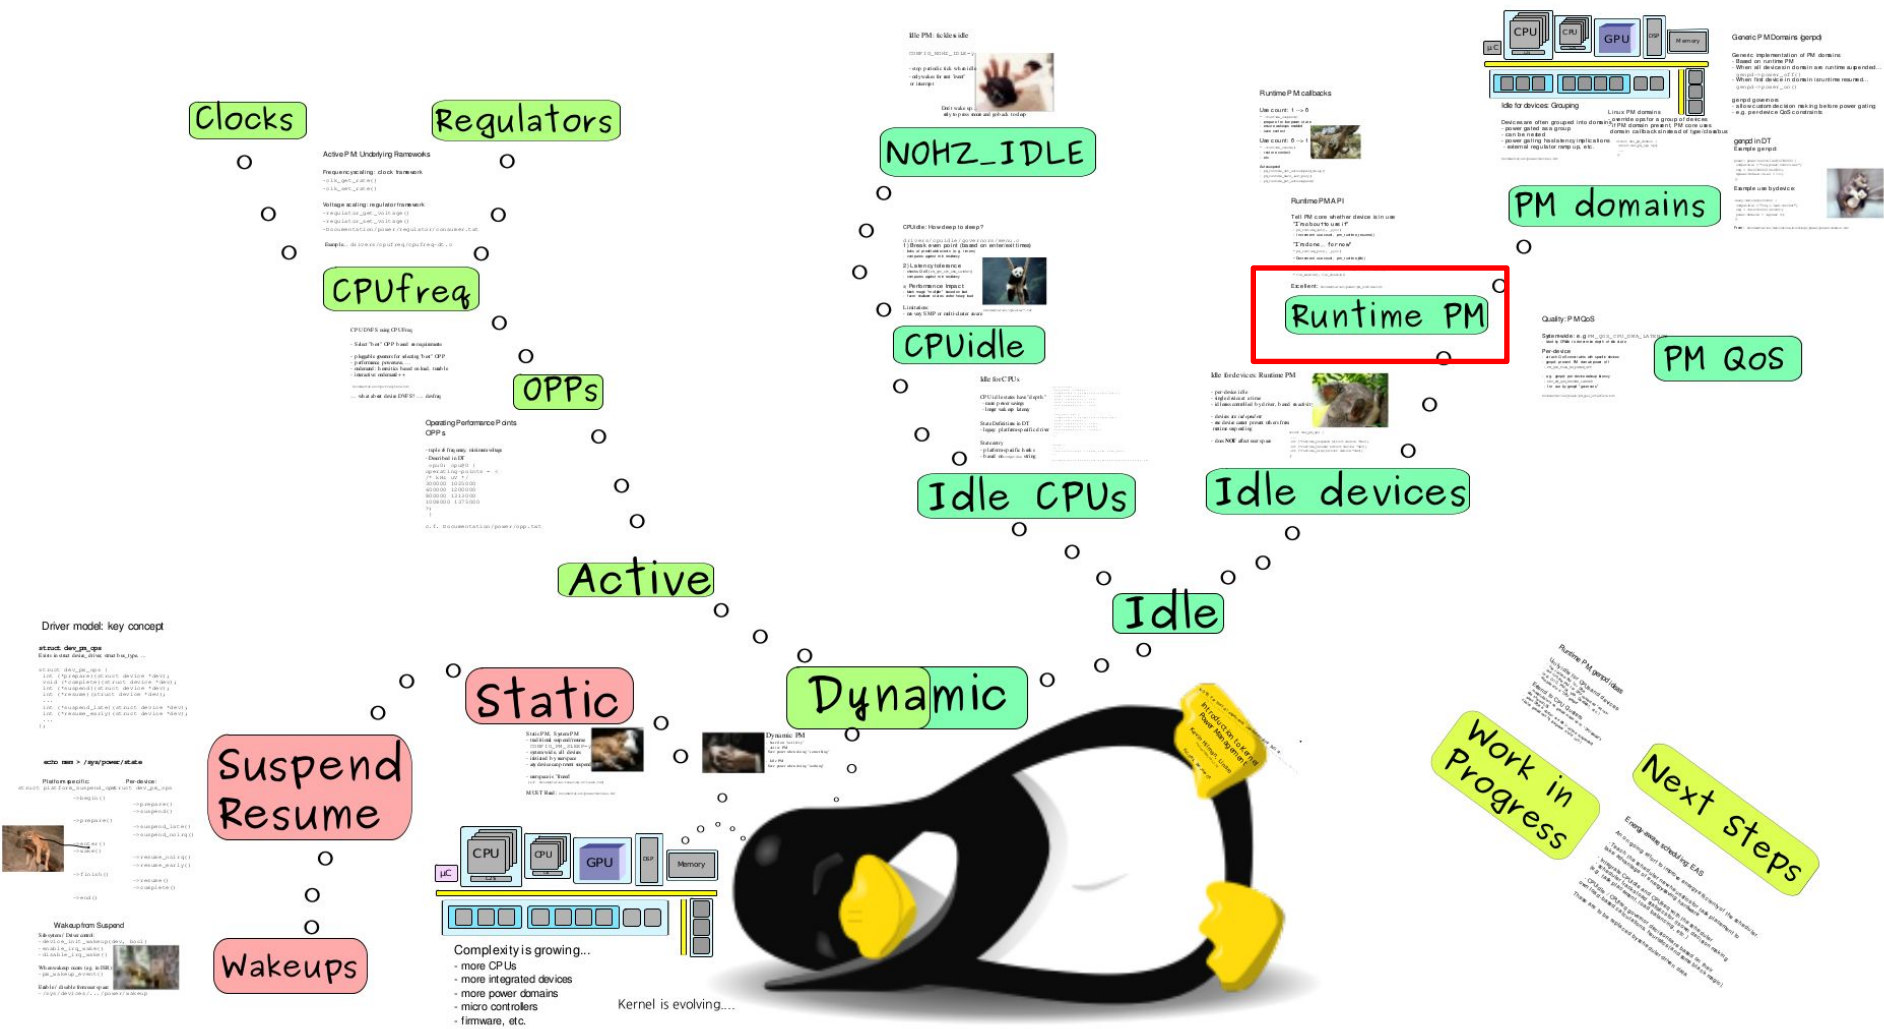
# Menu governor: How deep to sleep?

- **Energy break**
  - next timer event
  - historic behavior
- **Performance impact**
  - nr\_iowaiters, cpu\_load
- **Latency tolerance**
  - checks QoS (PM\_QOS\_CPU\_DMA\_LATENCY)



# **Per-device Power Management**

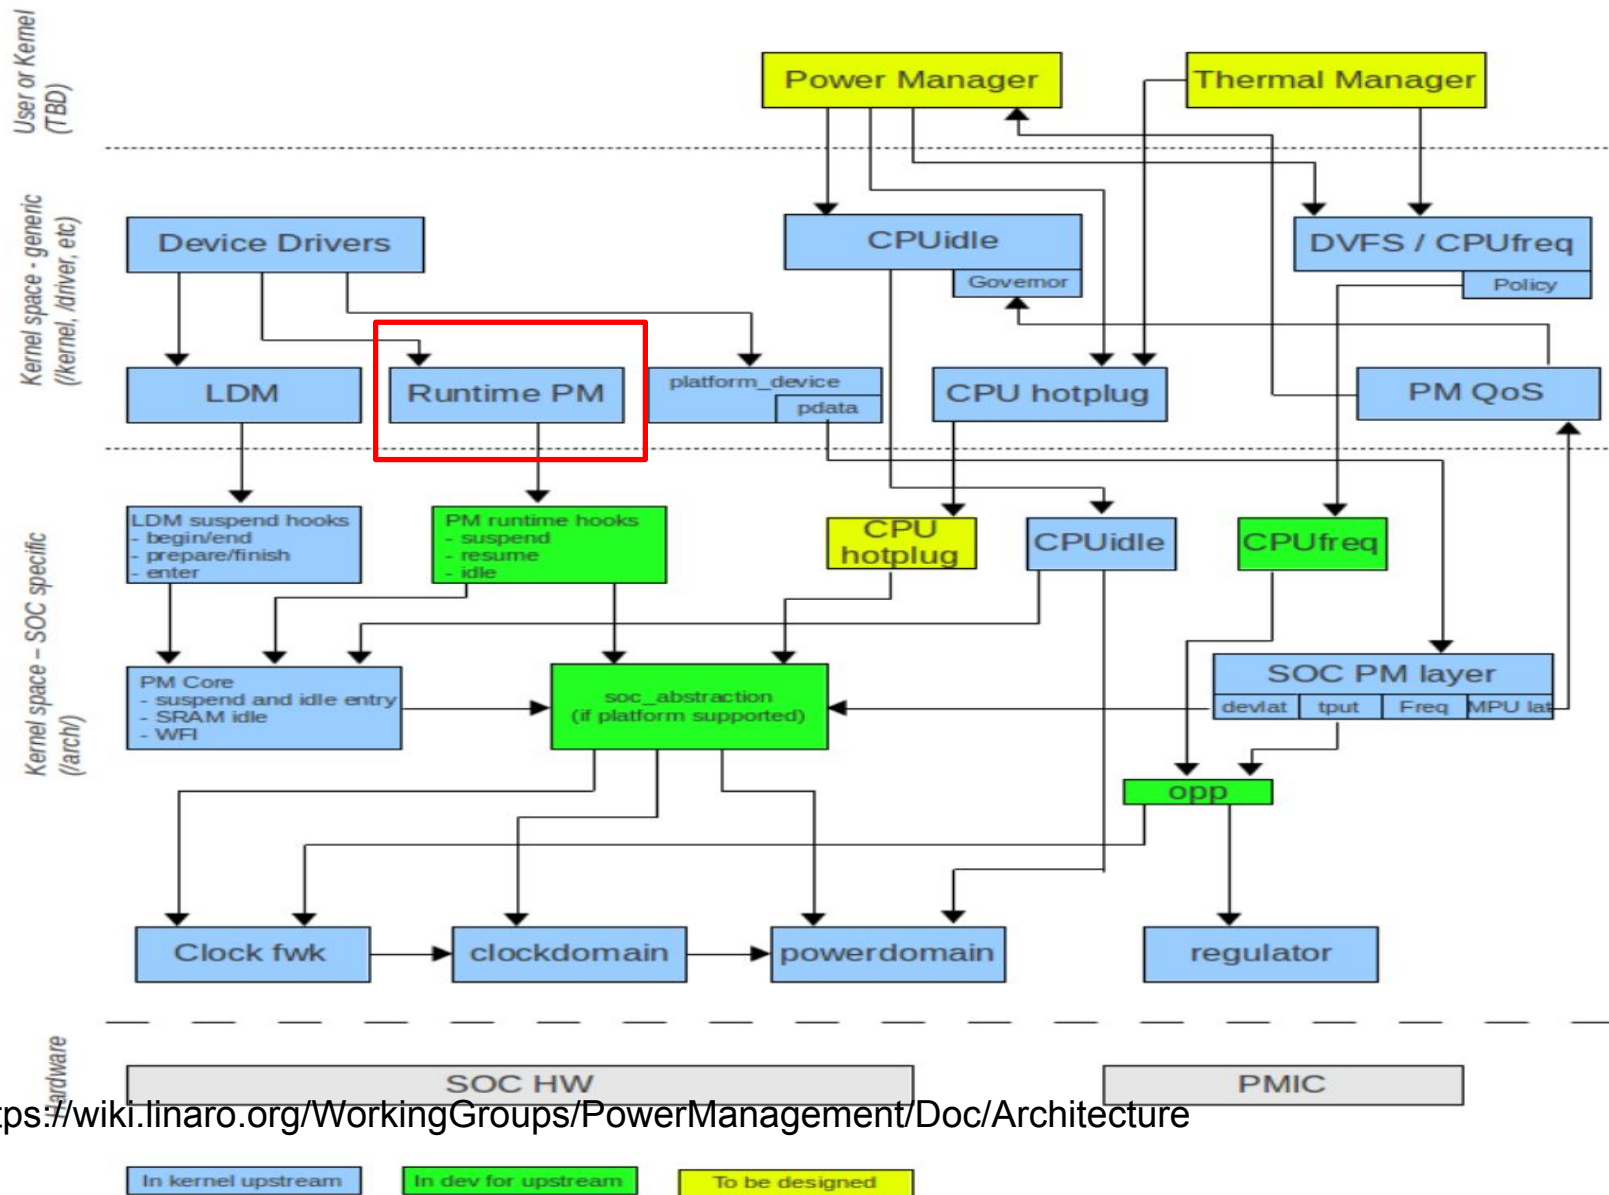
# Tree view - Dynamic and Static PM



Reference from "Introduction to Kernel Power Management" Kevin Hilman, Linaro



# Linux Power Framework



# Runtime Power Management

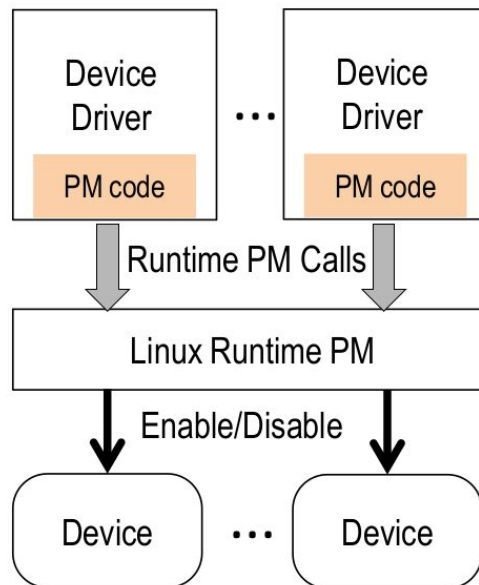
- Why's Runtime PM
- What's Runtime PM
- How's Runtime PM

# Problem

- **System sleep is not enough to decrease runtime energy consumption**
- **Devices may depend on another device**

# Why's Runtime PM

- **System sleep is not enough to decrease runtime energy consumption**
- **Devices may depend on another device**
- **Solution**
  - Runtime PM



# Runtime PM callback and Helper Functions

- **ksrc/include/linux/pm.h :**

```
struct dev_pm_ops {  
    ...  
    int (*runtime_suspend)(struct device *dev);  
    int (*runtime_resume)(struct device *dev);  
    int (*runtime_idle)(struct device *dev);  
};
```

- **Ksrc/include/linux/pm\_runtime.h**

- pm\_runtime\_xxx();



# Linux Support for Device Runtime PM

- Track the number of concurrent users of device
- use a per-device reference counter
  - `_get()`, `_put()`

# Runtime PM API: `_get()`, `_put()`

- **Tell PM core whether device is in use**
- **I need the device**
  - `pm_runtime_get()`, `_sync()`, `_noresume()`
  - Increment use count, `pm_runtime_resume()`
- **I'm done**
  - `pm_runtime_put()`, `_sync`, `_noidle()`
  - Decrement use count, `pm_runtime_idle()`

# Autosuspend

- **Some device should not allow devices to be suspend**
  - until they have been inactive for some minimum period.
- **A common heuristic**
  - similarly a way of PM qos
- **Init**
  - `pm_runtime_set_suspended(dev);`
  - `pm_runtime_use_autosuspend(dev);`
  - `pm_runtime_set_autosuspend_delay(dev, DELAY);`
- **I need the device**
  - `pm_runtime_get_sync(dev);`
- **I'm done**
  - `pm_runtime_put_sync_autosuspend(dev);`

# However, Real world - Runtime PM

- Exynos SPI driver
- Runtime PM often supported after long delay.

```
1 int s3c64xx_spi_probe(platform_device
    *pdev)
2 {
3     /* allocate controller resources...*/
4     pm_runtime_get_sync(dev);
5     /*initialize the controller...*/
6 }
7
8 int s3c64xx_spi_remove(platform_device
    *pdev)
9 {
10    /* deinitialize controller...*/
11    pm_runtime_put(dev);
12    /* free controller resources...*/
13 }
```



```
1 void s3c64xx_spi_work(work_struct *work)
2 {
3     pm_runtime_get_sync(dev);
4     while (!list_empty(queue)) {
5         /* transmitting message... */
6     }
7     pm_runtime_put(dev);
8 }
9
10 int s3c64xx_spi_setup(spi_device *spi)
11 {
12     pm_runtime_get_sync(dev);
13     /* set up SPI, like tx rate... */
14     pm_runtime_put(dev);
15 }
```

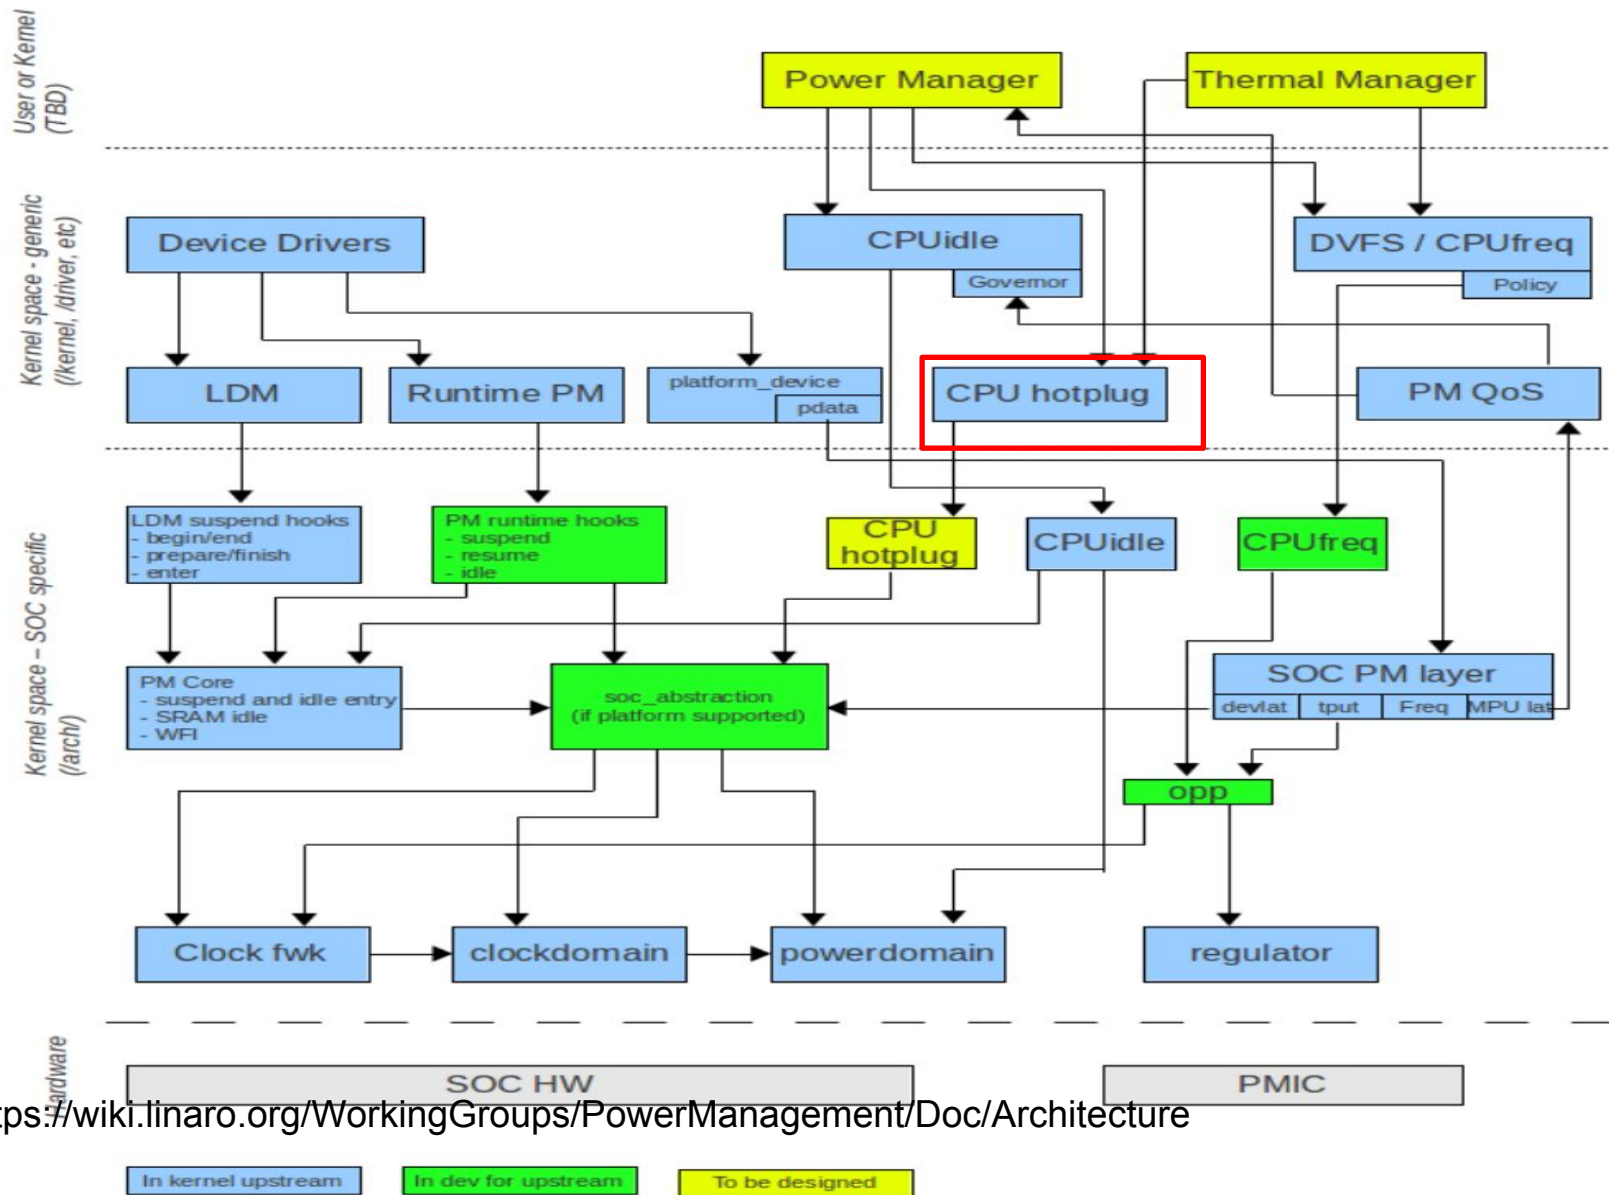
Hand-tuned PM in the Exynos SPI driver

# However, Real world - Runtime PM

- **Complex drivers make it hard to do runtime PM.**
  - It is similar with a Lock.
- **The impact of bad drivers.**
  - Because clock, power, voltage depends on the domains.
  - sharing the same clock source.
  - sharing the same power
  - sharing the same voltage

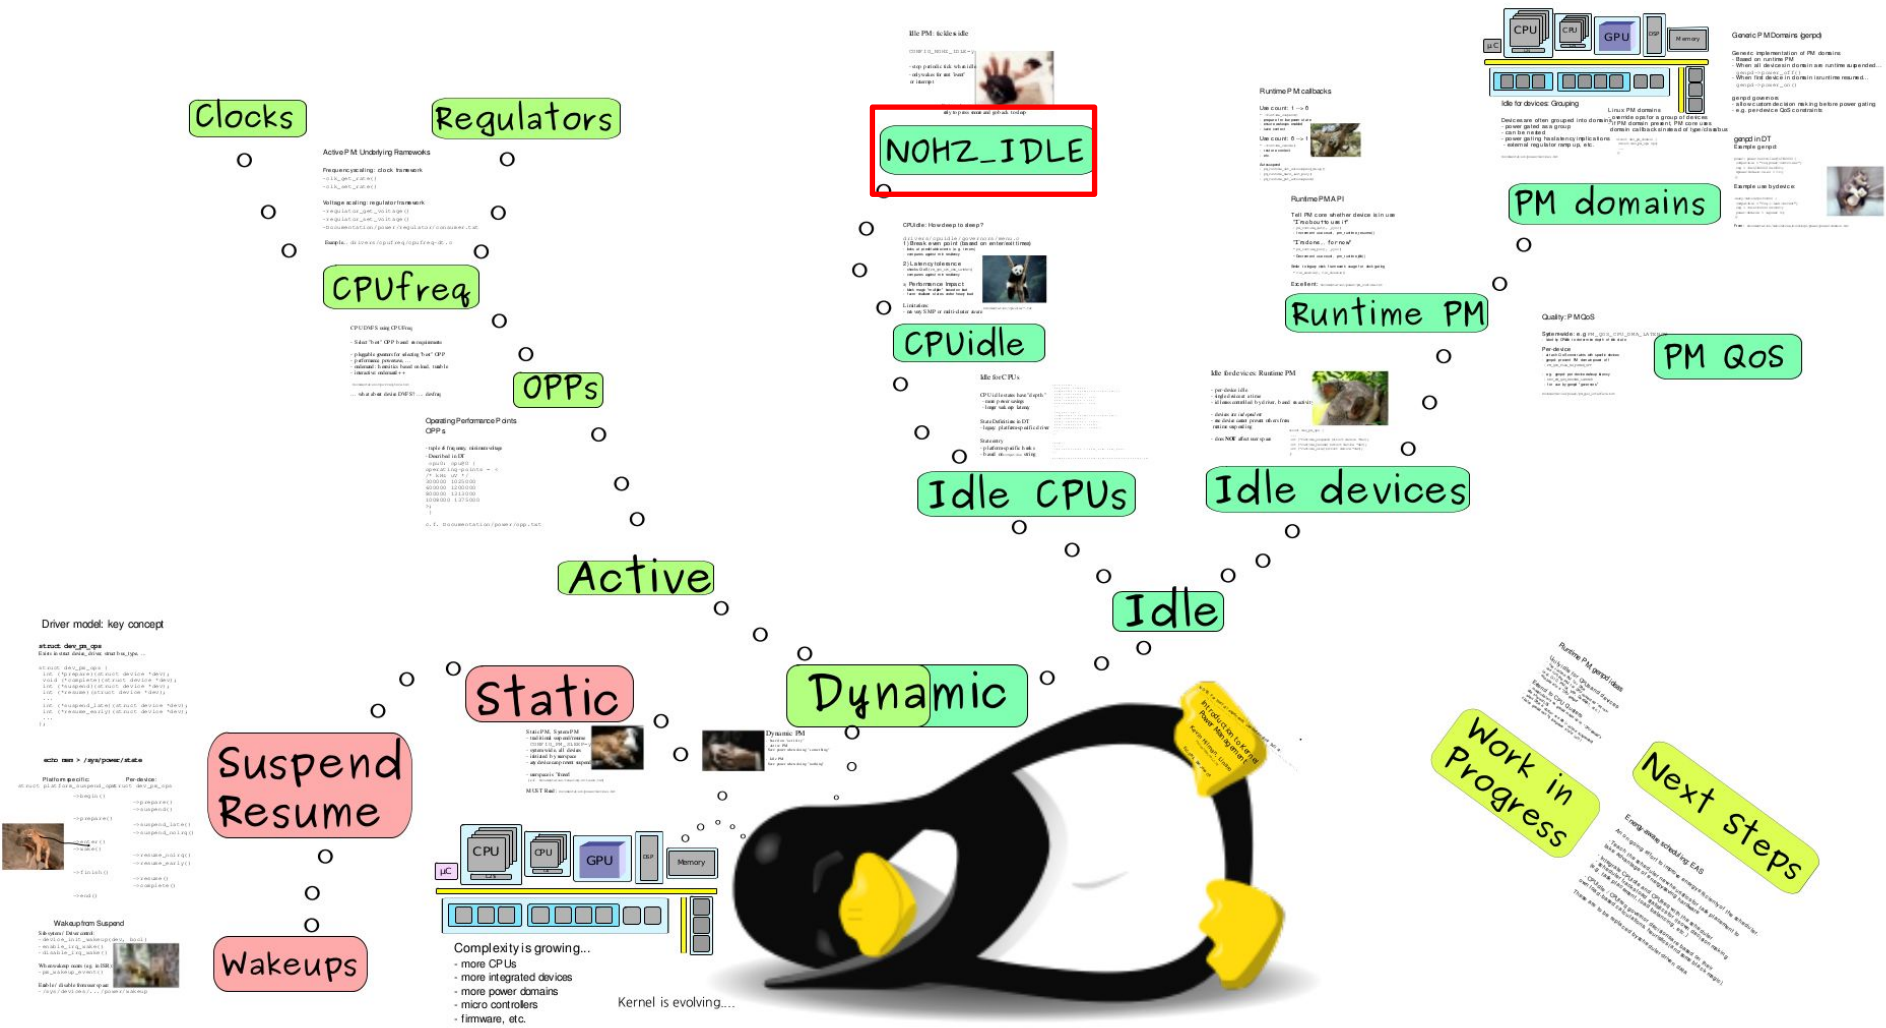


# Hotplug



<https://wiki.linaro.org/WorkingGroups/PowerManagement/Doc/Architecture>

# NOHZ\_IDLE



# Queue work on power efficient wq

- **Problem**

- Work-queues can be performance or power oriented.
- keep them running on a single cpu
  - it remains cache hot
- but consider big.LITTLE platform

# Queue work on power efficient wq

- **Problem**

- Work-queues can be performance or power oriented.
- keep them running on a single cpu
  - it remains cache hot
- but consider big.LITTLE platform

- **Solution**

- Power oriented WQ
- **Give scheduler the liberty**
  - to choose target cpu for running work handler.

- **Power efficient WQ**

- workqueue is allocated with WQ\_UNBOUND flag
- Idle time for many cpus has increased considerably on **big.LITTLE platform**

# Thermal - iphone, galaxy






# Thermal - G6



# Thermal

Problem?





# Thermal

## Spread out



# Today

**process/thread -> performance -> multi-core  
performance scalability-> lock-> bottleneck-> cache  
cohere system problem-> per-core partition  
approach**

## Next Step.

# Energy-aware scheduling: EAS

1. **CFS scheduler(User level, Tools, Kernel level)**
2. **Load Balancer(Group Scheduling, Bandwidth Control, PELT)**
3. **EAS features**





# Reference

- <https://pdos.csail.mit.edu/6.828/2016/schedule.html>
- <http://web.mit.edu/6.033>
- <http://www.rdrop.com/~paulmck/>
- "Is Parallel Programming Hard, And If So, What Can You Do About It?"
- Davidlohr Bueso. 2014. Scalability techniques for practical synchronization primitives. *Commun. ACM* 58  
<http://queue.acm.org/detail.cfm?id=2698990>
- "CPUFreq and The Scheduler Revolution in CPU Power Management", Rafael J. Wysocki
- <https://sites.google.com/site/embedwiki/oses/linux/pm/pm-qos>
- <https://intl.aliyun.com/forum/read-916>
- User-level threads : co-routines  
<http://www.gamedevforever.com/291>  
[https://www.youtube.com/watch?v=YYtzQ355\\_Co](https://www.youtube.com/watch?v=YYtzQ355_Co)
- Scheduler Activations
  - <https://cgi.cse.unsw.edu.au/~cs3231/12s1/lectures/SchedulerActivations.pdf>
- [https://en.wikipedia.org/wiki/FIFO\\_\(computing\\_and\\_electronics\)](https://en.wikipedia.org/wiki/FIFO_(computing_and_electronics))
- <http://jake.dothome.co.kr/>
- <http://www.linuxjournal.com/magazine/completely-fair-scheduler?page=0.0>
- [https://www2.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/6\\_CPU\\_Scheduling.html](https://www2.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/6_CPU_Scheduling.html)
- "Energy Aware Scheduling", Byungchul Park, LG Electronic
- "Update on big.LITTLE scheduling experiments", ARM
- "EAS Update" 2015 september ARM
- "EAS Overview and Integration Guide", ARM TR
- "Drowsy Power Management", Matthew Lentz, SOSP 2015
- <https://www.slideshare.net/nanik/learning-aosp-android-hardware-abstraction-layer-hal>
- <https://www.youtube.com/watch?v=oTGQXqD3CNI>
- <https://www.youtube.com/watch?v=P80NcKUKpuo>
- <https://lwn.net/Articles/398470/>
- "SCHED\_DEADLINE: It's Alive!", ARM, 2017