

Simplified 리눅스 스케줄러(FIFO, RR)

국민대학교 임베디드 연구실
경 주 현


Outline

- **Simplified Linux scheduler(FIFO, RR)**
- **Application : 스케줄러 System call 사용**
- **트레이스 실습**
- **Simplified Linux scheduler(FIFO, RR) 소스**

The objectives of this pt is to learn

- How to add a pluggable CPU scheduler in Linux kernel.
- How to apply different scheduling policies to individual processes.
- Learn the internal working of a CPU scheduler.

Kernel-level thread




The key decisions made in the scheduler

*“how to determine a thread’s **timeslice**? and how to pick the **next thread** to run”*

FIFO

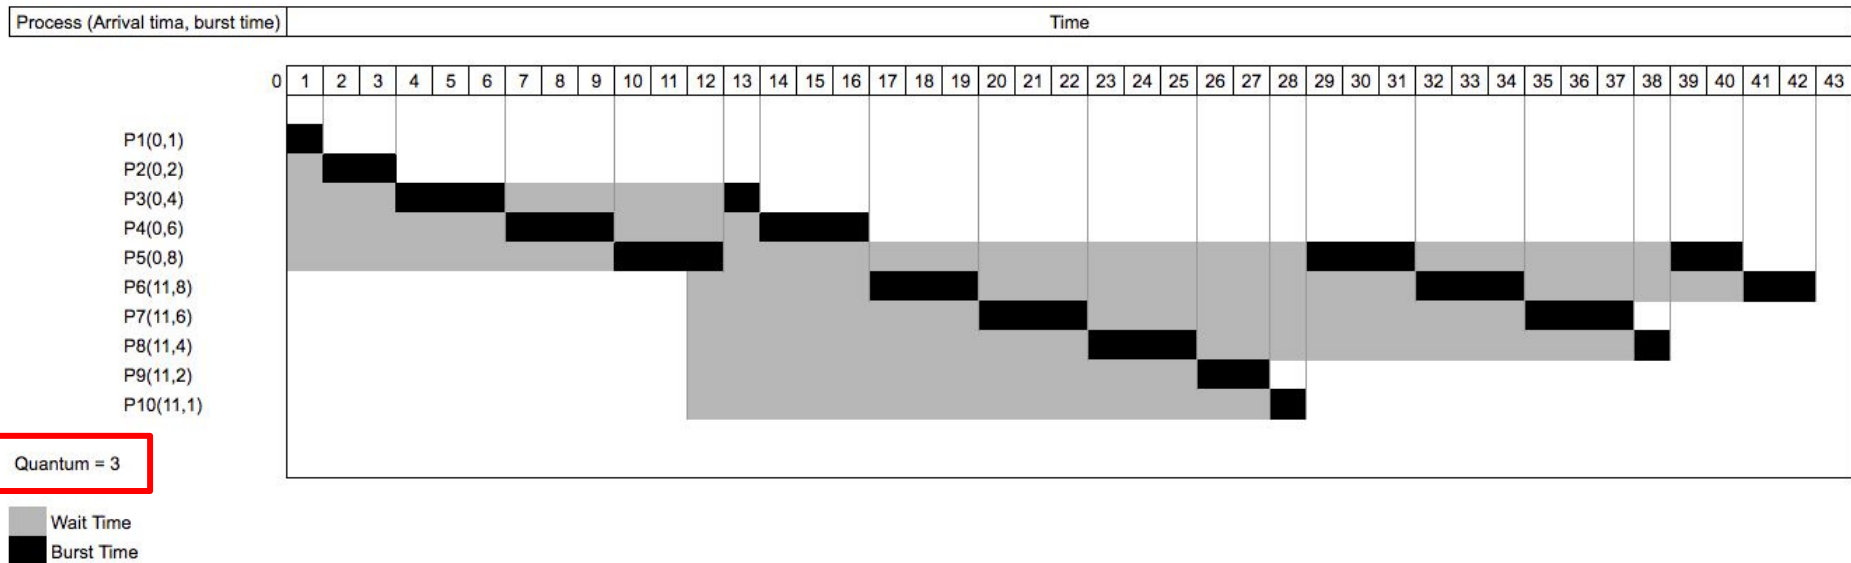
- first in, first out
- Problem : starvation



[https://en.wikipedia.org/wiki/FIFO_\(computing_and_electronics\)](https://en.wikipedia.org/wiki/FIFO_(computing_and_electronics))

Round Robin

- Simple, easy to implement, and starvation-free.



커널 소스 위치

- **/home/lkc/Desktop/linux-linaro-stable**
- **Rootfs 위치**
 - /home/lkc/Desktop/linux-linaro-stable/rootfs/initrd

커널 컴파일

- `cd ~/Desktop/linux-linaro-stable`
- `make -j2 bzImage`

```
lkc@lkc-VM:~/Desktop/linux-linaro-stable$ make -j2 bzImage
CHK      include/config/kernel.release
CHK      include/generated/uapi/linux/version.h
CHK      include/generated/utsrelease.h
CHK      include/generated/timeconst.h
CHK      include/generated/bounds.h
CHK      include/generated/asm-offsets.h
CALL     scripts/checksyscalls.sh
CHK      include/generated/compile.h
Kernel: arch/x86/boot/bzImage is ready (#1)
```

커널 실행

- `cd ~/Desktop/linux-linaro-stable`
- `make qemu`

```
lkc@lkc-VM:~/Desktop/linux-linaro-stable$ pwd
/home/lkc/Desktop/linux-linaro-stable
lkc@lkc-VM:~/Desktop/linux-linaro-stable$ make qemu
./qemu/i386-softmmu/qemu-system-i386 -smp 2 -kernel arch/x86/boot/bzImage -initrd i
0.000000] Initializing cgroup subsys cpuset
0.000000] Initializing cgroup subsys cpu
0.000000] Initializing cgroup subsys cpuacct
0.000000] Linux version 4.4.71 (lkc@lkc-VM) (gcc version 4.8.4 (Ubuntu 4.8.4-2ub
0.000000] x86/fpu: Legacy x87 FPU detected.
0.000000] x86/fpu: Using 'lazy' FPU context switches.
0.000000] e820: BIOS-provided physical RAM map:
0.000000] BIOS-e820: [mem 0x0000000000000000-0x0000000000009fbff] usable
0.000000] BIOS-e820: [mem 0x0000000000009fc00-0x0000000000009ffff] reserved
0.000000] BIOS-e820: [mem 0x000000000000f0000-0x000000000000ffffff] reserved
0.000000] BIOS-e820: [mem 0x00000000000100000-0x000000000007fdffff] usable
0.000000] BIOS-e820: [mem 0x000000000007fe000-0x000000000007ffffff] reserved
0.000000] BIOS-e820: [mem 0x00000000000fffc0000-0x00000000000ffffffff] reserved
0.000000] Notice: NX (Execute Disable) protection missing in CPU!
0.000000] SMBIOS 2.8 present.
0.000000] e820: last_pfn = 0x7fe0 max_arch_pfn = 0x100000
0.000000] MTRR: Disabled
```

User-level application

- FIFO application

`/home/lkc/Desktop/linux-linaro-stable/rootfs/initrd/root/sched/fifo_app.c`

- RR application

`/home/lkc/Desktop/linux-linaro-stable/rootfs/initrd/root/sched/rr_app.c`


FIFO scheduler 트레이스

- trace-cmd 실행
- nfs 폴더로 이동

```
/nfs # trace-cmd record -e sched /root/sched/fifo_app
```


kernelshark

kernelshark 실행 후 분석




RR scheduler 트레이스

- trace-cmd 실행
- nfs 폴더로 이동

```
/nfs # trace-cmd record -e sched /root/sched/rr_app
```

kernelshark

kernelshark 실행 후 분석



Simplified 리눅스 스케줄러

- `kernel/sched/rt.c`
- 기존 리눅스 `real-time scheduler`는 복잡함
- 분석을 위해 단순화 시킴

Key Functions in Scheduling Class

- **enqueue_task**
 - put the task into the run queue
 - increment the nr_running variable
- **dequeue_task**
 - remove the task from the run queue
 - decrement the nr_running variable
- **yield_task**
 - relinquish the CPU
- **check_preempt_curr**
 - check whether the currently running task can be preempted by a new task
- **pick_next_task**
 - choose the most appropriated task
- **load_balance**
 - trigger load balancing code

Simplified 리눅스 스케줄러

- `kernel/sched/rt.c`

enqueue_task_rt()

kernel/sched/rt.c

dequeue_task_rt()

kernel/sched/rt.c

check_preempt_curr_rt()

kernel/sched/rt.c

task_tick_rt()

kernel/sched/rt.c

Next Step.

Energy-aware scheduling: EAS

1. CFS scheduler - Kernel level
2. Load Balancer(Group Scheduling, Bandwidth Control, PELT)
3. EAS features



Reference

- <https://pdos.csail.mit.edu/6.828/2016/schedule.html>
- <http://web.mit.edu/6.033>
- <http://www.rdrop.com/~paulmck/>
- "Is Parallel Programming Hard, And If So, What Can You Do About It?"
- Davidlohr Bueso. 2014. Scalability techniques for practical synchronization primitives. *Commun. ACM* 58

<http://queue.acm.org/detail.cfm?id=2698990>

- "CPUFreq and The Scheduler Revolution in CPU Power Management", Rafael J. Wysocki
- <https://sites.google.com/site/embedwiki/oses/linux/pm/pm-qos>
- <https://intl.aliyun.com/forum/read-916>
- User-level threads : co-routines

<http://www.gamedevforever.com/291>

https://www.youtube.com/watch?v=YYtzQ355_Co

- Scheduler Activations
 - <https://cgi.cse.unsw.edu.au/~cs3231/12s1/lectures/SchedulerActivations.pdf>
- [https://en.wikipedia.org/wiki/FIFO_\(computing_and_electronics\)](https://en.wikipedia.org/wiki/FIFO_(computing_and_electronics))
- <http://jake.dothome.co.kr/>
- <http://www.linuxjournal.com/magazine/completely-fair-scheduler?page=0.0>
- https://www2.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/6_CPU_Scheduling.html
- "Energy Aware Scheduling", Byungchul Park, LG Electronic
- "Update on big.LITTLE scheduling experiments", ARM
- "EAS Update" 2015 september ARM
- "EAS Overview and Integration Guide", ARM TR
- "Drowsy Power Management", Matthew Lentz, SOSP 2015
- <https://www.slideshare.net/nanik/learning-aosp-android-hardware-abstraction-layer-hal>
- <https://www.youtube.com/watch?v=oTGQXqD3CNI>
- <https://www.youtube.com/watch?v=P80NcKUKpuo>
- <https://lwn.net/Articles/398470/>
- "SCHED_DEADLINE: It's Alive!", ARM, 2017