

# CFS 스케줄러

국민대학교 임베디드 연구실  
경 주 현

# Outline

- **CFS scheduler**
- **CFS Load balancer**

# Linux scheduler history

- **Linux 2.4: global queue,  $O(N)$** 
  - Simple
  - Poor performance on multiprocessor/core
  - Poor performance when  $n$  is large
- **Linux 2.5:  $O(1)$  scheduler, per-CPU run queue**
  - Solves performance problems in the old scheduler
  - Complex, error prone logic to boost interactivity

# O(1)

- The Linux scheduler was overhauled completely with the release of kernel 2.6
- O(1) scheduler relies on active and expired arrays of processes
  - To achieve constant scheduling time
- Problem

# O(1)

- The Linux scheduler was overhauled completely with the release of kernel 2.6
- Time slice == nice value
- To achieve constant scheduling time
- Problem
  - No guarantee of fairness
  - the complex heuristics
    - interactive or non-interactive

# Linux scheduler implementations

- **Linux 2.6~4.x: completely fair scheduler (CFS)**
  - Fair
  - Naturally boosts interactivity
- **CFS scheduler defines a fixed time interval during which each thread in the system **must run at least once.****

# Scheduler Policy

- The key decisions made in the scheduler are :

*“how to determine a thread’s **timeslice**? and how to pick the **next thread** to run”*

- **Previous studies : FIFO, Round Robin**
  - Problem : Starvation and unfair

# On a single-CPU system

- CFS is very **simple**.
- An implementation of the weighted fair queueing(WFQ)
- CFS scheduler defines a fixed time interval during which each thread in the system **must run at least once**.
- The interval is divided among **thread's weights**.
  - We call the *time slice*.
  - $\text{Time slice} \leftarrow \text{period} / \text{weight's rate}$



*“how to determine a thread’s **timeslice**? and how to pick the **next thread** to run”*

# CFS's Time slice

- Completely Fair Scheduler (CFS)
- Provide each task CPU time proportional to its weight

***time slice = (Weight of task / Total weight) x period***

# Task Weight

- **priority number must be mapped**
  - kernel by a number between 100 and 139
- **A priority number of 120(nice value 0) = 1024 load**
  - see prio\_to\_weight table

```
struct load_weight {  
    unsigned long weight;  
}; // found in struct sched_entity
```

```
static const int prio_to_weight[40] = {  
    /* -20 */ 88761,    71755,    56483,    46273,    36291,  
    /* -15 */ 29154,    23254,    18705,    14949,    11916,  
    /* -10 */ 9548,     7620,     6100,     4904,     3906,  
    /* -5 */ 3121,     2501,     1991,     1586,     1277,  
    /* 0 */ 1024,      820,      655,      526,      423,  
    /* 5 */ 335,       272,       215,       172,       137,  
    /* 10 */          110,        87,        70,        56,        45,  
    /* 15 */          36,         29,         23,         18,        15,  
};
```

# Nice Values and Task Priority

- **Non-real-time priority**
  - Nice value (-20~19, default 0)
  - A large nice value corresponds to a lower priority
- **Real-time priority**
  - Priority range: 0~99
  - Priority for real-time tasks
    - SCHED\_FIFO, SCHED\_RR
  - A smaller value corresponds to a higher priority



# Task Weight Example

- **A priority: 120, normal task, load: 1024**
- **A and B, running at a priority of 120**
- **Available CPU time**
  - $1024 / (1024 * 2) = 0.5$
- **Increased by one level to 121**
  - $820 ((1024/1.25))$
  - Task A :  $820 / (1024 + 820) = \sim 0.45$
  - Task B :  $(1024 / (1024 + 820)) = \sim 0.55$
- **10% decrease in the CPU time share for Task A.**

# Time slice

- **se.load.weight**
  - weight of scheduling entity
- **cfs\_rq.load.weight**
  - accumulation of weights of all tasks on its run queue


$\text{time\_slice} = (\text{se.load.weight}) / \text{cfs\_rq.load.weight} * \text{sched\_period}();$

# **sched\_slice()**

**kernel/sched/fair.c**

# CFS problem

- **Real-world**
  - CFS basically models an “ideal, precise multitasking CPU” on real hardware.”
  - But in real-world, ideal CPU is nonexistent.
  - Time-sharing system




<http://www.linuxjournal.com/magazine/completely-fair-scheduler?page=0,0>



# CFS problem

- **Real-world**
  - CFS basically models an “ideal, precise multitasking CPU” on real hardware.”
  - But in real-world, ideal CPU is nonexistent.



<http://www.linuxjournal.com/magazine/completely-fair-scheduler?page=0,0>

- **Solution**
  - Virtual Runtime

# vruntime

- How long a process has run

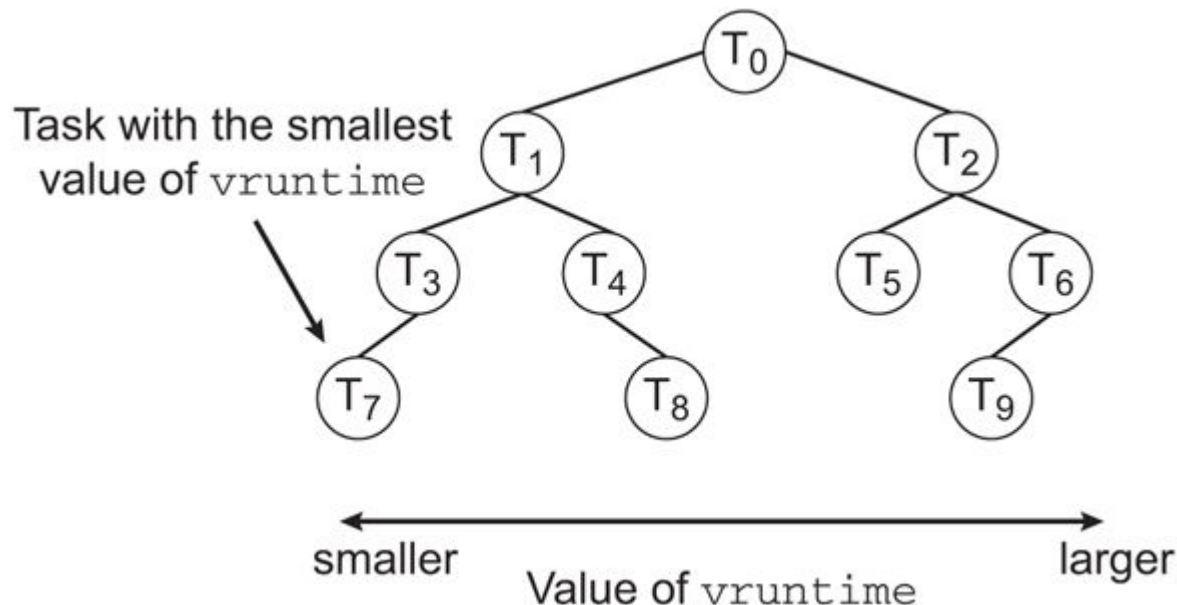
**vruntime +=**

**delta\_exec \* (NICE\_0\_LOAD / curr->load.weight);**

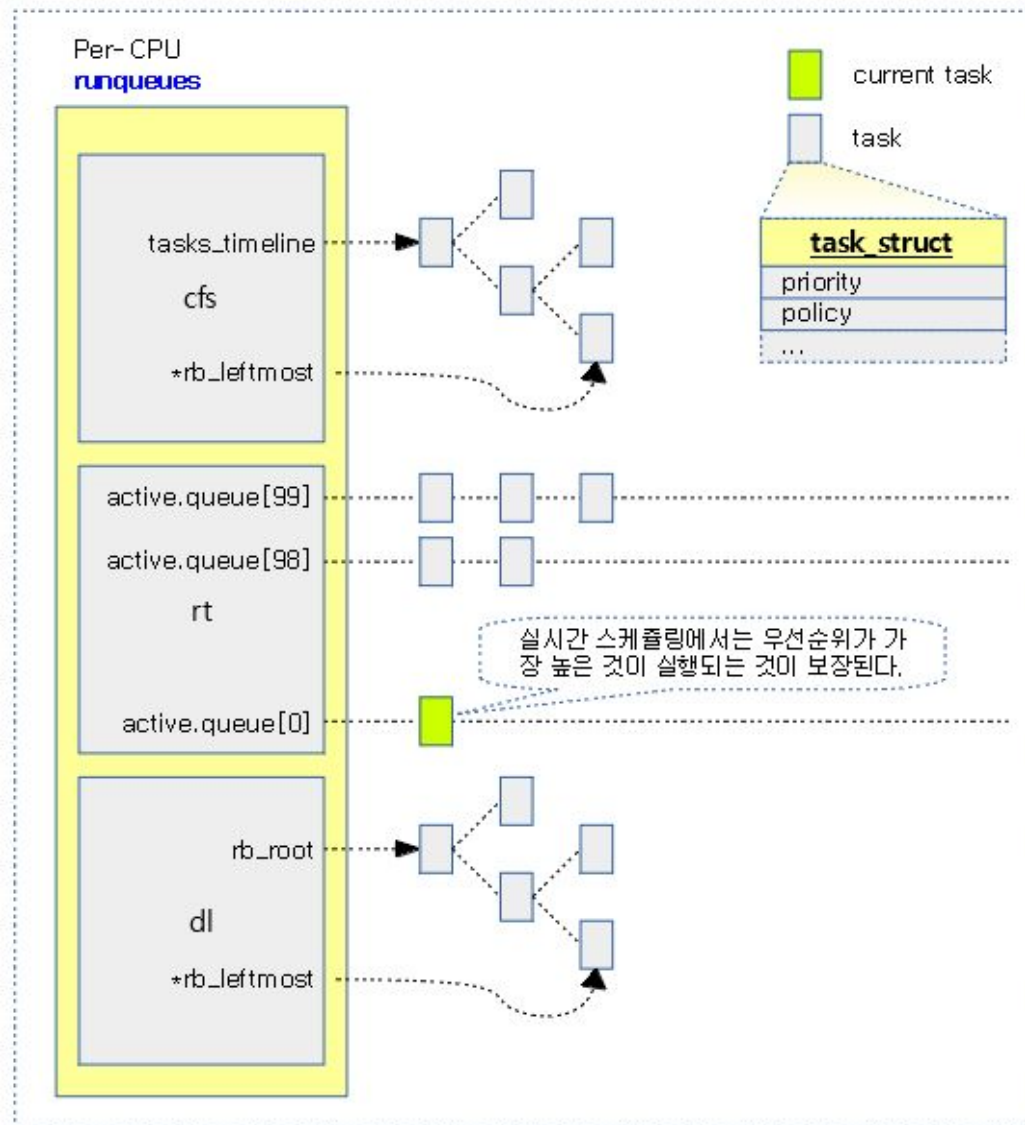
- **delta\_exec**
  - The time spent by the task since the last time vruntime was updated.
- CFS scheduler defines a fixed time interval during which each thread in the system **must run at least once.**

# Red-black tree

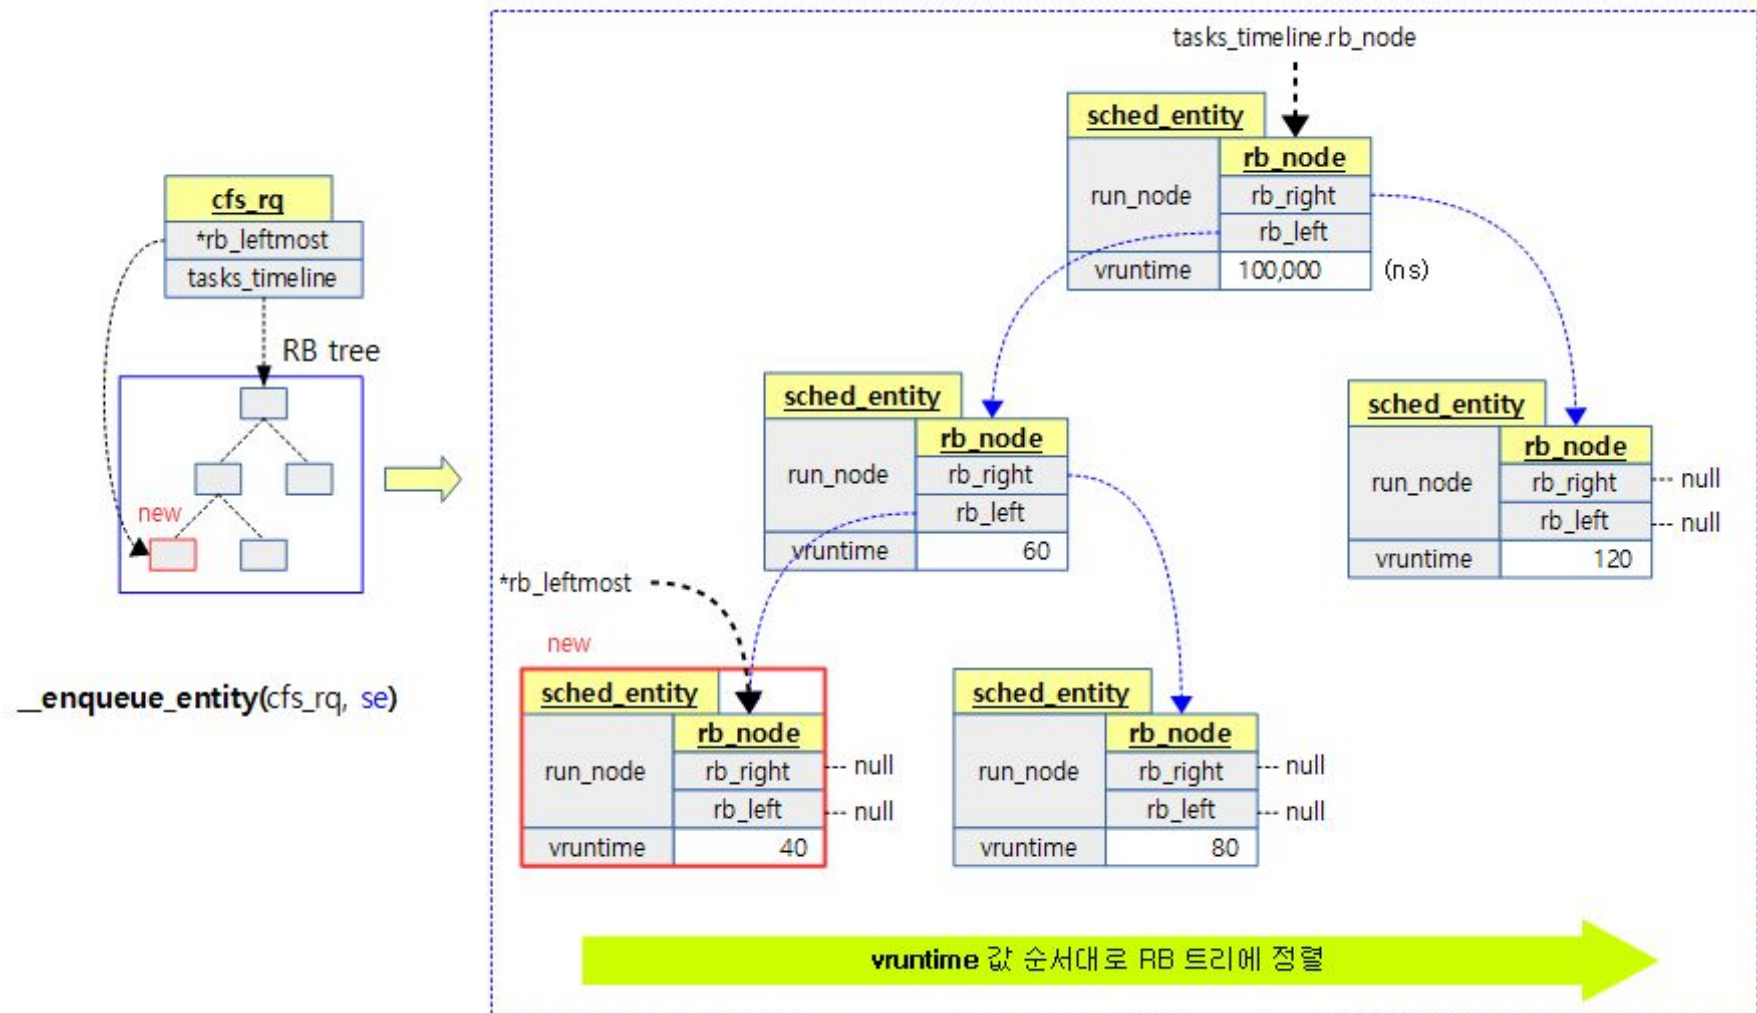
- **Each runnable task is placed in a red-black tree**
  - A balanced binary search tree whose key is based on the value of vruntime
- **The leftmost node has the smallest key value**
  - the task with the highest priority



# Red-black tree



# Red-black tree



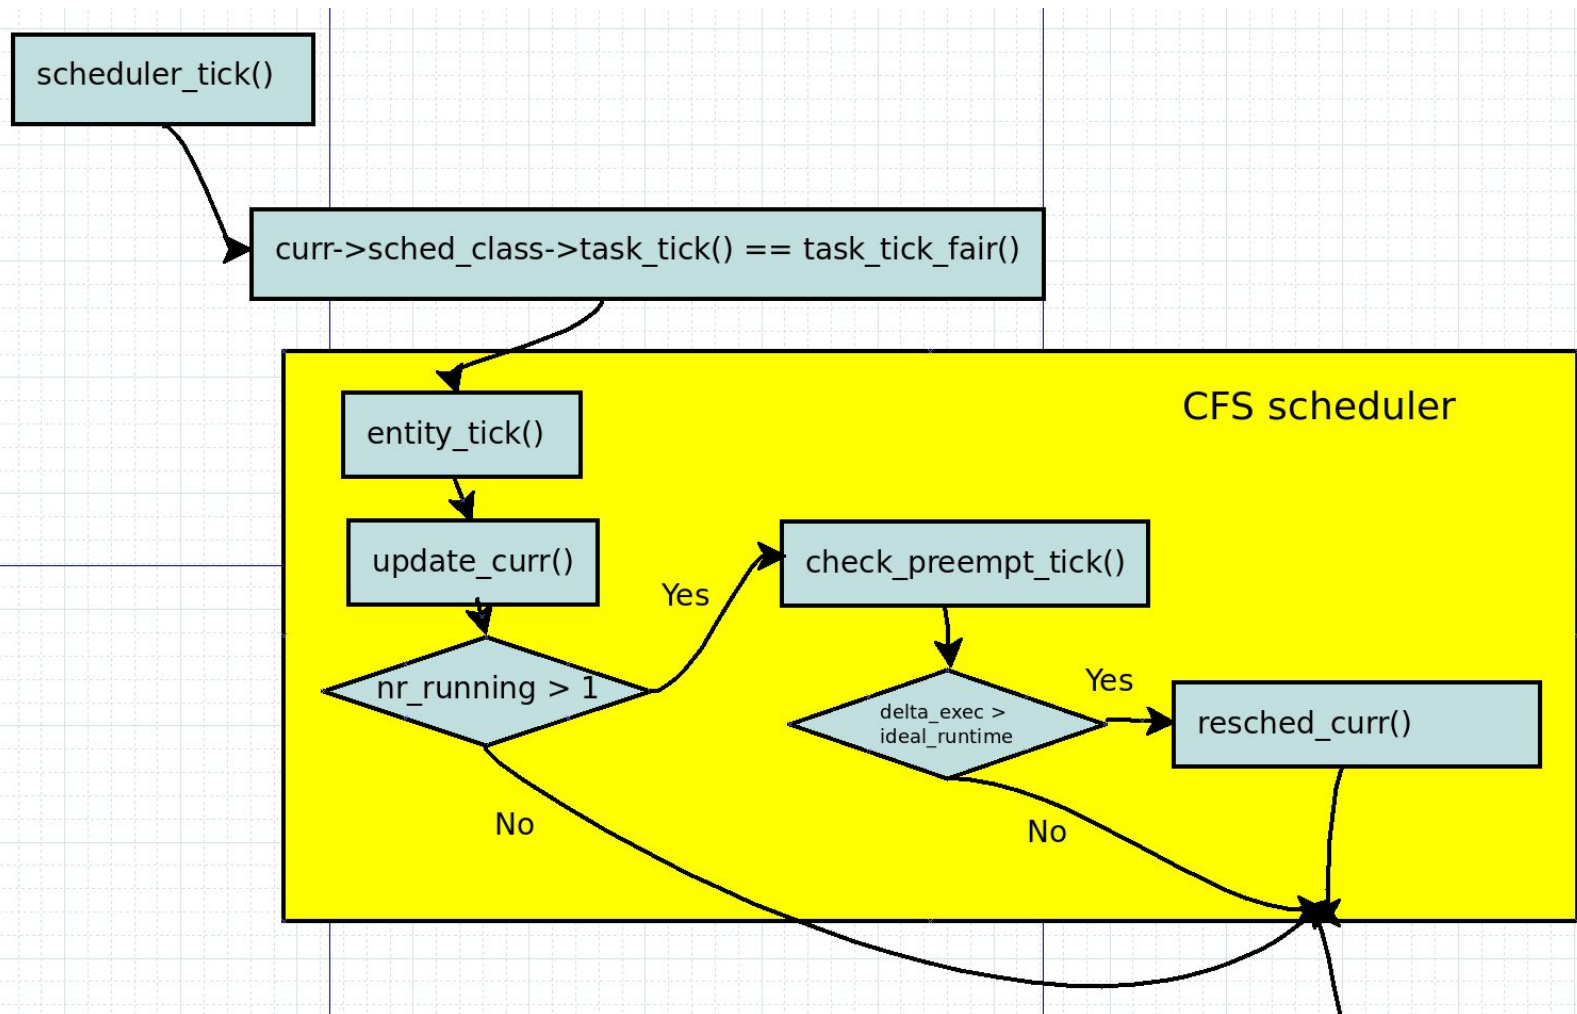
# Linux multiplexing

- **Linux multiplexes by two situations.**
- **Periodically forces a switch**
  - Tick interrupt
- **Sleep and wakeup mechanism**
  - process wait for device or pipe I/O



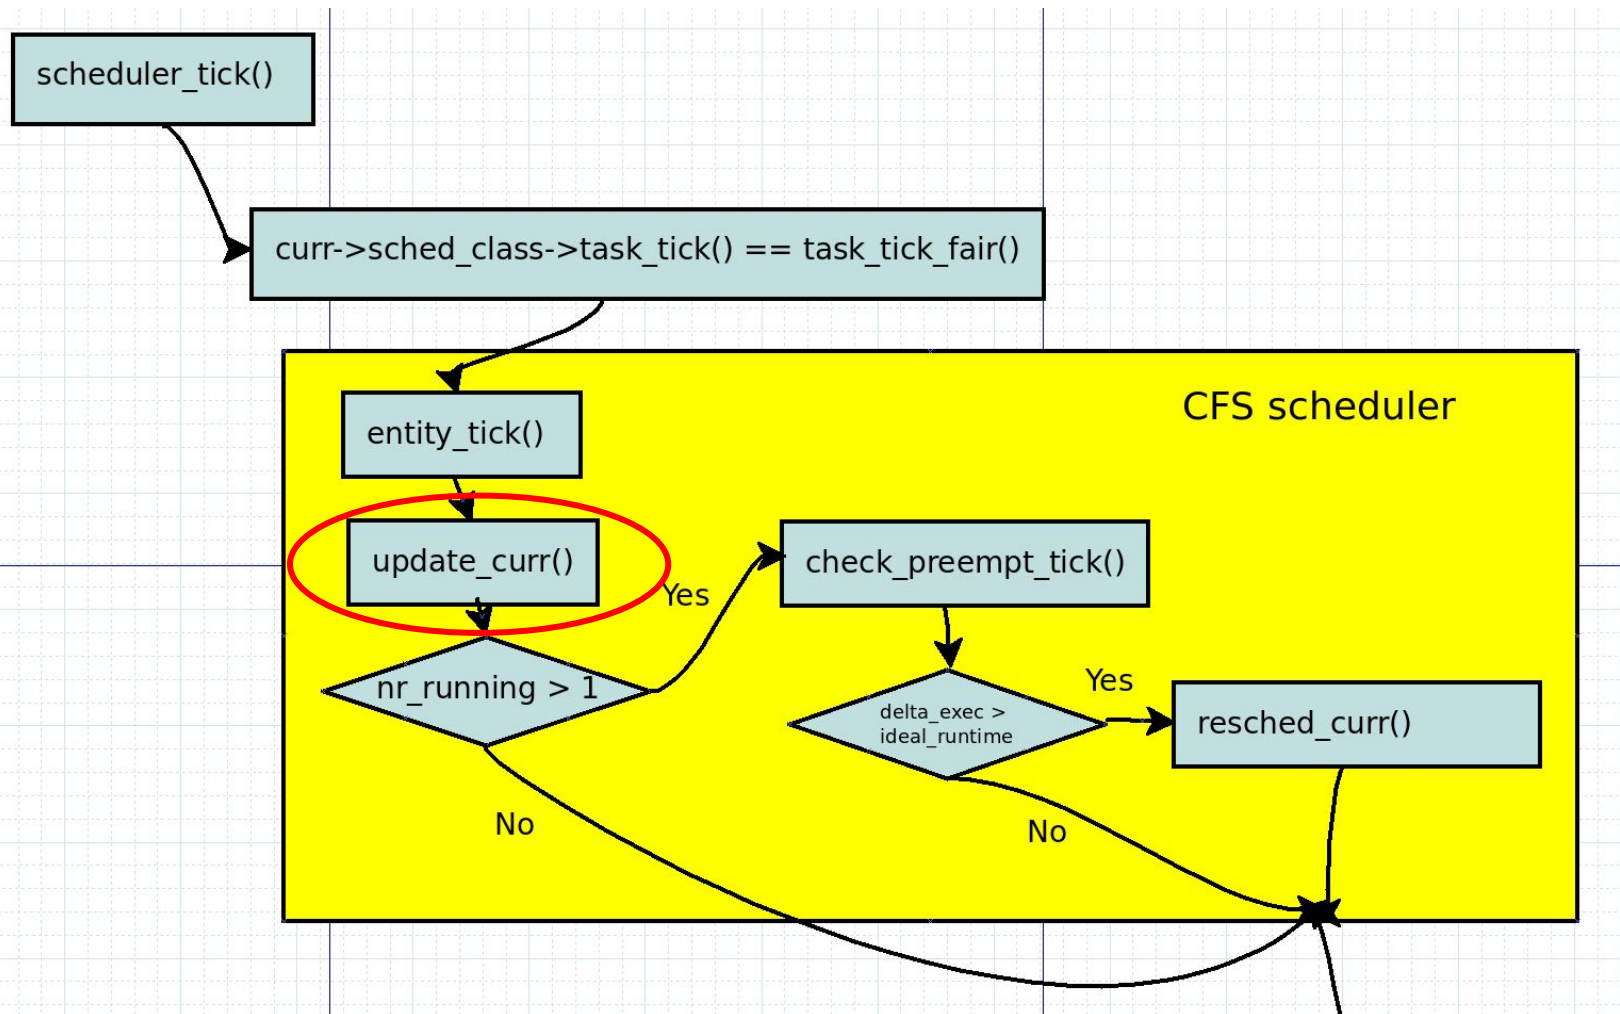
# When schedule a task

- Scheduler Tick



# When schedule a task

- Scheduler Tick



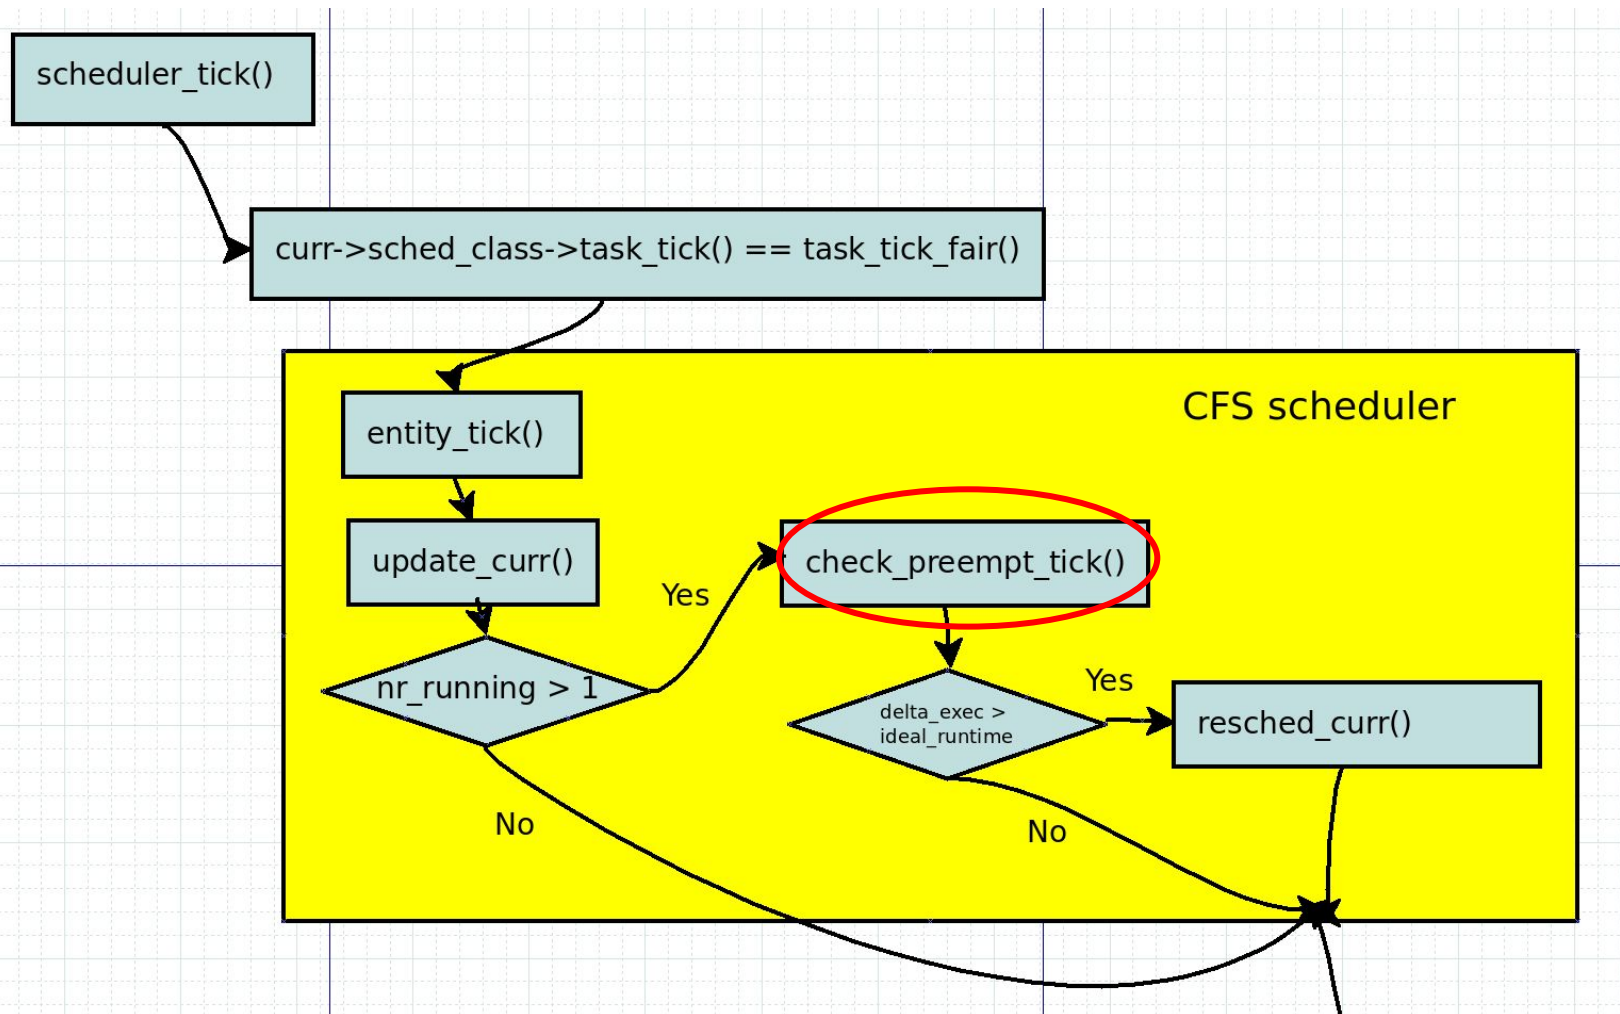


# **update\_curr()**

**kernel/sched/fair.c**

# When schedule a task

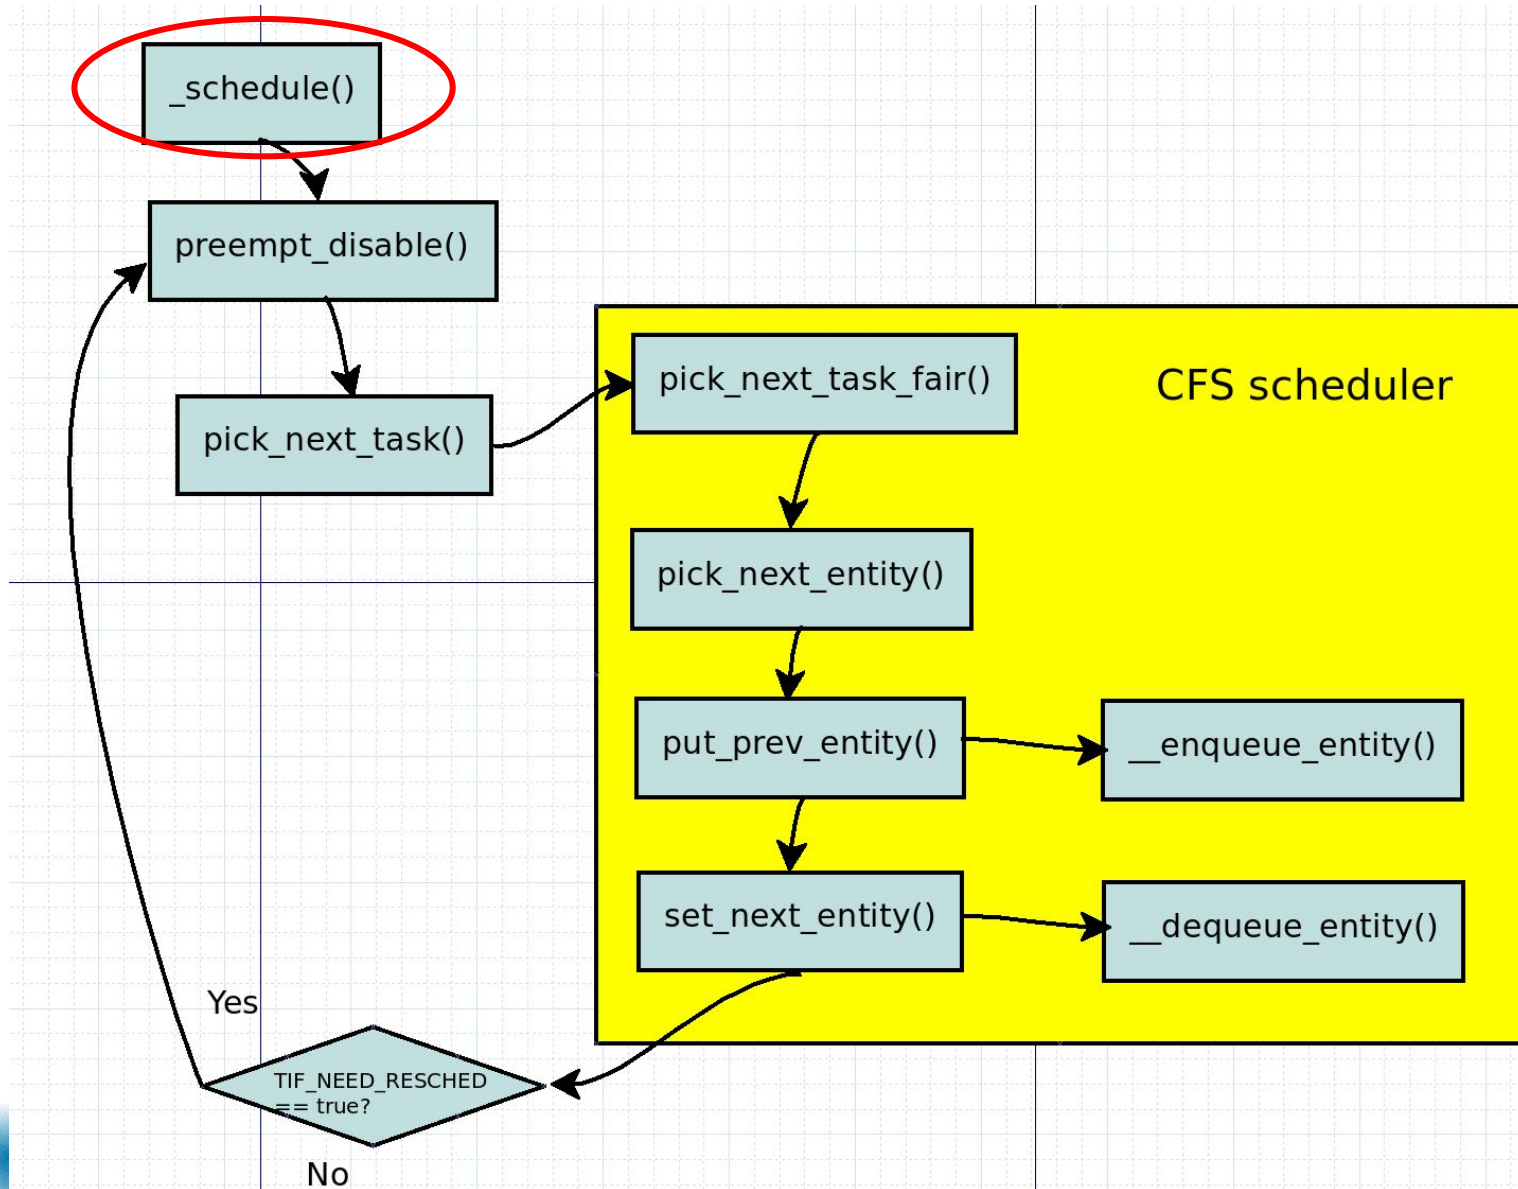
- Scheduler Tick



# **check\_preempt\_tick()**

**kernel/sched/fair.c**

# Scheduling by schedule()



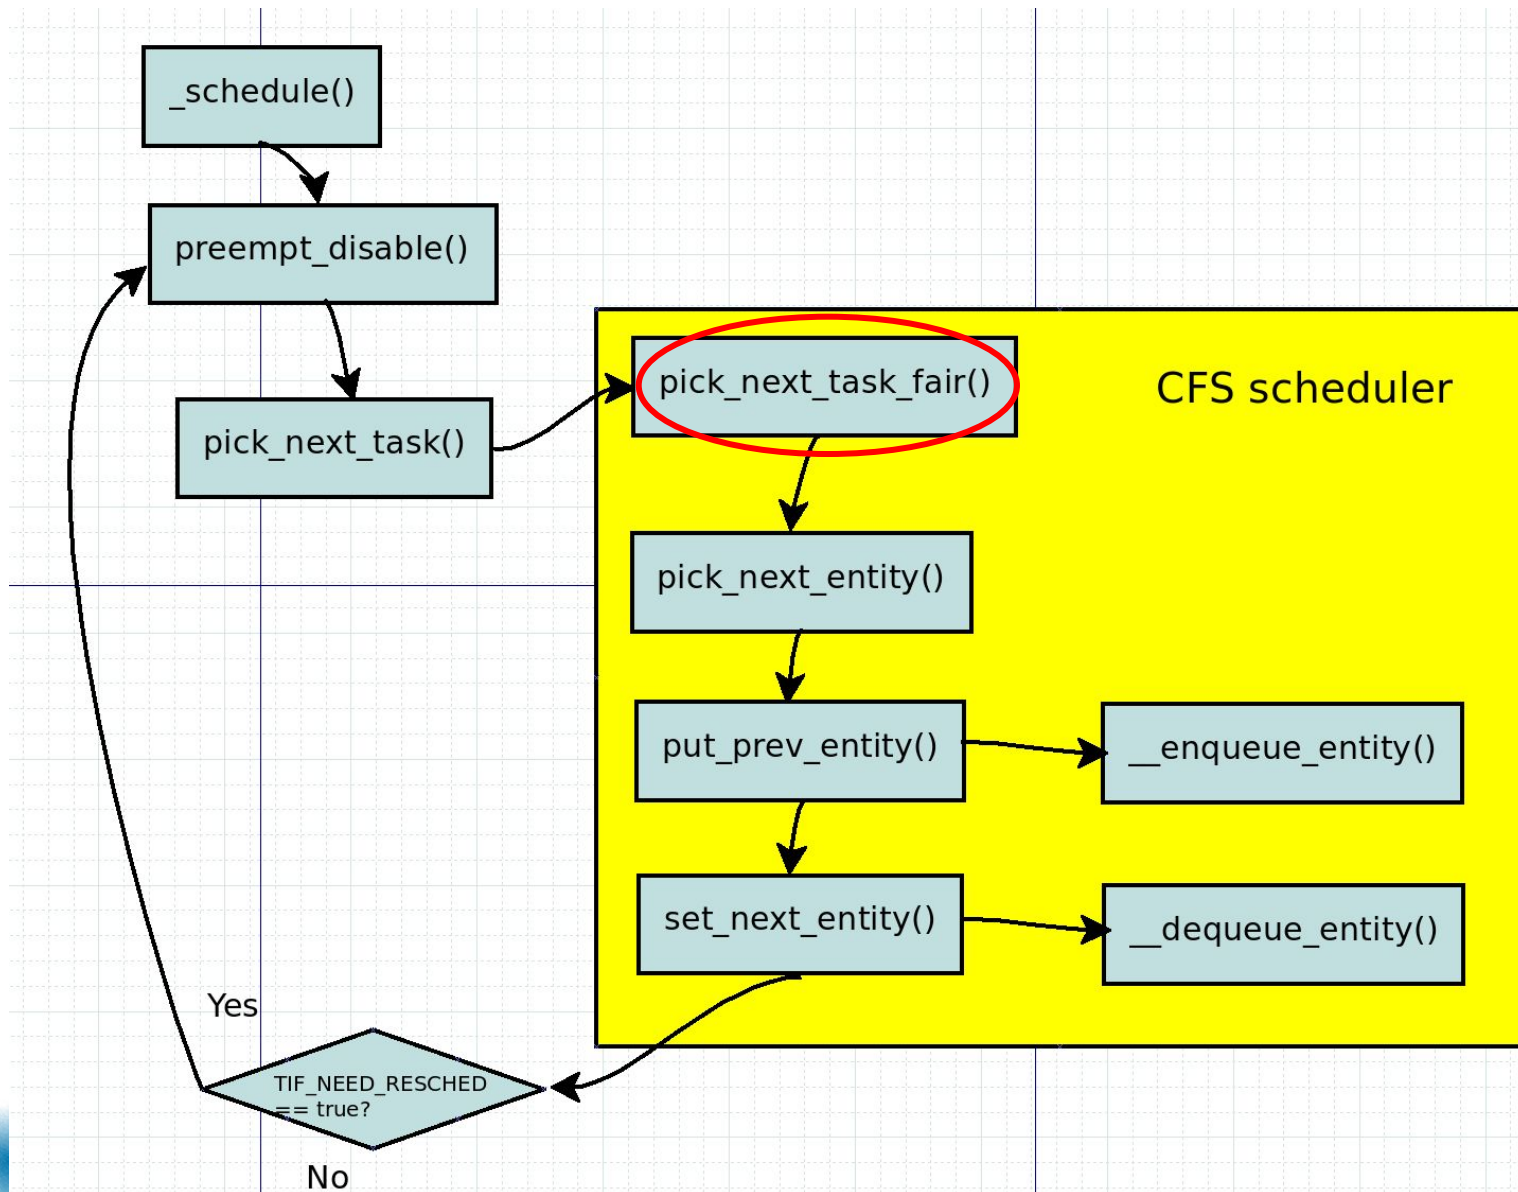
# Linux multiplexing

- Linux multiplexes by two situations.
- Periodically forces a switch
  - Tick interrupt
- Sleep and wakeup mechanism
  - process wait for device or pipe I/O

*“how to determine a thread’s **timeslice**? and how to pick  
the **next thread** to run”*



# Scheduling by schedule()

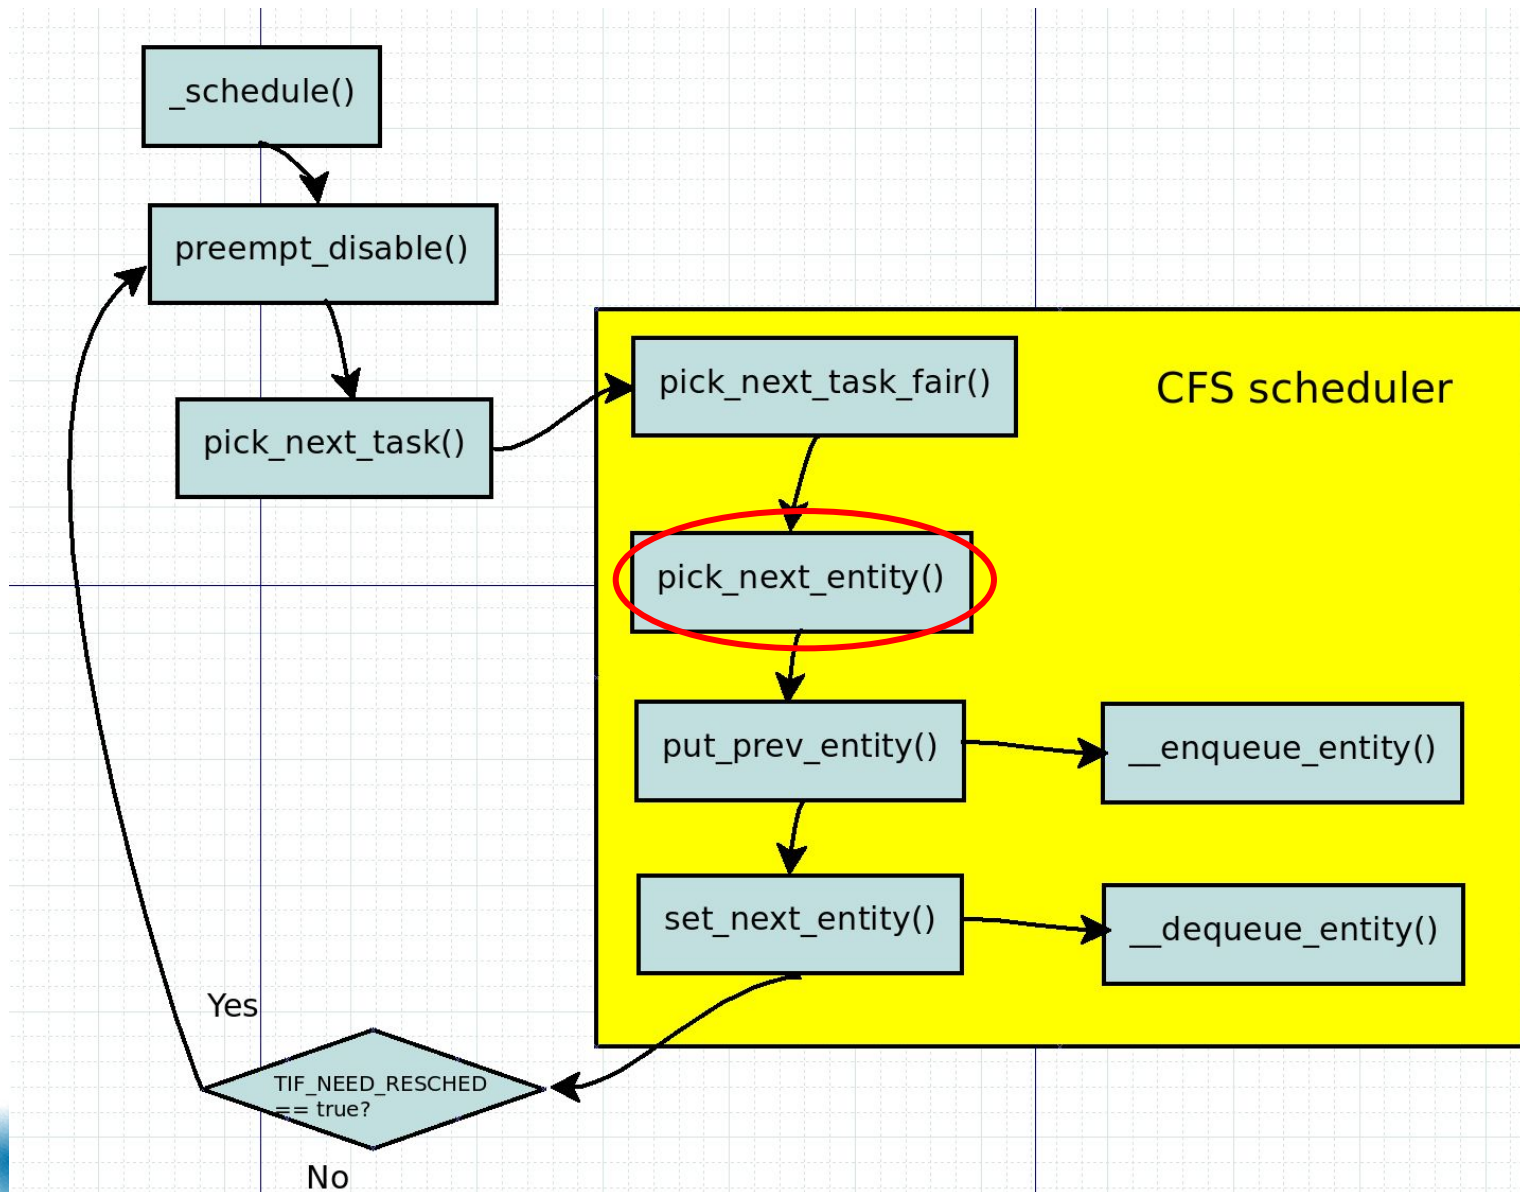


# **pick\_next\_task\_fair()**

**kernel/sched/fair.c**



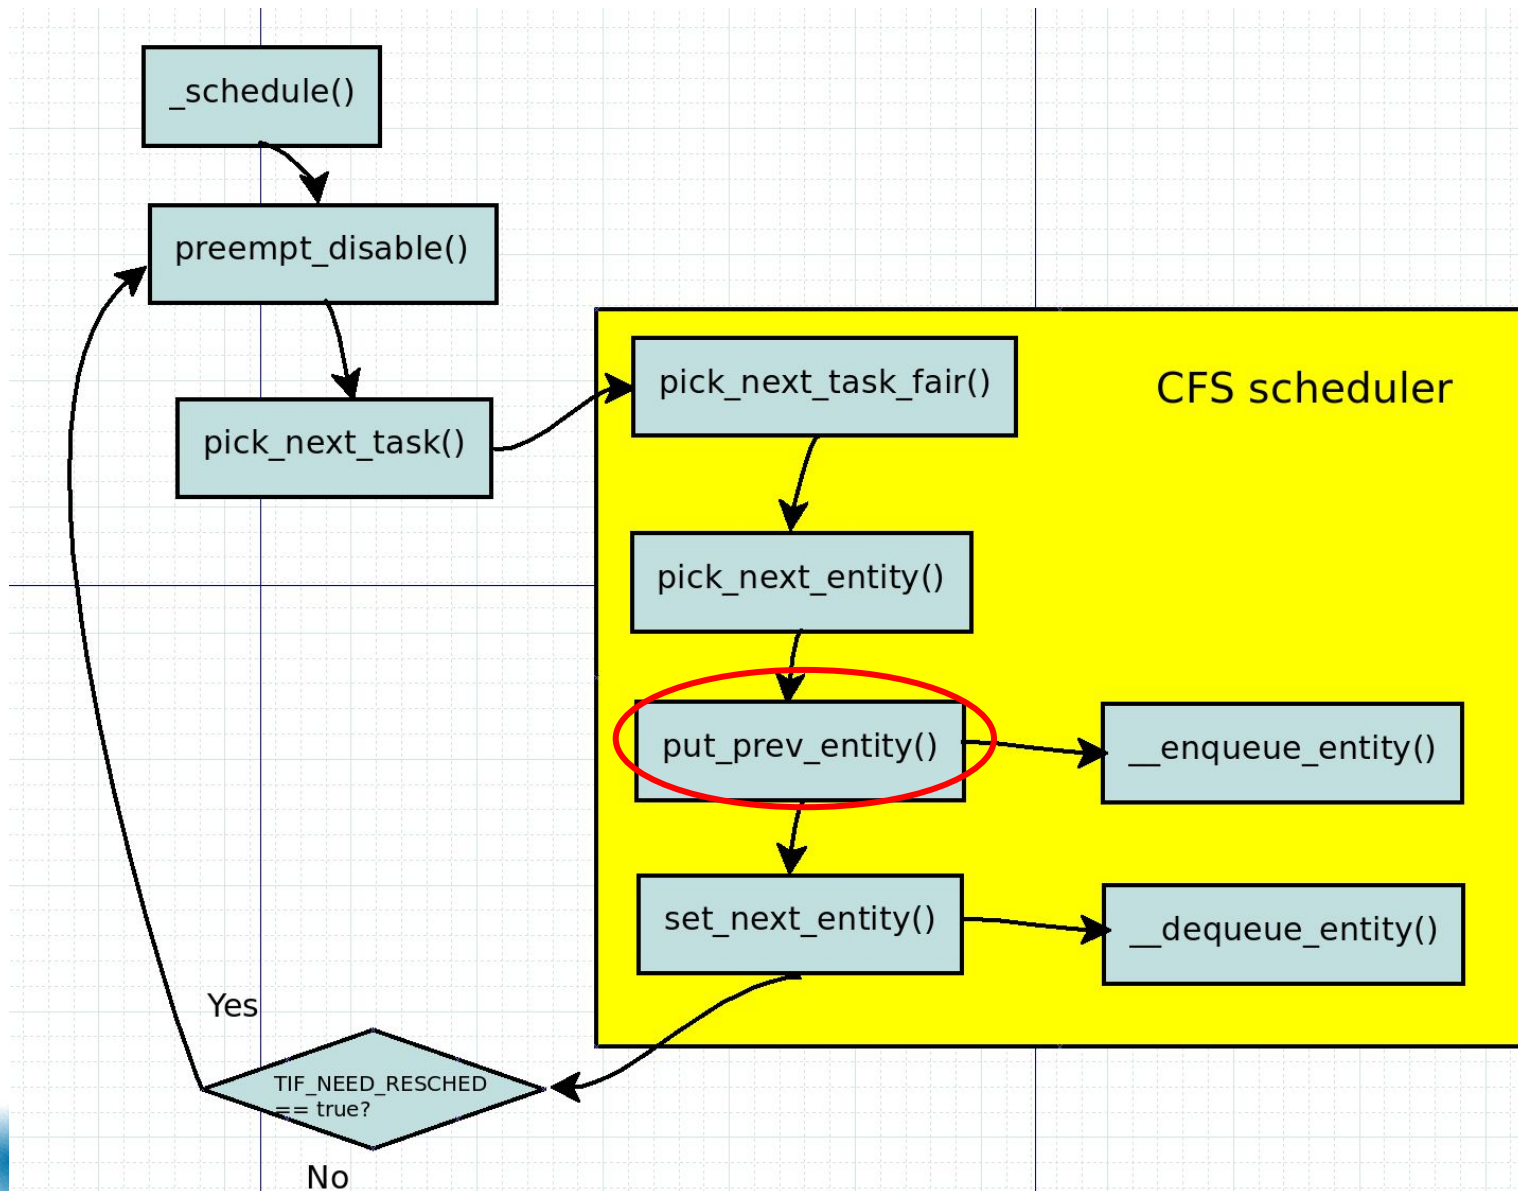
# Scheduling by schedule()



# **pick\_next\_entity()**

**kernel/sched/fair.c**

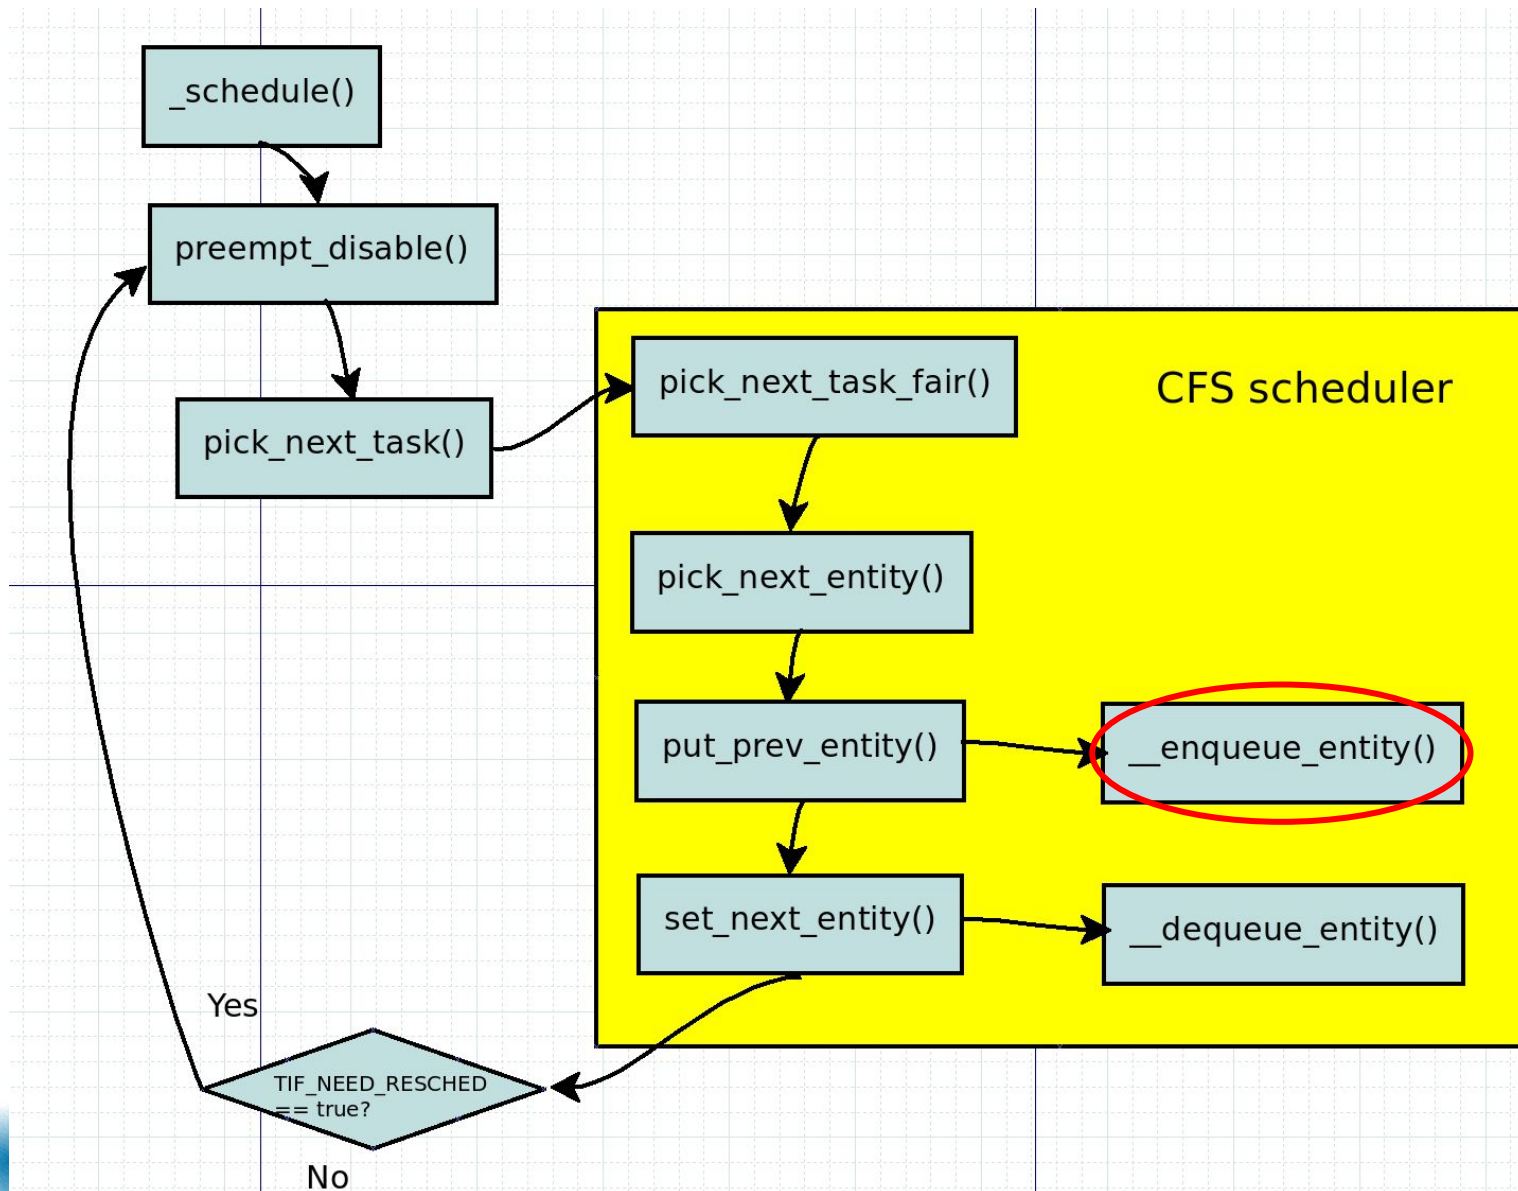
# Scheduling by schedule()



# **put\_prev\_entity()**

**kernel/sched/fair.c**

# Scheduling by schedule()

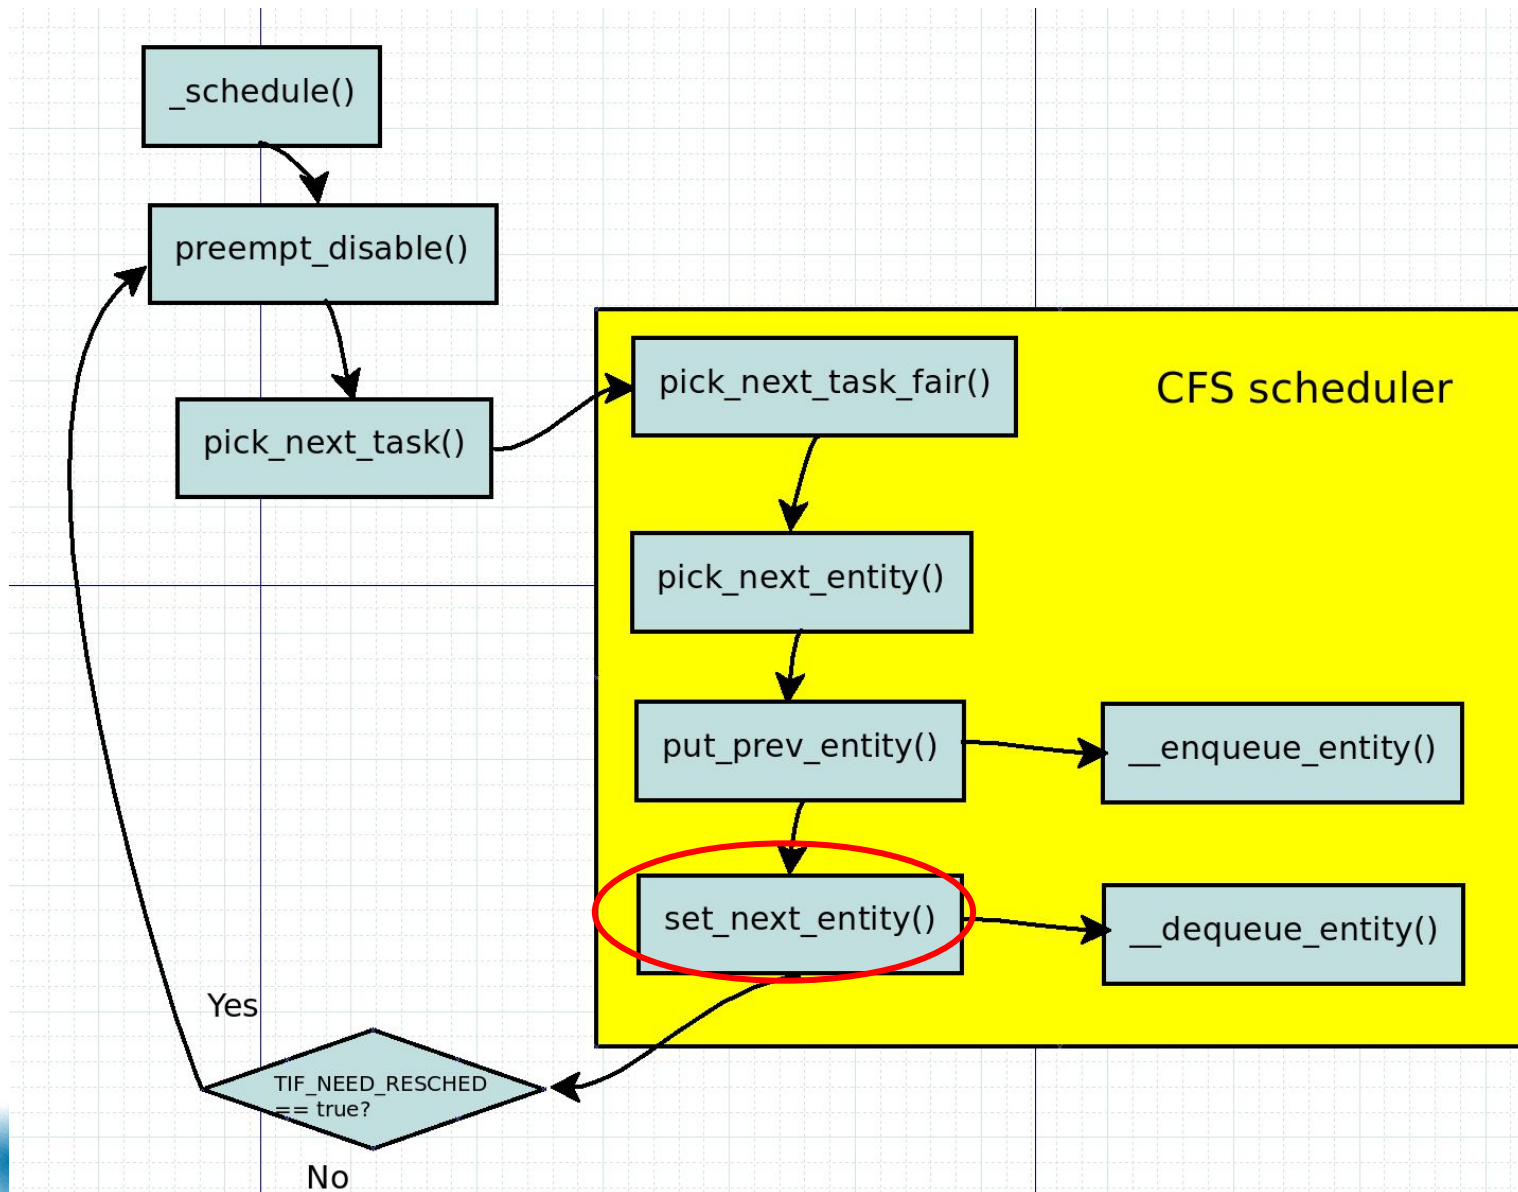




# **\_\_enqueue\_entity()**

**kernel/sched/fair.c**

# Scheduling by schedule()



# **set\_next\_entity()**

**kernel/sched/fair.c**



## Next Step.

# Energy-aware scheduling: EAS

1. CFS scheduler - Kernel level
2. Load Balancer(Group Scheduling, Bandwidth Control, PELT)
3. EAS features



# Reference

- <https://pdos.csail.mit.edu/6.828/2016/schedule.html>
- <http://web.mit.edu/6.033>
- <http://www.rdrop.com/~paulmck/>
- "Is Parallel Programming Hard, And If So, What Can You Do About It?"
- Davidlohr Bueso. 2014. Scalability techniques for practical synchronization primitives. *Commun. ACM* 58

<http://queue.acm.org/detail.cfm?id=2698990>

- "CPUFreq and The Scheduler Revolution in CPU Power Management", Rafael J. Wysocki
- <https://sites.google.com/site/embedwiki/oses/linux/pm/pm-qos>
- <https://intl.aliyun.com/forum/read-916>
- User-level threads : co-routines

<http://www.gamedevforever.com/291>

[https://www.youtube.com/watch?v=YYtzQ355\\_Co](https://www.youtube.com/watch?v=YYtzQ355_Co)

- Scheduler Activations
  - <https://cgi.cse.unsw.edu.au/~cs3231/12s1/lectures/SchedulerActivations.pdf>
- [https://en.wikipedia.org/wiki/FIFO\\_\(computing\\_and\\_electronics\)](https://en.wikipedia.org/wiki/FIFO_(computing_and_electronics))
- <http://jake.dothome.co.kr/>
- <http://www.linuxjournal.com/magazine/completely-fair-scheduler?page=0.0>
- [https://www2.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/6\\_CPU\\_Scheduling.html](https://www2.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/6_CPU_Scheduling.html)
- "Energy Aware Scheduling", Byungchul Park, LG Electronic
- "Update on big.LITTLE scheduling experiments", ARM
- "EAS Update" 2015 september ARM
- "EAS Overview and Integration Guide", ARM TR
- "Drowsy Power Management", Matthew Lentz, SOSP 2015
- <https://www.slideshare.net/nanik/learning-aosp-android-hardware-abstraction-layer-hal>
- <https://www.youtube.com/watch?v=oTGQXqD3CNI>
- <https://www.youtube.com/watch?v=P80NcKUKpuo>
- <https://lwn.net/Articles/398470/>
- "SCHED\_DEADLINE: It's Alive!", ARM, 2017