

CFS Load balancer

Domain, group scheduling, Bandwidth control, PELT load tracking

국민대학교 임베디드 연구실
경 주 현

Outline

- **Load balancer**
- **CFS load balancer**
- **Schedule domain**
- **Group scheduling**
- **Bandwidth control**
- **PELT load tracking**

On multi-core systems

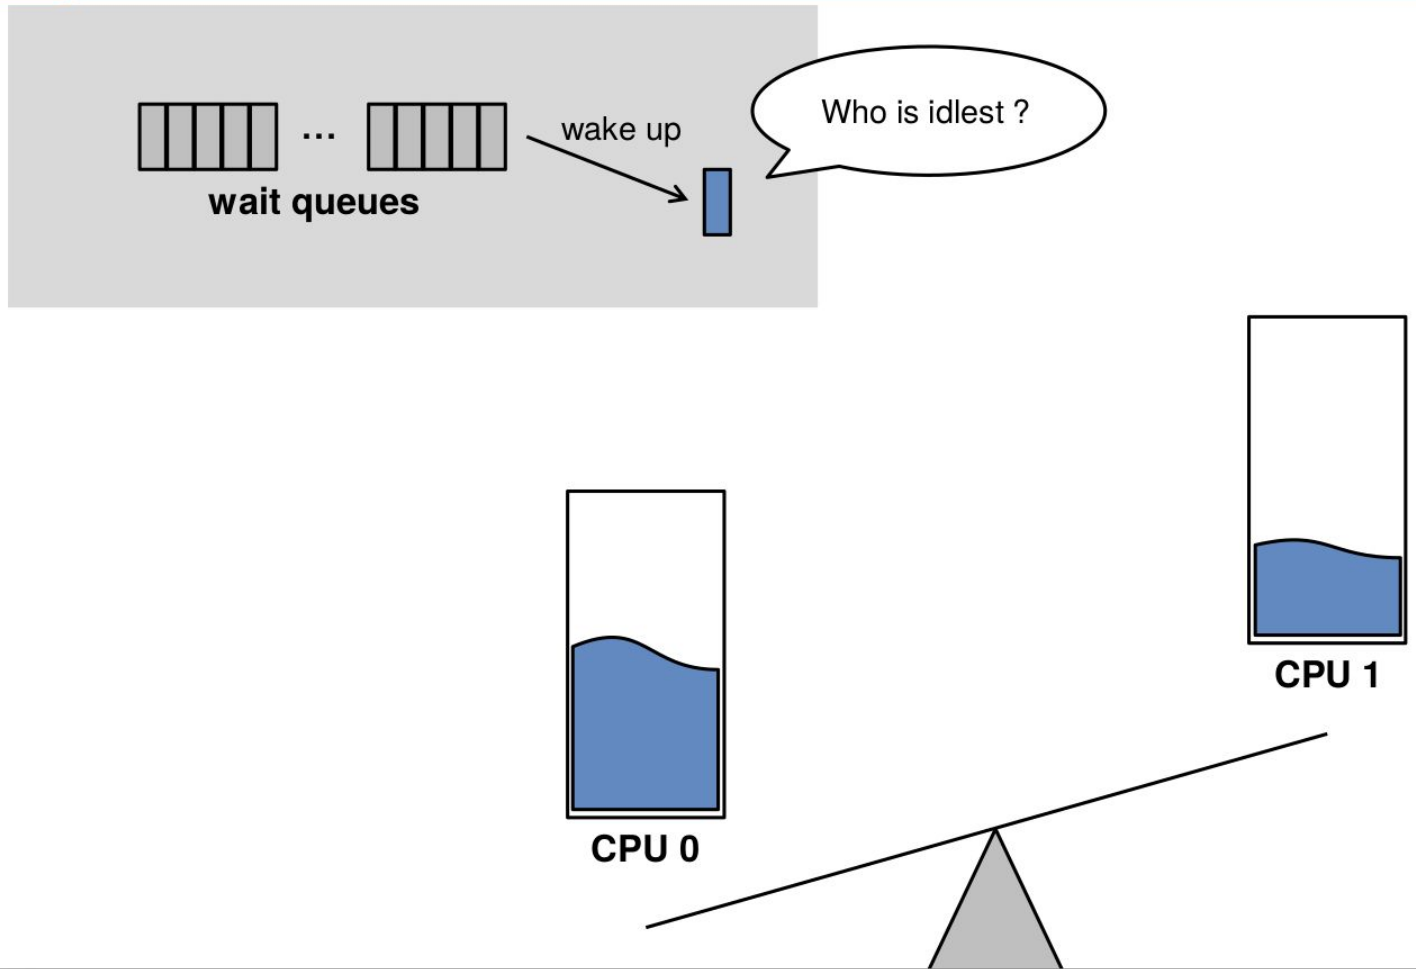
- On a single-CPU system

CFS is very **simple**.

- On multi-CPU systems


CFS becomes **quite complex**

Load Balancing - review



"Energy Aware Scheduling", Byungchul Park, LG Electronic

Load Balancing - review





"Energy Aware Scheduling", Byungchul Park, LG Electronic

Load Balancing - review




BALANCED !!!




"Energy Aware Scheduling", Byungchul Park, LG Electronic

Run over key point again

- **Global shared runqueue.**
 - expensive synchronized access for global data structure



- **per-cpu runqueue.**
 - must be kept balanced -> Load balancing



Load balancing

- Conceptually, load balancing is simple.
- But, modern multicore systems is complex.
 - Cache data structure, power, big.little, NUMA
- To understanding the load balancing algorithm, the **load tracking metric** must be known.

The load tracking metric.

- Load balancer uses to track load.
- Easy metric method

The load tracking metric.

- Load balancer uses to track load.
- Easy metric method
 - Balancing with **Number of threads**
- Number of threads metric
 - problem

The load tracking metric.

- **Load balancer uses to track load.**
- **Easy metric method**
 - Balancing with Number of threads
- **Number of threads metric**
 - problem :
 - queue 1 : high-priority thread queue
 - queue 2 : low-priority thread queue

The load tracking metric.

- **Load balancer uses to track load.**
- **Easy metric method**
 - Balancing with Number of threads
- **Number of threads based load-balancing**
 - problem :
 - queue 1 : high-priority thread queue
 - queue 2 : low-priority thread queue
- **Solution**
 - To balance the queues based on thread's **weights**, not their number.

Balancing with Number of threads

- **Problem**

- One thread is high priority and nine threads are of low priority.
- Queue 1: one thread (high priority)
- Queue 2: nine threads (low priority)
- Work-stealing
 - When high priority thread often sleep.
 - It may frequently steal work from queue2 to queue1.

Balancing with thread's weights

- **Problem**

- One thread is high priority and nine threads are of low priority.
- Queue 1: one thread (high priority)
- Queue 2: nine threads (low priority)
- Work-stealing
 - When high priority thread often sleep.
 - It may frequently steal work from queue2 to queue1.

- **Solution**

Balancing with thread's weights

- **Problem**

- One thread is high priority and nine threads are of low priority.
- Queue 1: one thread (high priority)
- Queue 2: nine threads (low priority)
- Work-stealing
 - When high priority thread often sleep.
 - It may frequently steal work from queue2 to queue1.

- **Solution**

- Do not just based on weights, but based on combination of **thread's weight and its average CPU utilization**.

- **CFS's load**

- Thread's **weight** and its **average CPU utilization**

CFS load balancer

- **Problem**

- Consider**

- multithreading in different processes.
 - process 1 : lots of threads.
 - process 2 : few threads.

- **Result**

CFS load balancer

- **Problem**

- Consider**

- multithreading in different processes.
 - process 1 : lots of threads.
 - process 2 : few threads.

- **Result**

- process 1 : receive a lot more CPU time than process 2
 - process 2 : starvation
 - This would be unfair.

CFS load balancer

- **Problem**

- Consider**

- multithreading in different processes.
 - process 1 : lots of threads.
 - process 2 : few threads.

- **Result**

- process 1 : receive a lot more CPU time than process 2
 - process 2 : starvation
 - This would be unfair.

- **Solution**

- Linux added a **group scheduling feature** to bring fairness between groups of threads

Linux Group Scheduling

- **Group scheduling is enabled**
 - CONFIG_FAIR_GROUP_SCHED
- **A group of tasks is called a “scheduling entity”**

```
struct sched_entity {  
    struct load_weight load;  
    struct sched_entity *parent;  
    struct cfs_rq *cfs_rq;  
    struct cfs_rq *my_rq;  
    struct sched_avg avg;  
    /* ... */  
};
```

Group Scheduling Data Structure

- Scheduling is always at the granularity of `sched_entity`
- A single task becomes a scheduling entity

```
struct task_struct {  
    struct sched_entity se;  
    /* ... */  
};
```

- Each scheduling entity contains a runqueue.

```
struct cfs_rq {  
    struct load_weight load;  
    unsigned long runnable_load_avg;  
    unsigned long blocked_load_avg;  
    unsigned long tg_load_contrib;  
    /* ... */  
};
```

Group Scheduling Data Structure

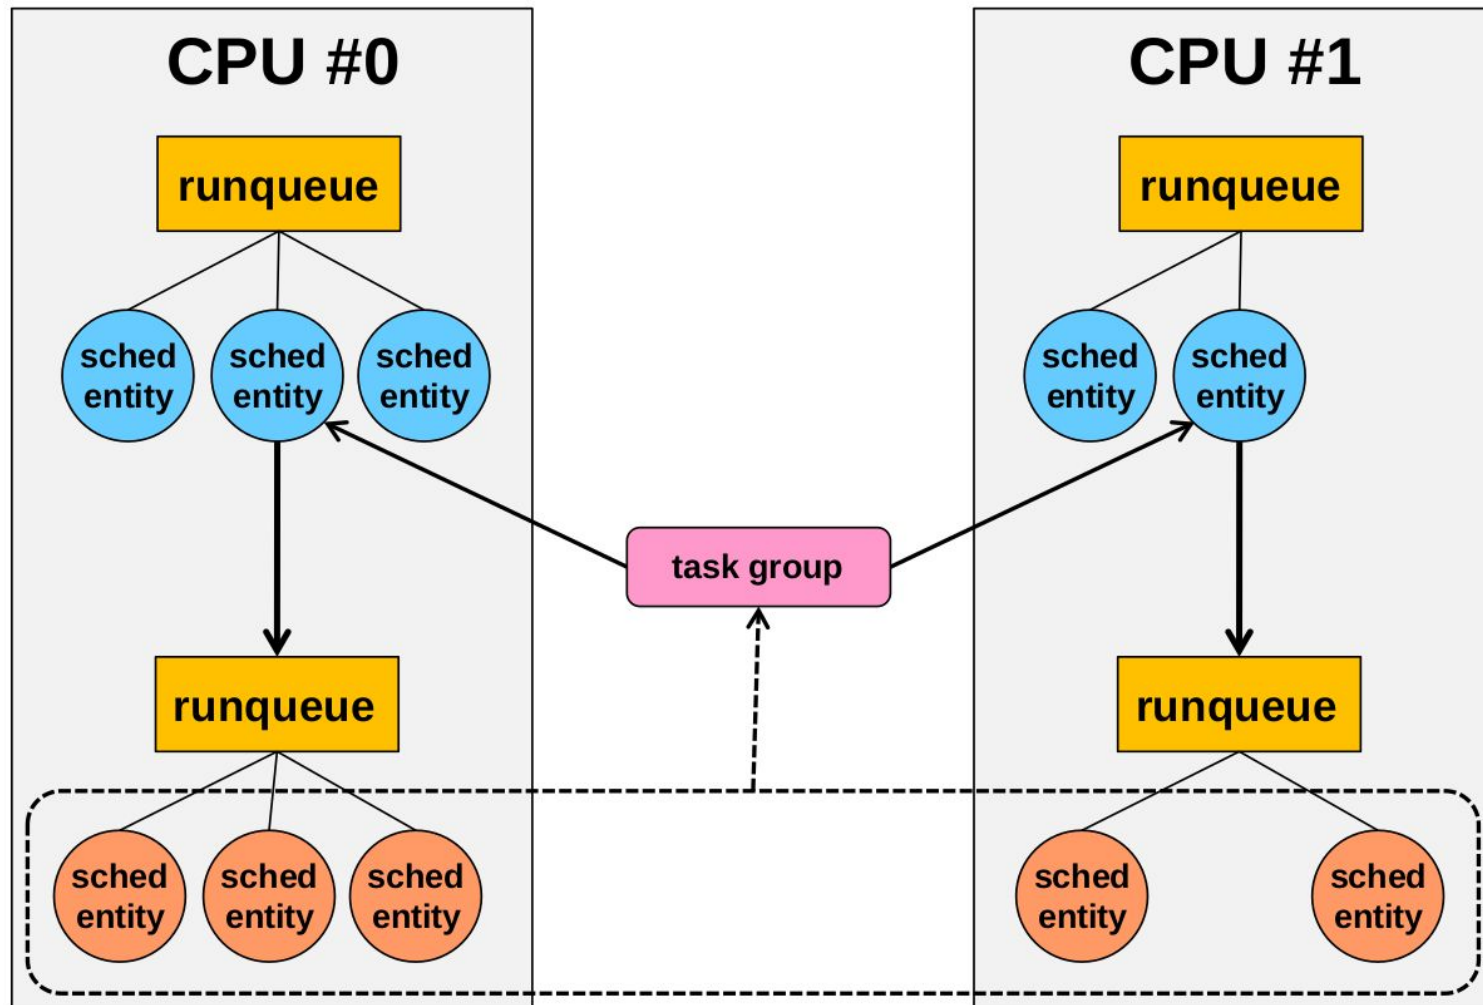
- **Extend the concept to multiprocessor**

- ```
struct task_group {
 struct sched_entity **se;
 struct cfs_rq **cfs_rq;
 unsigned long shares;
 atomic_long_t load_avg;
 /* ... */
};
```

- ***c* is CPU NUM**

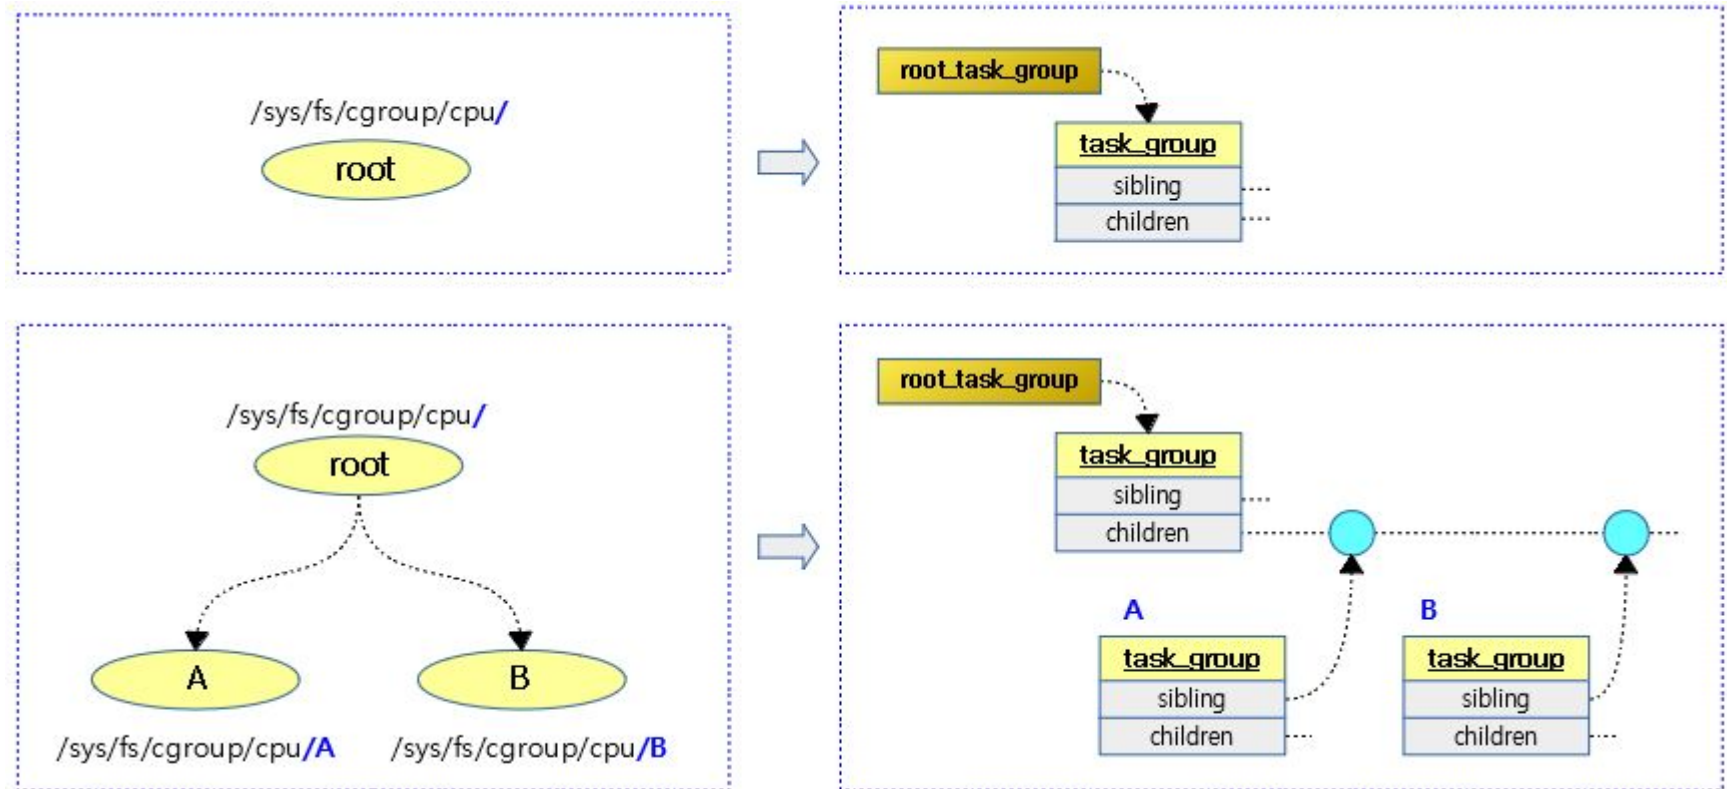
```
tg->se[c] = &se;
tg->cfs_rq[c] = &se->my_rq;
```

# Group Scheduling Data Structure



[http://rtcc.hanyang.ac.kr/rtccw/?page\\_id=1745](http://rtcc.hanyang.ac.kr/rtccw/?page_id=1745)

# Group Scheduling Data Structure

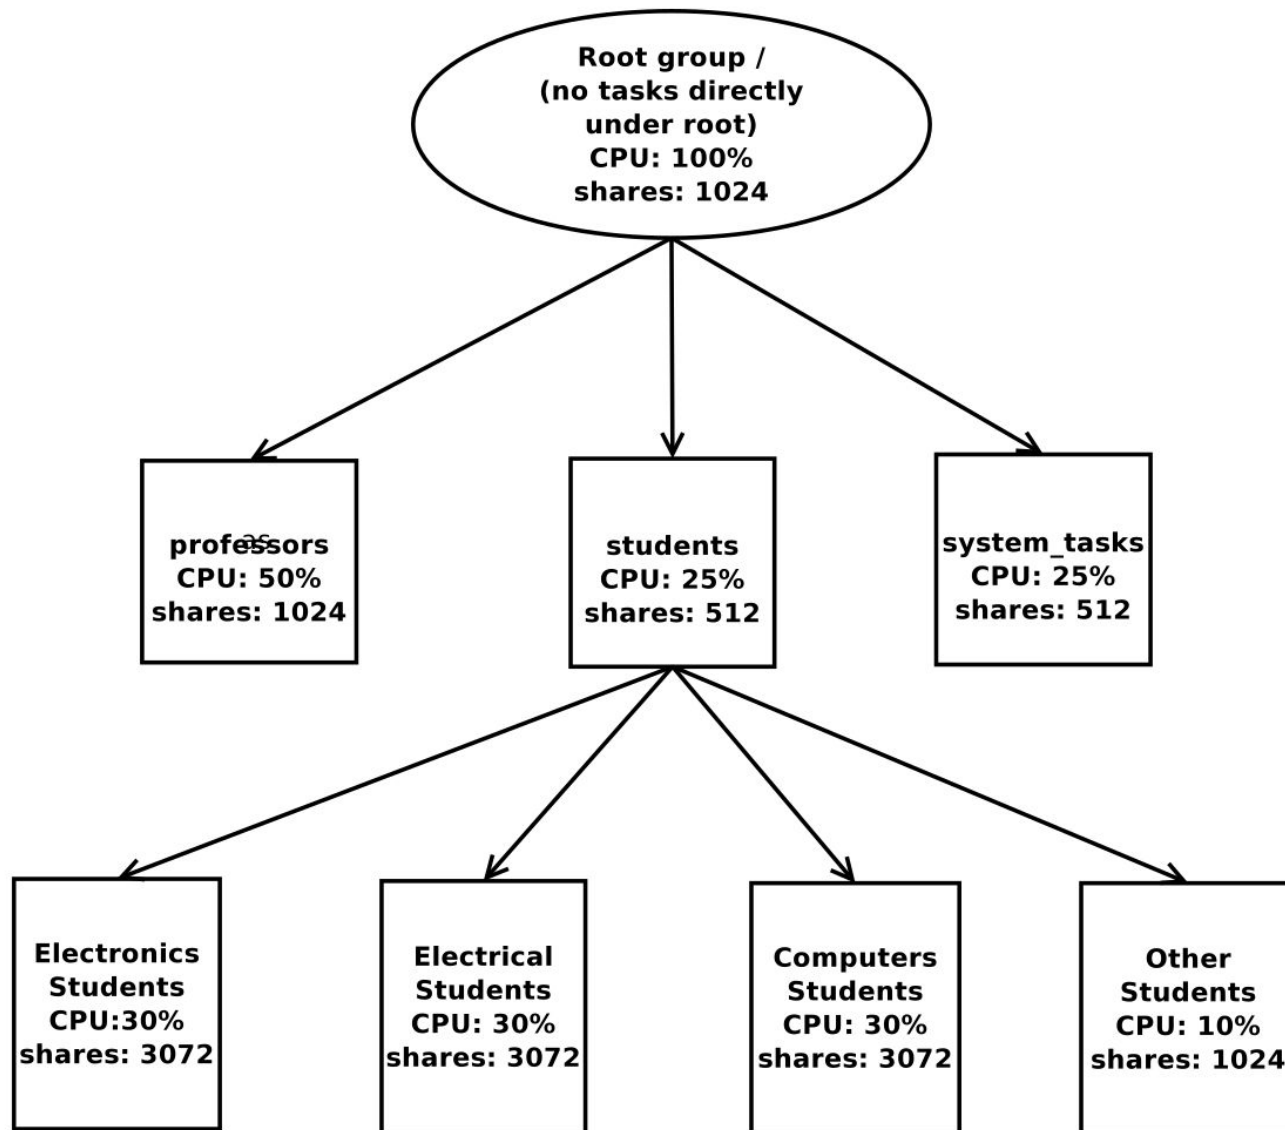








# Linux Group Scheduling



# CFS load balancer

- **Problem**

- When there is free CPU time available
  - The CPU go idle X
- The CFS scheduler will give any left-over time to other processes
- Sometimes system administrator may want to limit maximum share of CPU time
- The CFS scheduler cannot limit CPU time.

# CFS bandwidth control feature

- **Problem**

- When there is free CPU time available
  - The CPU go idle X
- The CFS scheduler will give any left-over time to other processes
- Sometimes system administrator may want to limit maximum share of CPU time
- The CFS scheduler cannot limit CPU time.

- **Solution**

- Linux added a **CFS bandwidth control feature** to limit  
Limiting the maximum share of CPU time that a process  
(or group of processes).


# CFS bandwidth control

- **Example use cases**
  - Virtual Machines
  - Pay-per-use

# CFS bandwidth control - two knobs

- **cpu.cfs\_period\_us :**
  - The period over which the group's CPU usage is to be regulated
- **cpu.cfs\_quota\_us :**
  - how much CPU time is available to the group over that period

# CFS bandwidth control - two knobs



- <http://jake.dothome.co.kr/>

# CFS load balancer

- **Problem**

- Hardware topologies are becoming more varied, accommodating different power/performance budgets:
  - SMP, NUMA, ARM big.LITTLE technology.
- Modern hardware composed of hierarchy levels of subsystems

# Scheduling domain

- **Problem**

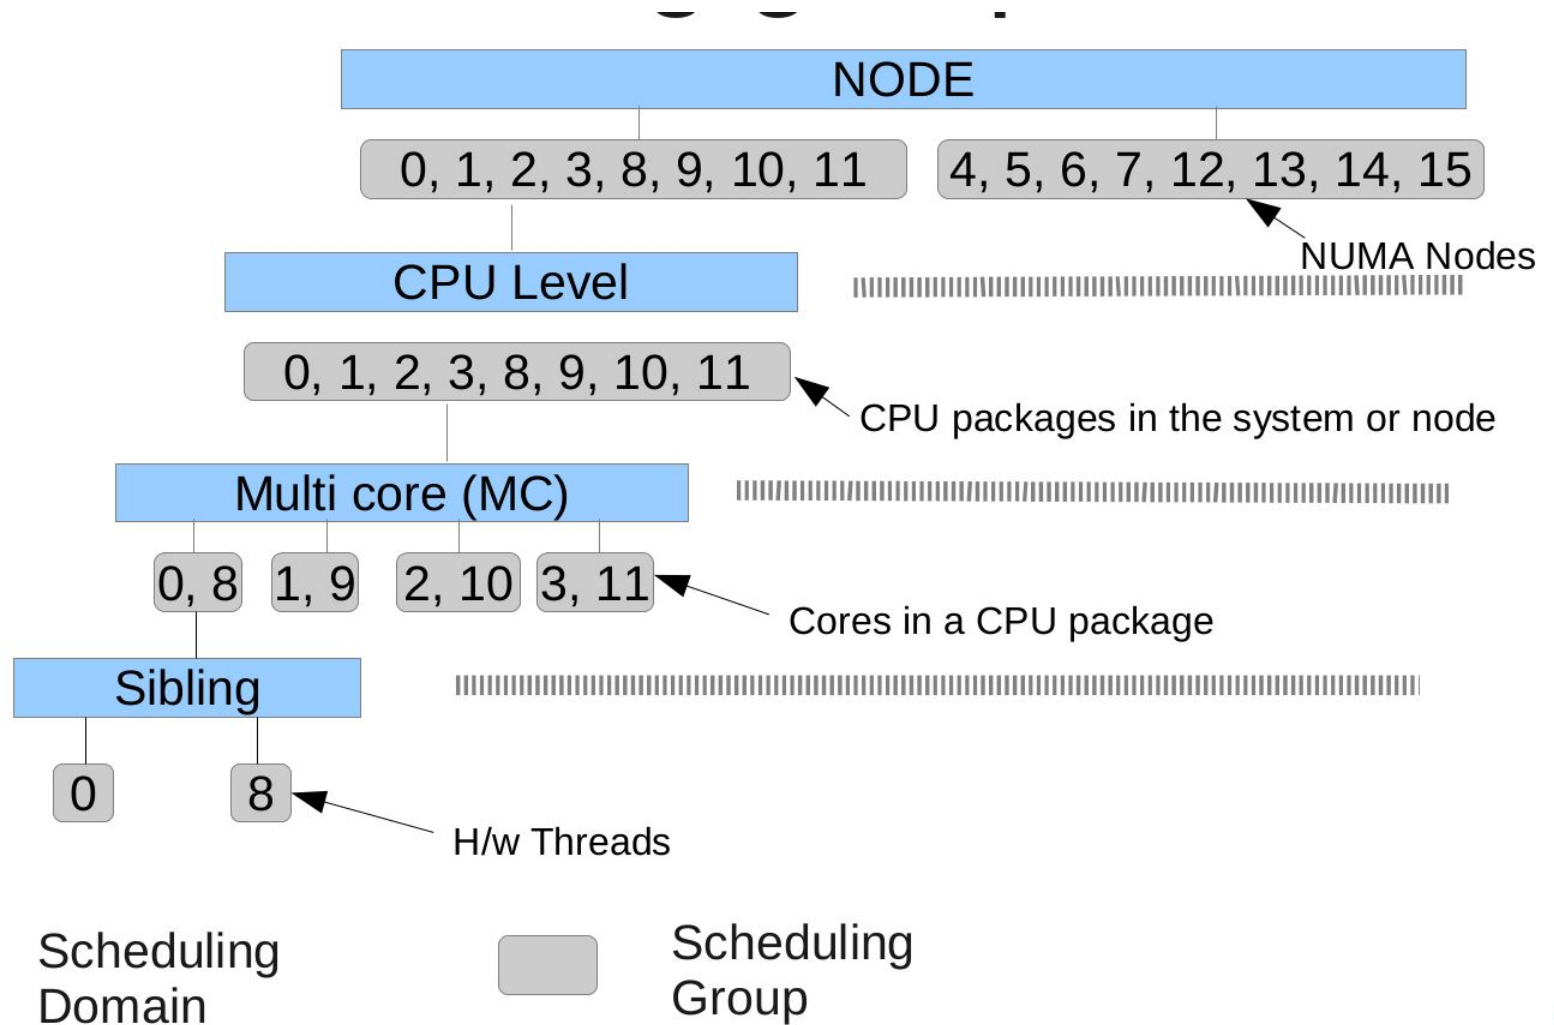
- Hardware topologies are becoming more varied, accommodating different power/performance budgets:
  - SMP, NUMA, ARM big.LITTLE technology.
- Modern hardware composed of hierarchy levels of subsystems

- **Solution**

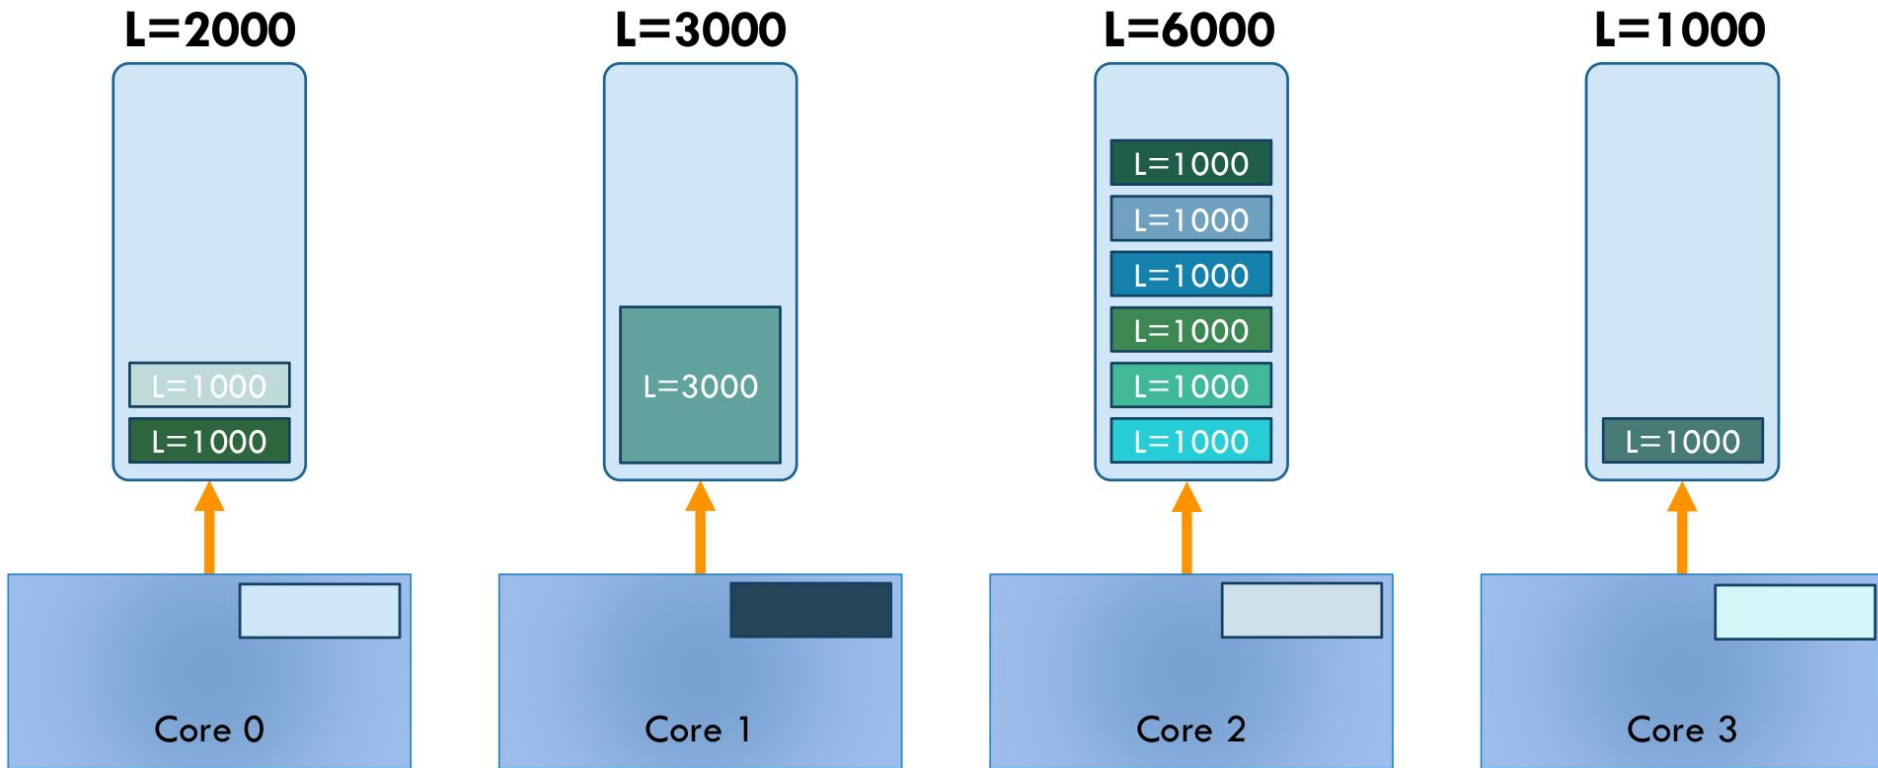
- Linux added a **scheduling domain feature** to cover various multi-core systems such as cache architecture, NUMA, heterogeneous systems.



# CFS Scheduling Domains

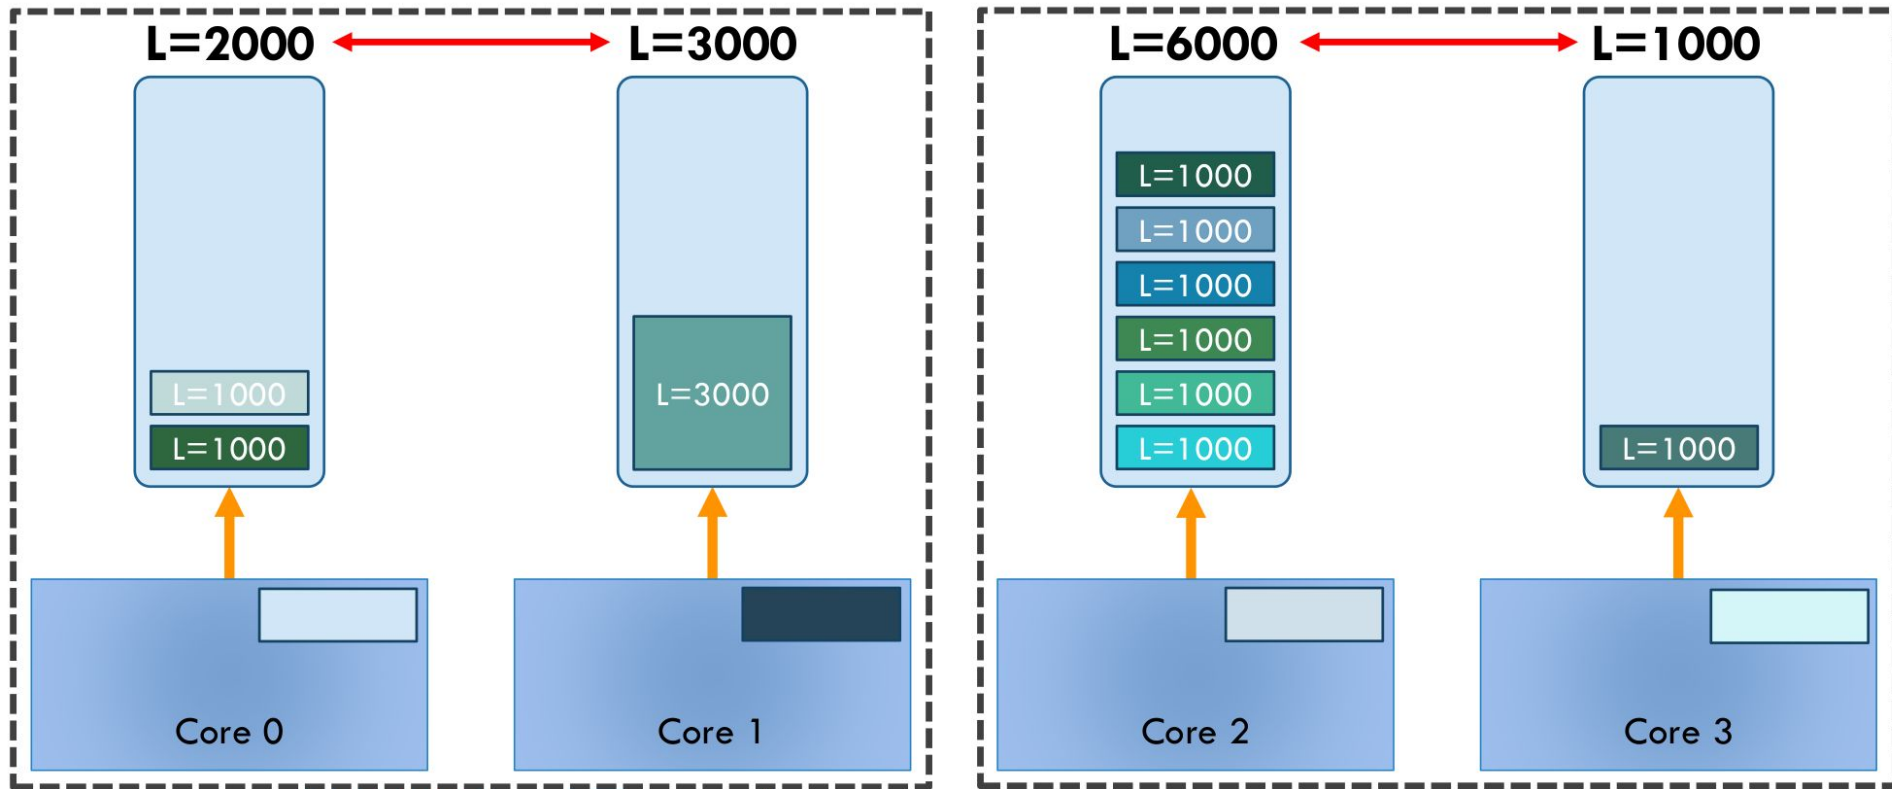


# HIERARCHICAL LOAD BALANCING



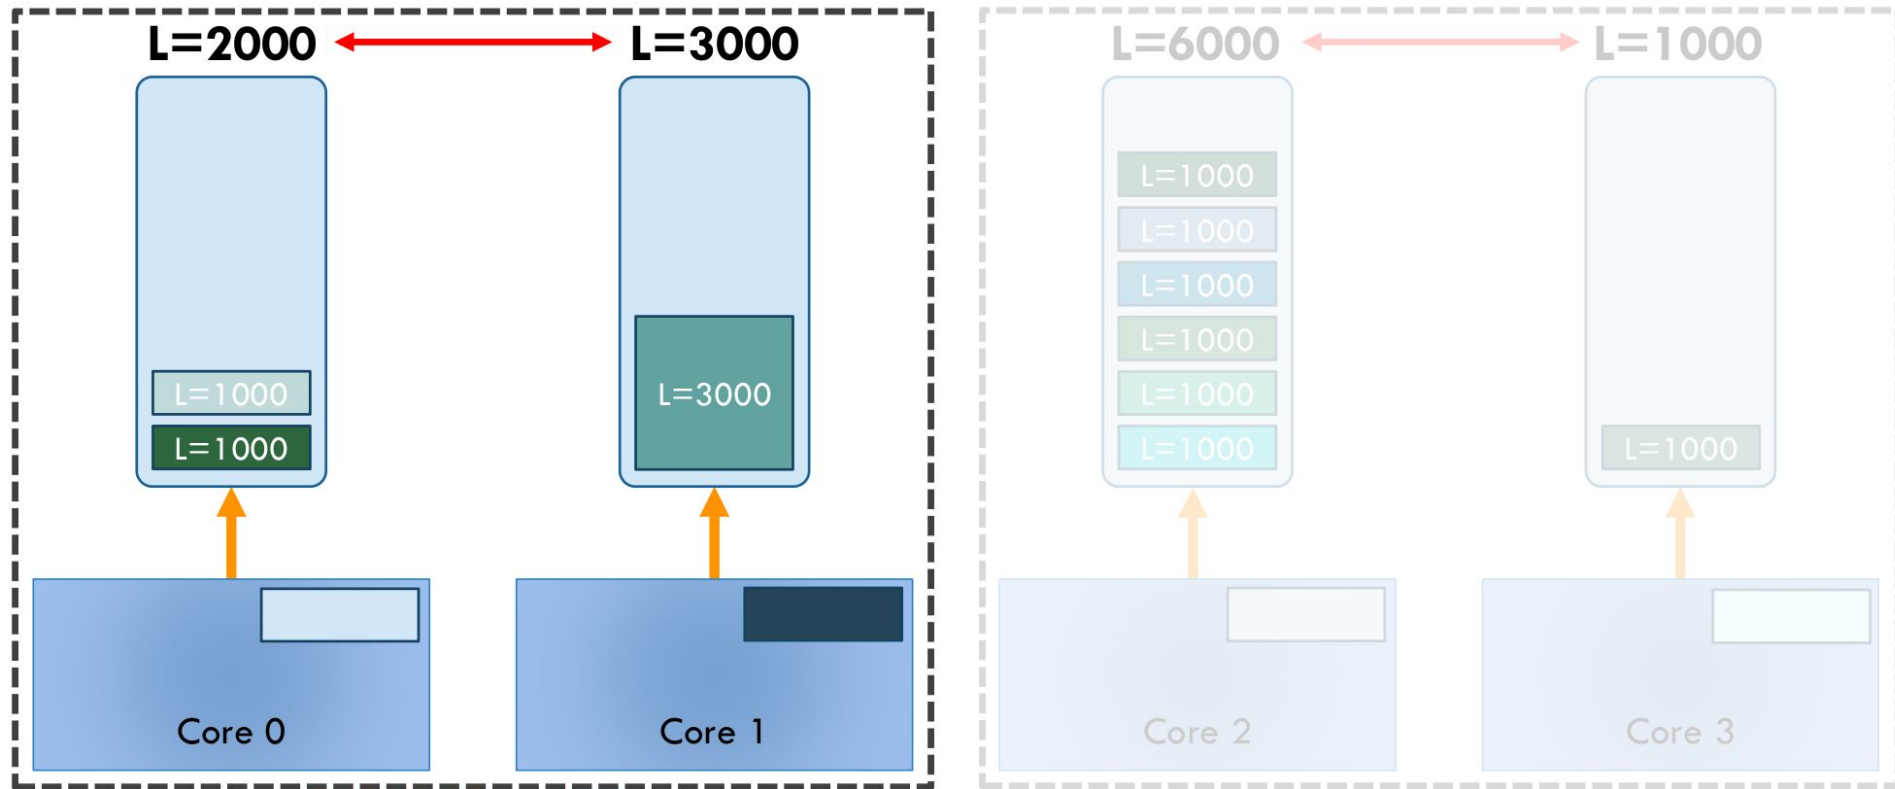
“THE LINUX SCHEDULER:  
A DECADE OF WASTED CORES”, EuroSys’16, Jean-Pierre Lozi, et.,al.

# HIERARCHICAL LOAD BALANCING



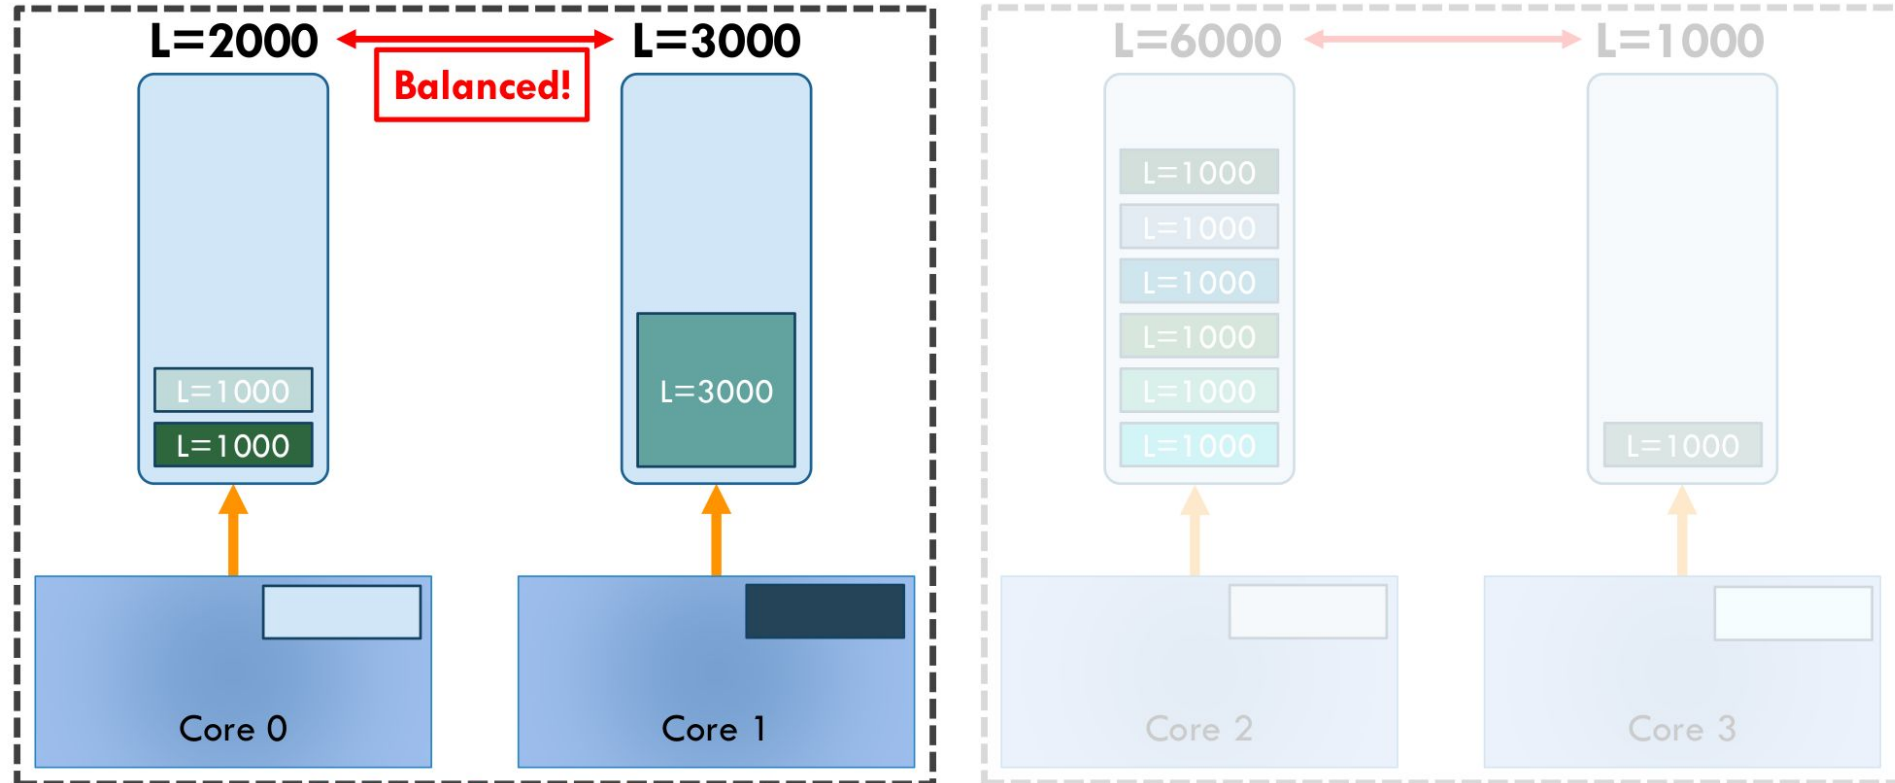
“THE LINUX SCHEDULER:  
A DECADE OF WASTED CORES”, EuroSys’16, Jean-Pierre Lozi, et.,al.

# HIERARCHICAL LOAD BALANCING



“THE LINUX SCHEDULER:  
A DECADE OF WASTED CORES”, EuroSys’16, Jean-Pierre Lozi, et.,al.

# HIERARCHICAL LOAD BALANCING





“THE LINUX SCHEDULER:  
A DECADE OF WASTED CORES”, EuroSys’16, Jean-Pierre Lozi, et.,al.

# **rebalance\_domains()**

**kernel/sched/fair.c**





# HIERARCHICAL LOAD BALANCING



"THE LINUX SCHEDULER:  
A DECADE OF WASTED CORES", EuroSys'16, Jean-Pierre Lozi, et.,al.


# HIERARCHICAL LOAD BALANCING



"THE LINUX SCHEDULER:  
A DECADE OF WASTED CORES", EuroSys'16, Jean-Pierre Lozi, et., al.

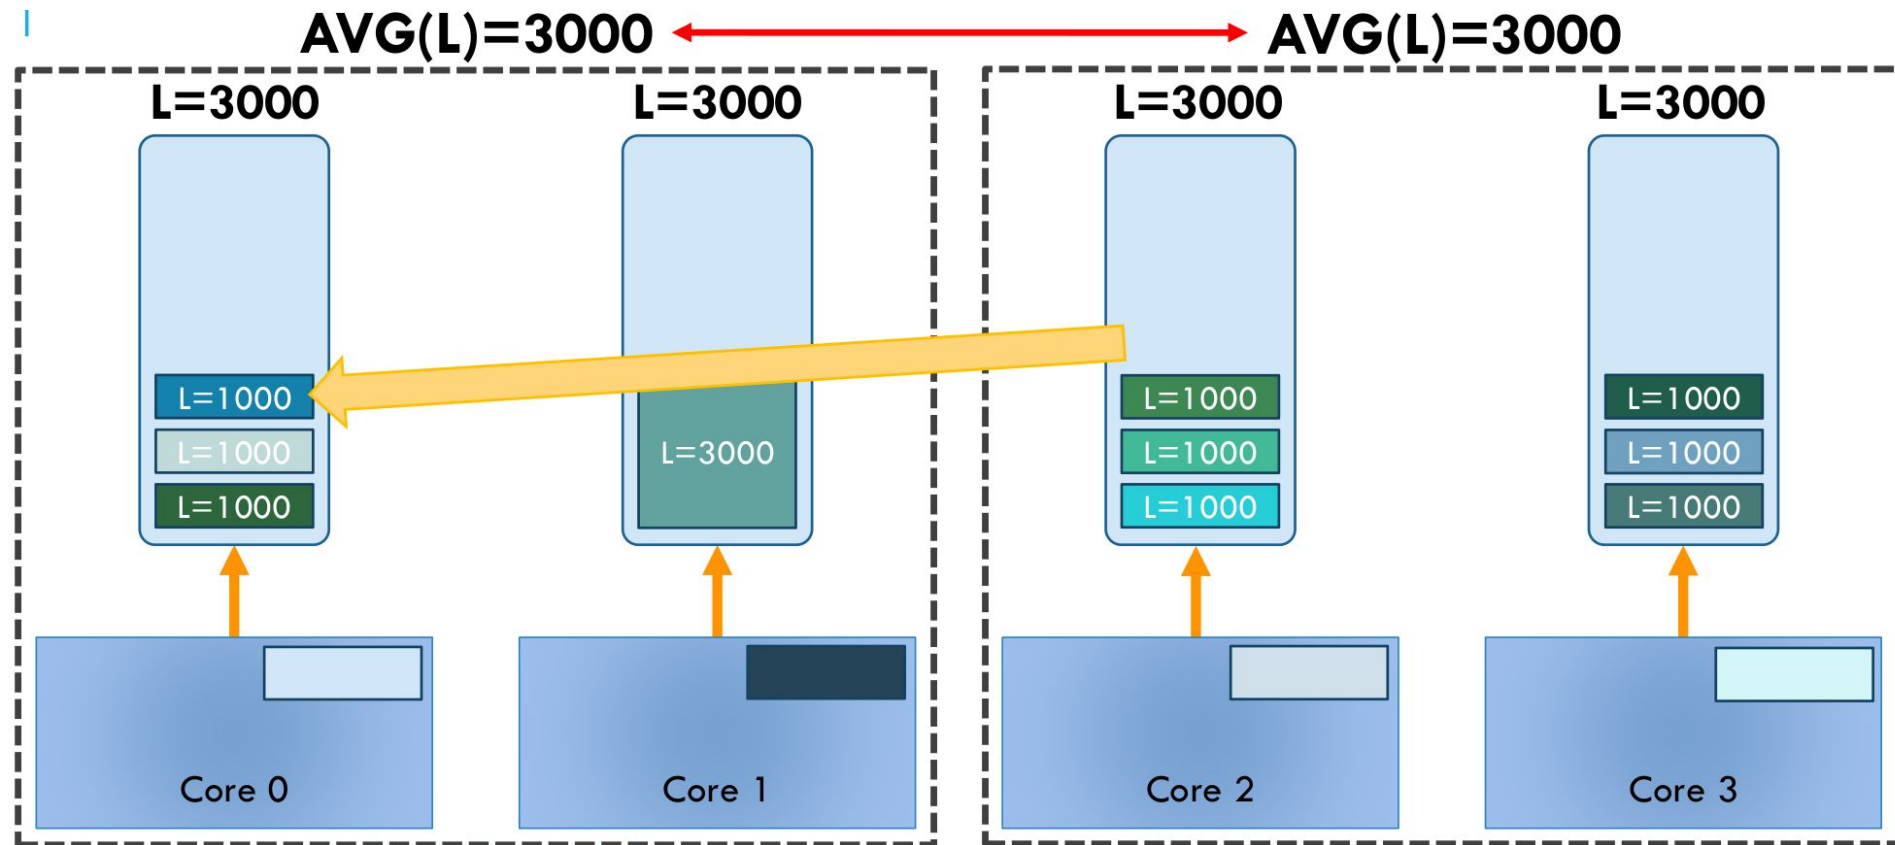


# HIERARCHICAL LOAD BALANCING



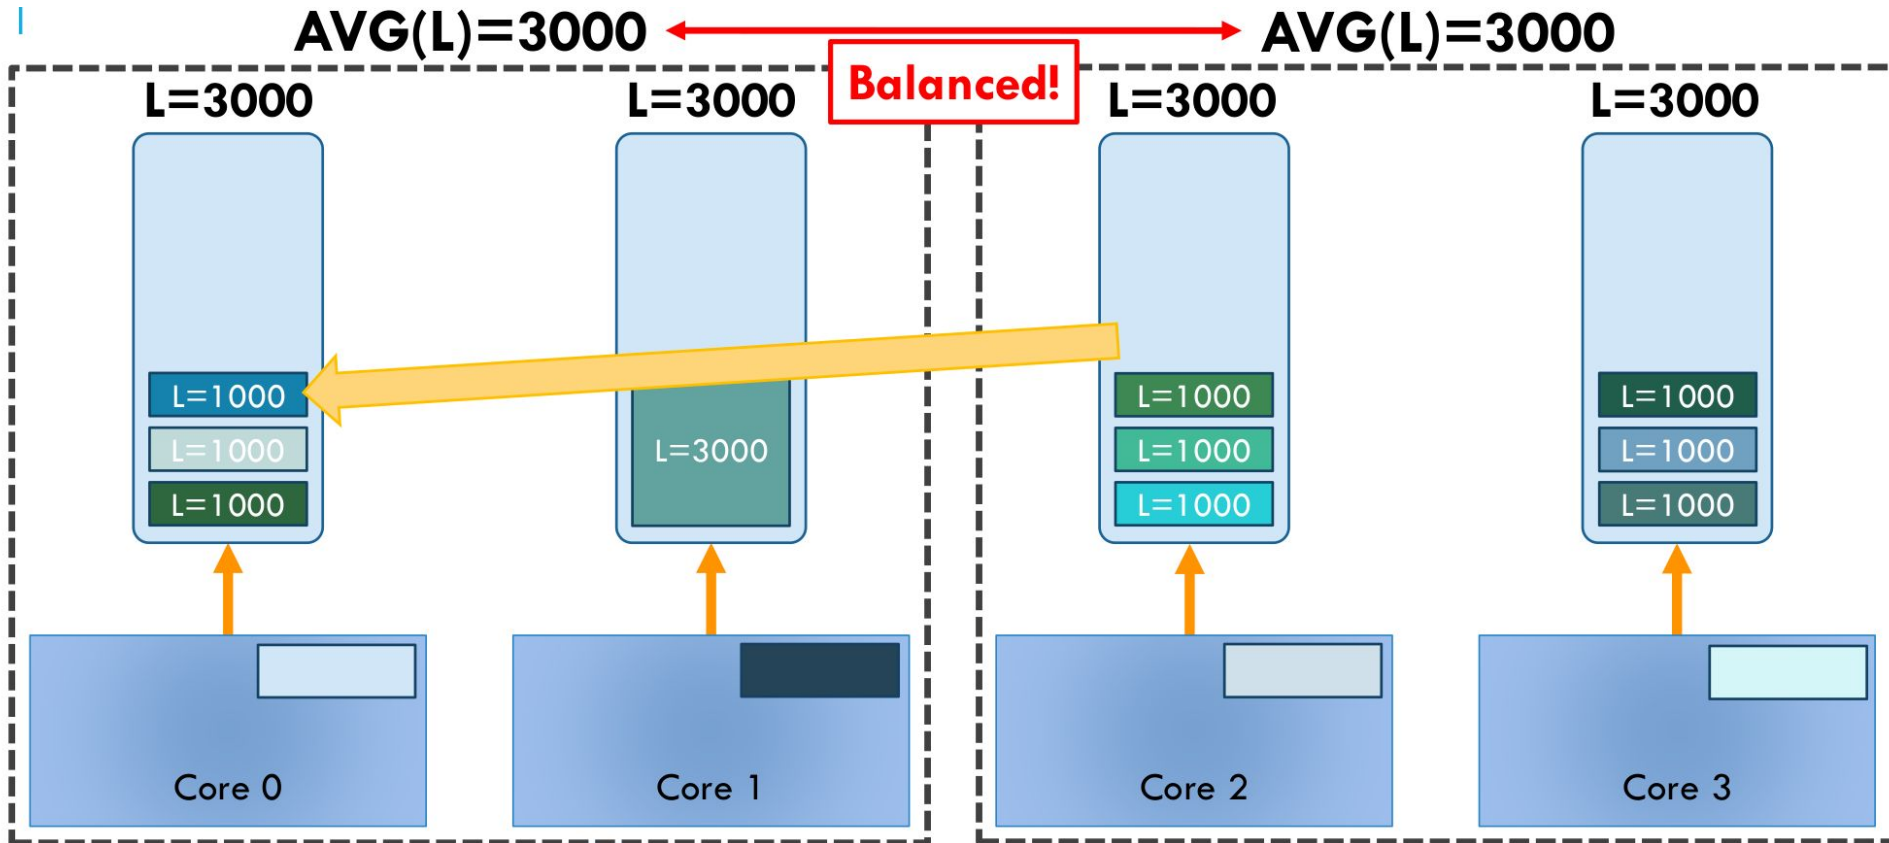
“THE LINUX SCHEDULER:  
A DECADE OF WASTED CORES”, EuroSys’16, Jean-Pierre Lozi, et.,al.

# HIERARCHICAL LOAD BALANCING



“THE LINUX SCHEDULER:  
A DECADE OF WASTED CORES”, EuroSys’16, Jean-Pierre Lozi, et.,al.

# HIERARCHICAL LOAD BALANCING



"THE LINUX SCHEDULER:  
A DECADE OF WASTED CORES", EuroSys'16, Jean-Pierre Lozi, et., al.

# **load\_balance()**

**kernel/sched/fair.c**

# CFS load balancer

- **Problem**

- How much is a process loading the system **right now**?
- a bursty or a steady task?
- a CPU-intensive or an I/O-bound task?
- It does not matter for scheduling
- It does matter for **load balancing**

# CFS Per-entity load tracking feature

- **Problem**

- How much is a process loading the system **right now**?
- a bursty or a steady task?
- a CPU-intensive or an I/O-bound task?
- It does not matter for scheduling
- It does matter for **load balancing**

- **Solution**


- Linux added a **CFS Per-entity load tracking feature** to estimate a task load.

# CFS load


- **Load of a process is important in scheduling**
  - which is required during load balancing.
- **CFS's load**
  - Thread's **weight** and its **average CPU utilization**
- **Load of CPU have been the sum of the load of all the scheduling entities.**



# CFS load



# CFS load



# Per-entity load avg

- **Scheduling is always at the granularity of sched\_entity**
- **A single task becomes a scheduling entity**

```
struct sched_avg {
 u64 last_update_time, load_sum;
 u32 util_sum, period_contrib;
 unsigned long load_avg, util_avg;
};
```

- **Each scheduling entity contains a runqueue.**

# load tracking

- Formula gives the most weight to the most recent load
- $y$  has been chosen so that  $y^{32}$  (32ms) is equal to 0.5
  - $y^0 = 1$ ,
  - $y^1 = 0.97852$
  - ...
  - $y^{32} = 0.5$
- Time is viewed as a sequence of 1ms (actually, 1024 $\mu$ s) periods

$$L = L_0 + L_1 * y + L_2 * y^2 + L_3 * y^3 + \dots$$

# load tracking

- PELT

$$L = L_0 + L_1 * y + L_2 * y^2 + L_3 * y^3 + \dots$$

- CFS's load

- Thread's **weight** and its **average CPU utilization**

- Old version

`sa.load_avg_contrib =`

`(sa.runnable_sum * se.load.weight) / sa.runnable_period;`

-> **Rewrite runnable load and utilization**

- New version

`load_avg : PELT(running time + runnable time) * weight`


`util_avg : PELT(running time) * CPU invariant`

**\_\_update\_load\_avg()**

**kernel/sched/fair.c**

# Per-entity load tracking

- The load is a history of time spent on the runqueue





“Update on big.LITTLE scheduling experiments”, ARM



# Example use cases

- HMP scheduler



**Next Step.**

# **Energy-aware scheduling: EAS**

## **1. EAS features**



# Reference

- <https://pdos.csail.mit.edu/6.828/2016/schedule.html>
- <http://web.mit.edu/6.033>
- <http://www.rdrop.com/~paulmck/>
- "Is Parallel Programming Hard, And If So, What Can You Do About It?"
- Davidlohr Bueso. 2014. Scalability techniques for practical synchronization primitives. *Commun. ACM* 58  
<http://queue.acm.org/detail.cfm?id=2698990>
- "CPUFreq and The Scheduler Revolution in CPU Power Management", Rafael J. Wysocki
- <https://sites.google.com/site/embedwiki/oses/linux/pm/pm-qos>
- <https://intl.aliyun.com/forum/read-916>
- User-level threads : co-routines  
<http://www.gamedevforever.com/291>  
[https://www.youtube.com/watch?v=YYtzQ355\\_Co](https://www.youtube.com/watch?v=YYtzQ355_Co)
- Scheduler Activations
  - <https://cgi.cse.unsw.edu.au/~cs3231/12s1/lectures/SchedulerActivations.pdf>
- [https://en.wikipedia.org/wiki/FIFO\\_\(computing\\_and\\_electronics\)](https://en.wikipedia.org/wiki/FIFO_(computing_and_electronics))
- <http://jake.dothome.co.kr/>
- <http://www.linuxjournal.com/magazine/completely-fair-scheduler?page=0.0>
- [https://www2.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/6\\_CPU\\_Scheduling.html](https://www2.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/6_CPU_Scheduling.html)
- "Energy Aware Scheduling", Byungchul Park, LG Electronic
- "Update on big.LITTLE scheduling experiments", ARM
- "EAS Update" 2015 september ARM
- "EAS Overview and Integration Guide", ARM TR
- "Drowsy Power Management", Matthew Lentz, SOSP 2015
- <https://www.slideshare.net/nanik/learning-aosp-android-hardware-abstraction-layer-hal>
- <https://www.youtube.com/watch?v=oTGQXqD3CNI>
- <https://www.youtube.com/watch?v=P80NcKUKpuo>
- <https://lwn.net/Articles/398470/>
- "SCHED\_DEADLINE: It's Alive!", ARM, 2017