

Android PM, Deadline(RM, EDF)

국민대학교 임베디드 연구실
경주현

Outline

- **Android PM**
- **Deadline Scheduler**



Android PM

Android PM

- **Android inherits the Linux PM**
- **Wakelocks**
- **Scheduling Wakeups**
- **Early Suspend**

Task may be blocked waiting on I/O

- **Problem:**
 - Task may be blocked waiting on I/O
 - Ex) Weather application
 - Waiting for a response from the remote server

Wakelocks

- **Problem:**
 - Task may be blocked waiting on I/O
 - Weather application
 - Waiting for a response from the remote server
- **Solution:**
 - Android uses wakelock
- **Wakelock**
 - While any wakelock is held the **system does not suspend**

Wakelock's problem

- **Wakelock's problem**
 - CPU can not go to suspend.
 - High power consumption

Scheduling Wakeups

- **Wakelock's problem**
 - CPU can not go to suspend.
 - High power consumption
- **Solution**
 - To allow applications to wake up at a specified time
- **Scheduling Wakeups**

Scheduling Wakeups

- To allow applications to wake up at a specified time
- Android introduces wakeup alarms
- /dev/alarm virtual device and AlarmManager
- Set next alarm -> wake up

Wake-lock for kernel

- **Problem**
 - User press the power button
 - Users want display and backlight power down
- **Solution**

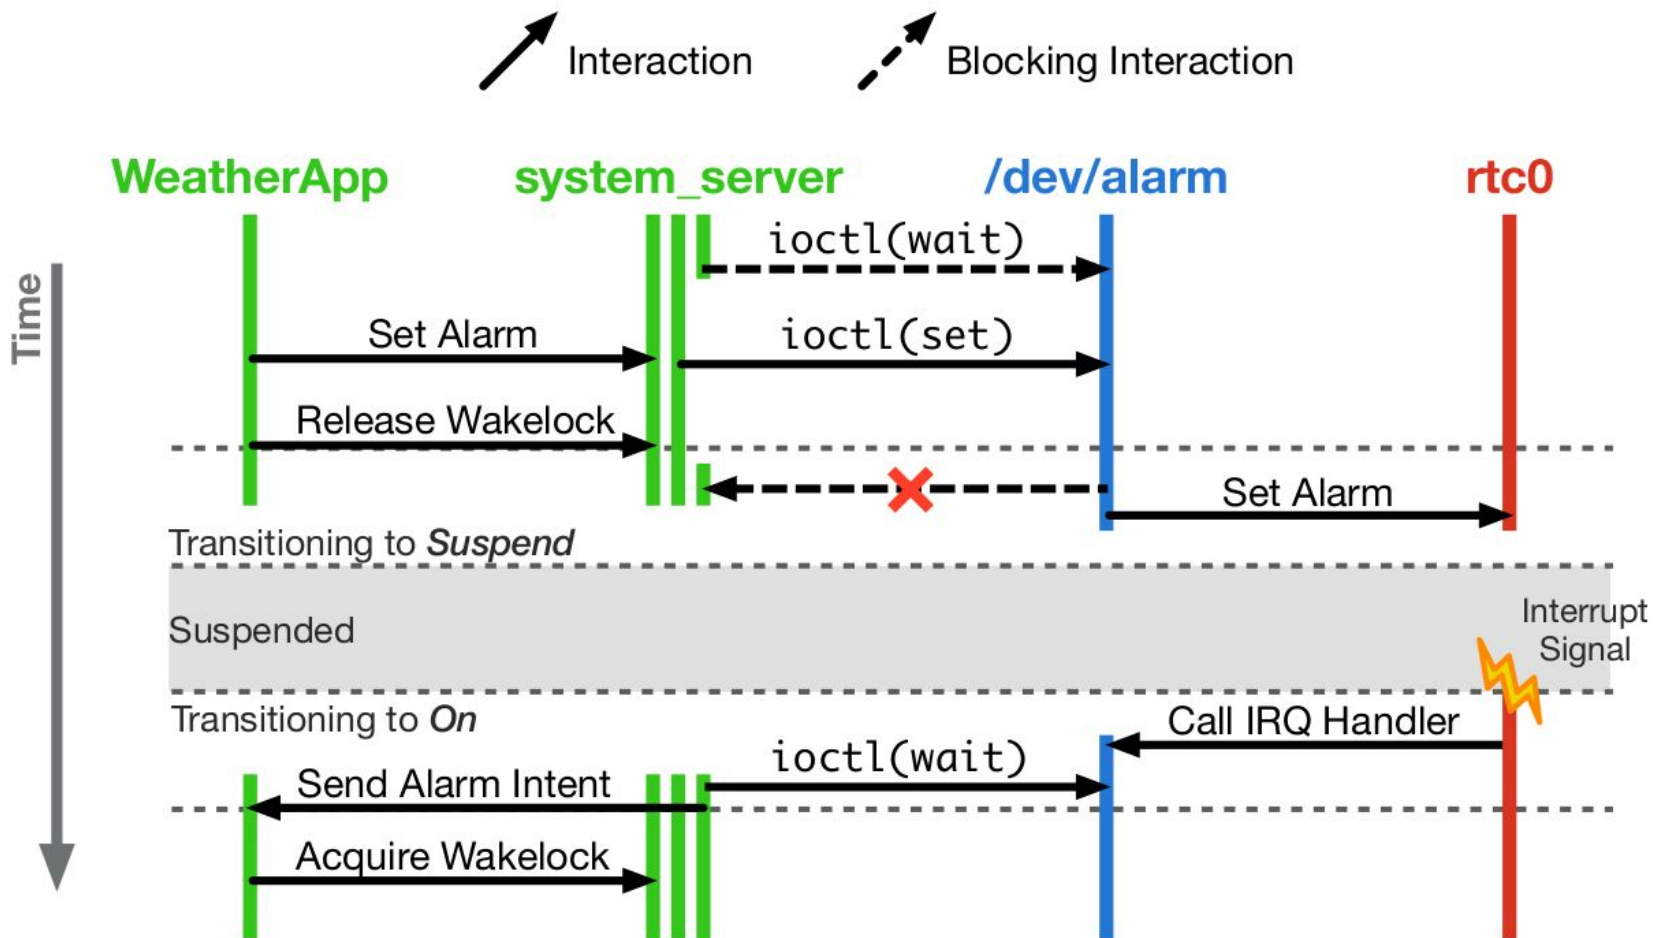
Wake-lock for kernel

- **Problem**
 - User press the power button
 - Users want display and backlight power down
- **Solution**
 - Early Suspend

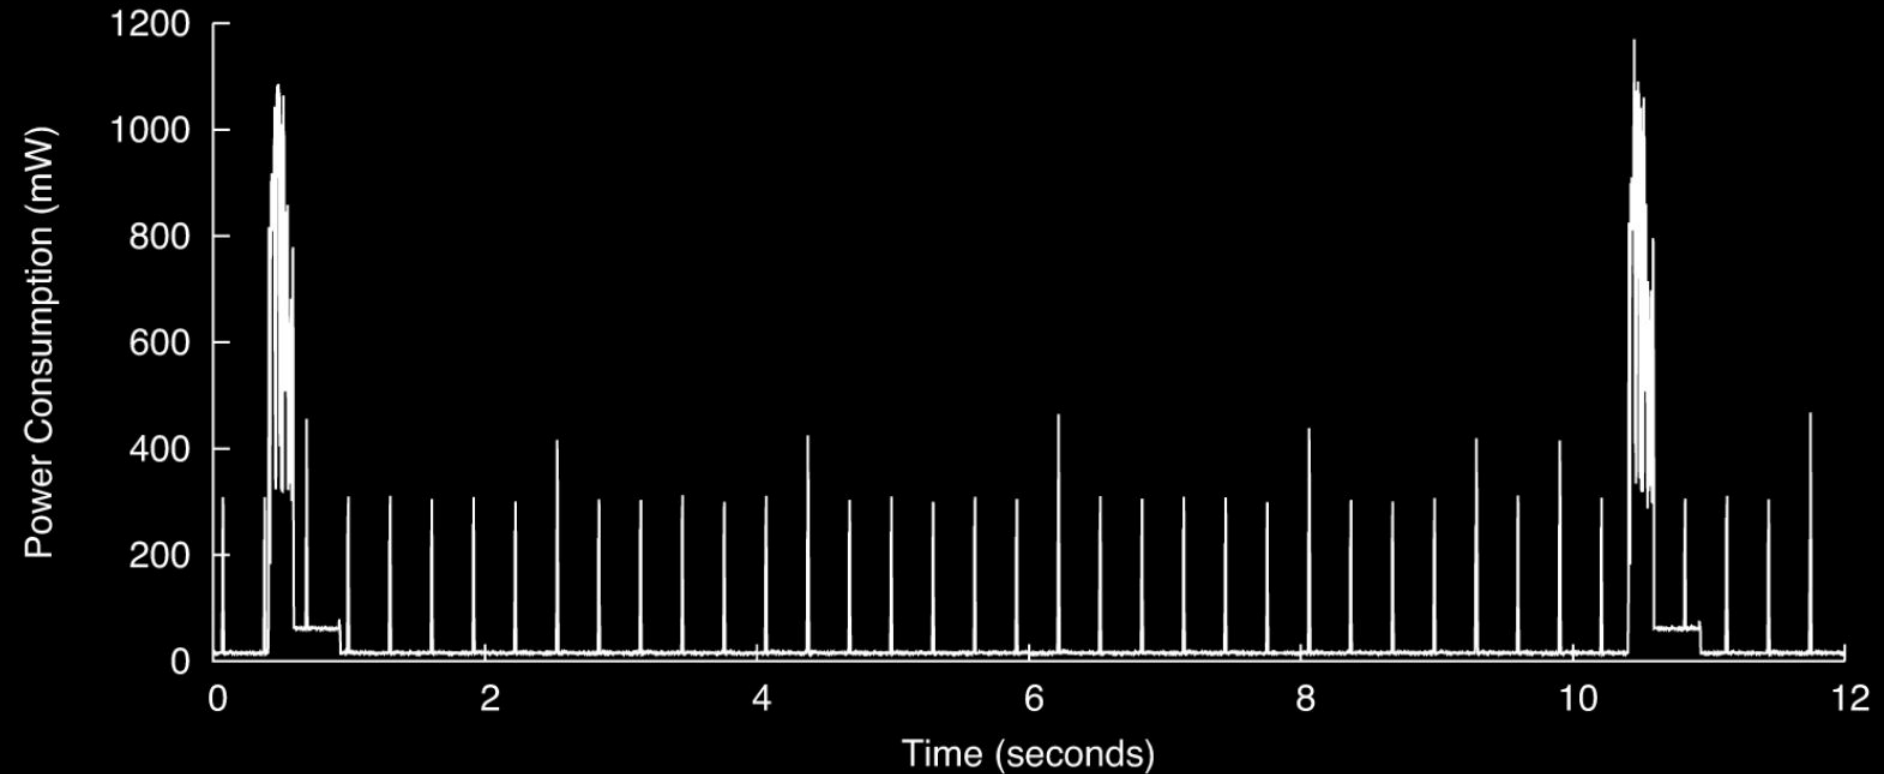
Early Suspend

- **Android register early suspend(and late resume)**
- **User presses the power button**
- **The system invokes all early suspend callbacks**
 - E.g., display and backlight power down

Message sequence of system interactions in Android

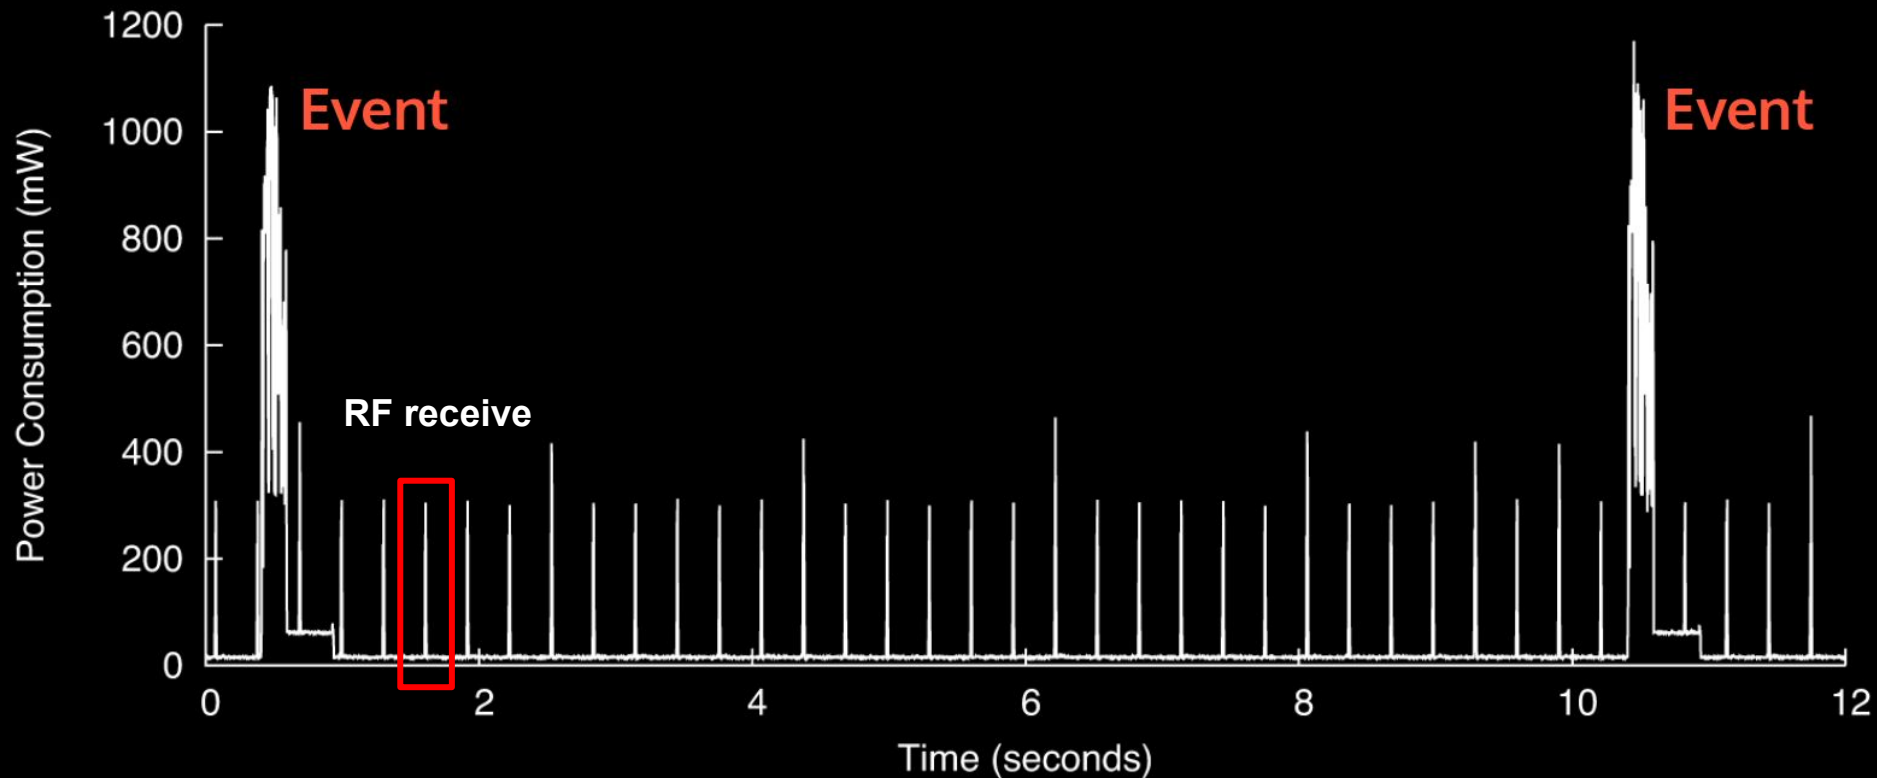


Android background

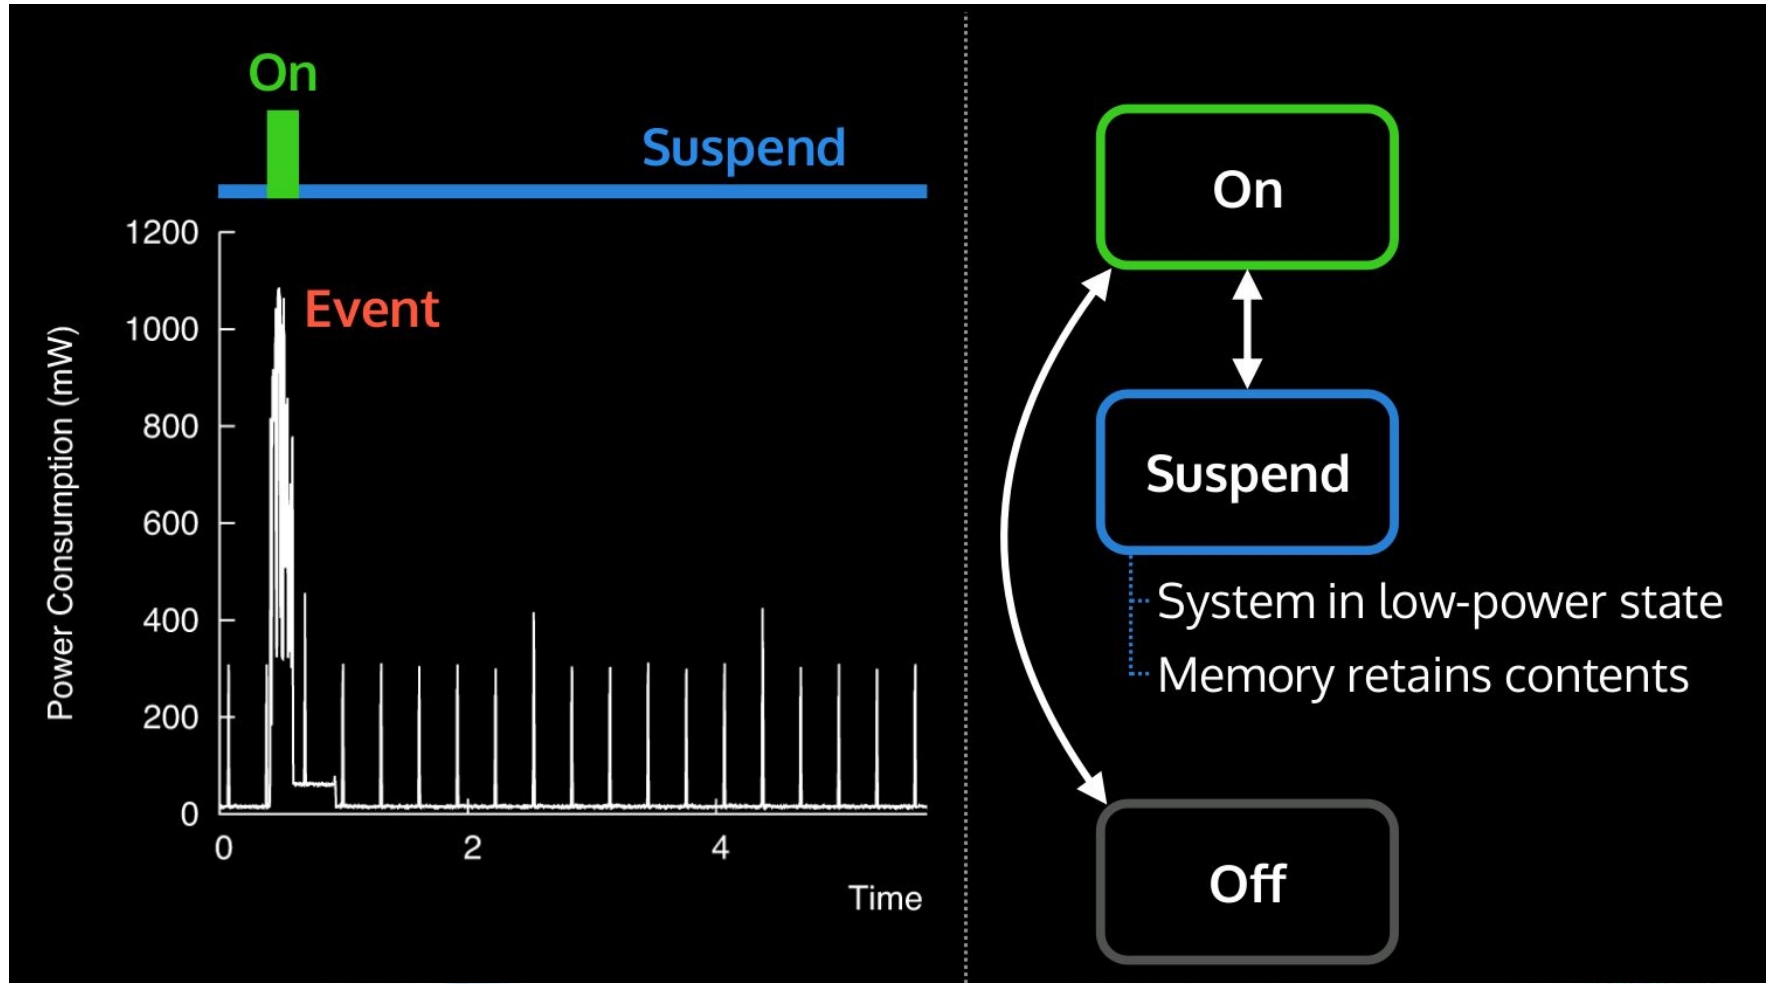


Background Energy Consumption

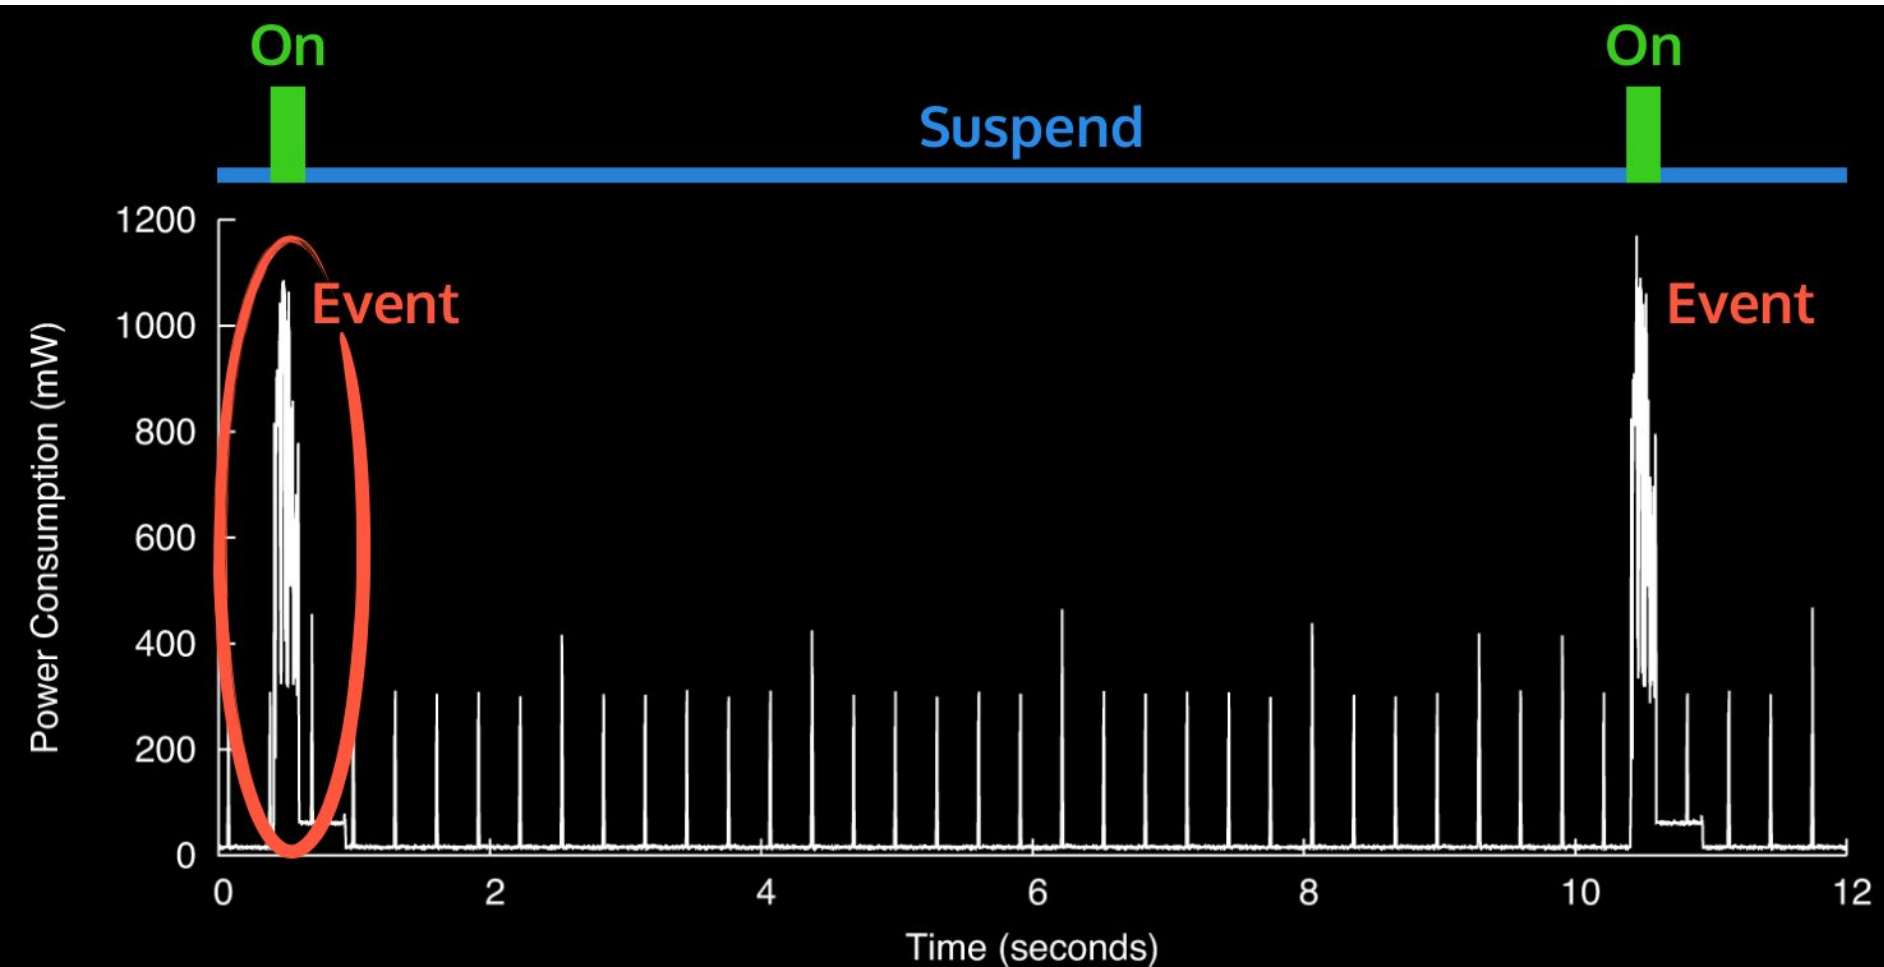
Event = Pull data from remote server over WiFi



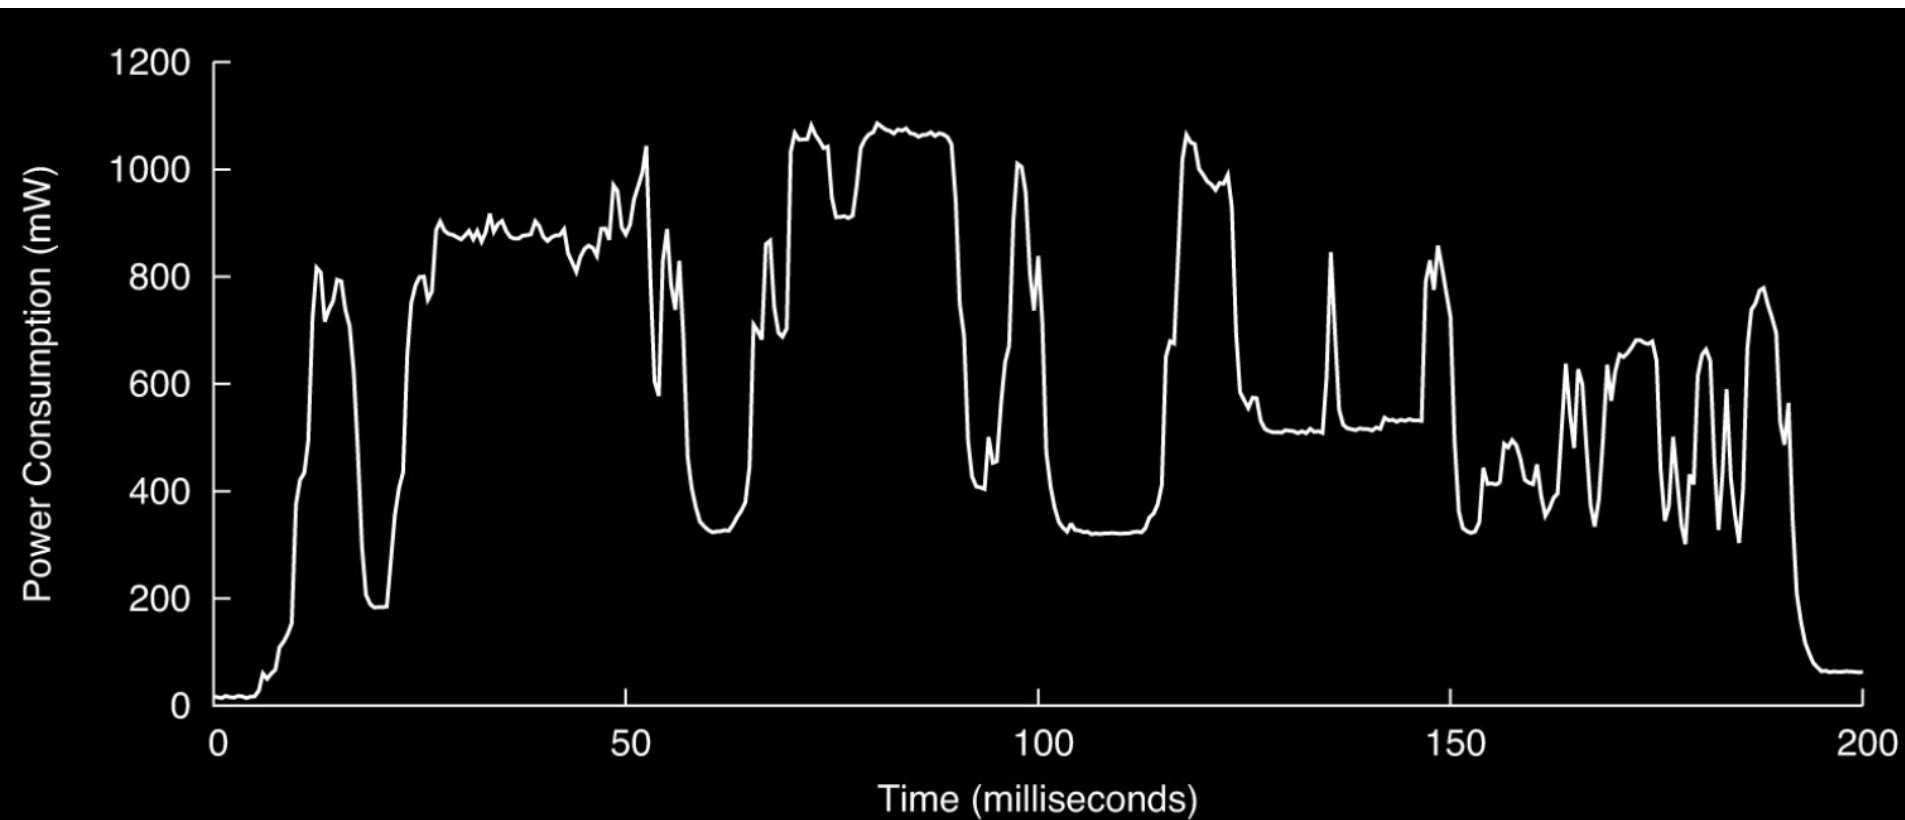
Background Energy Consumption



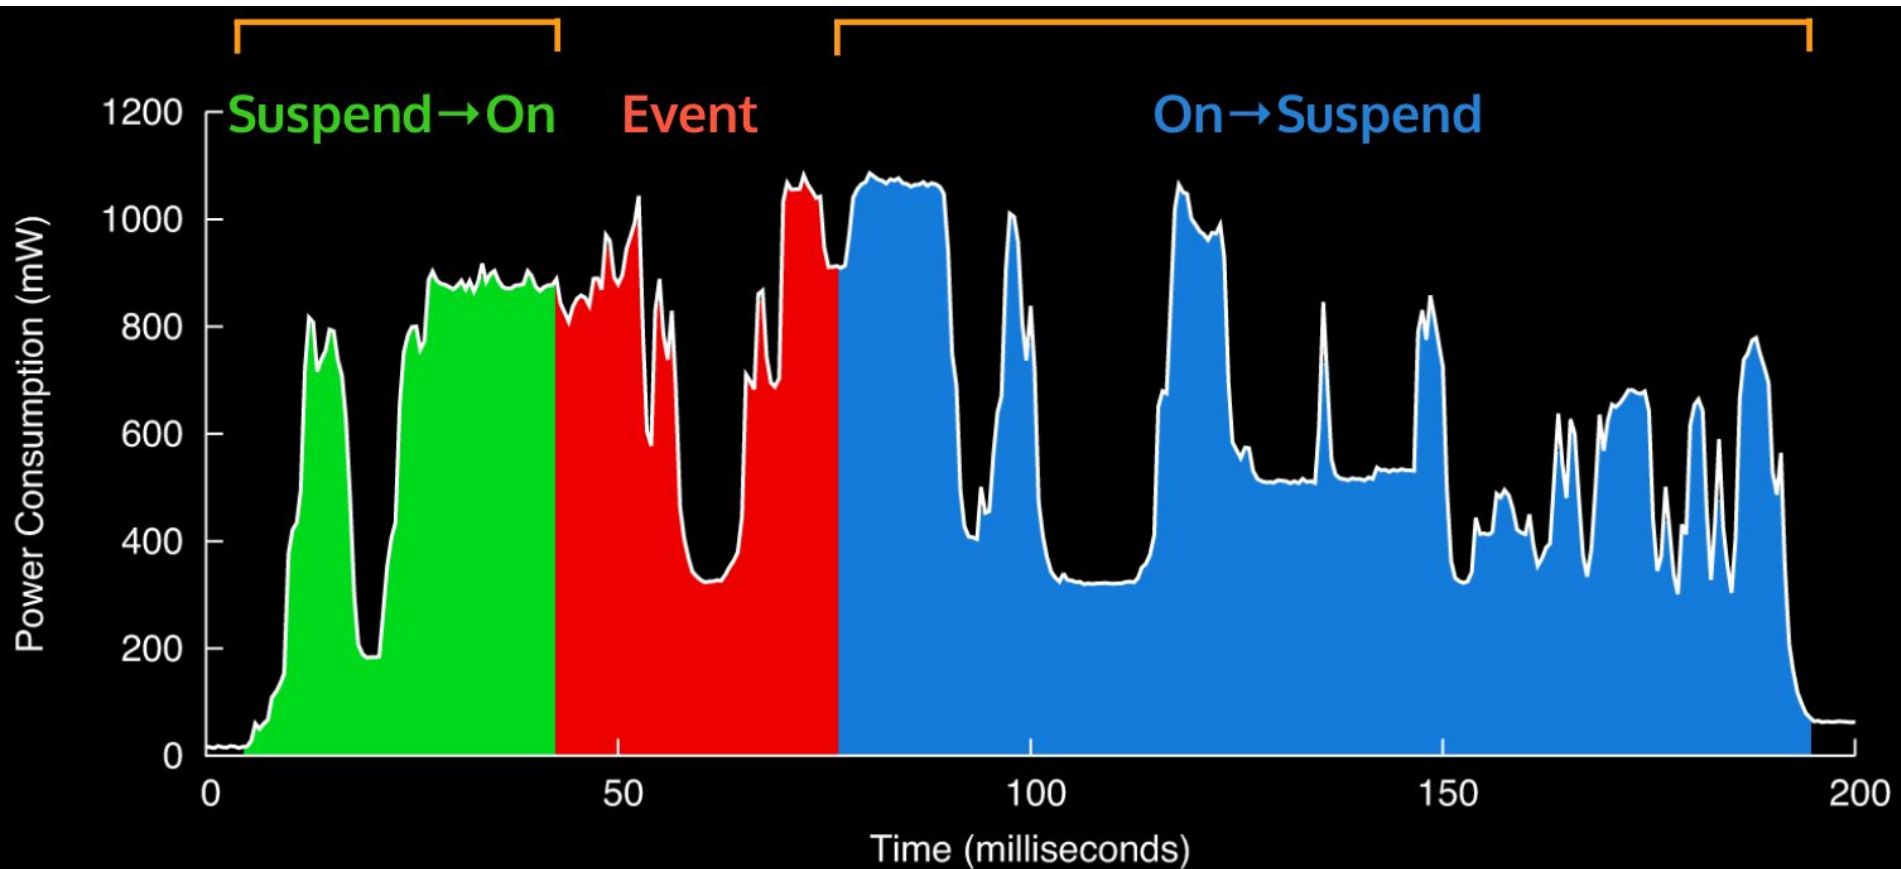
Background Energy Consumption



Event

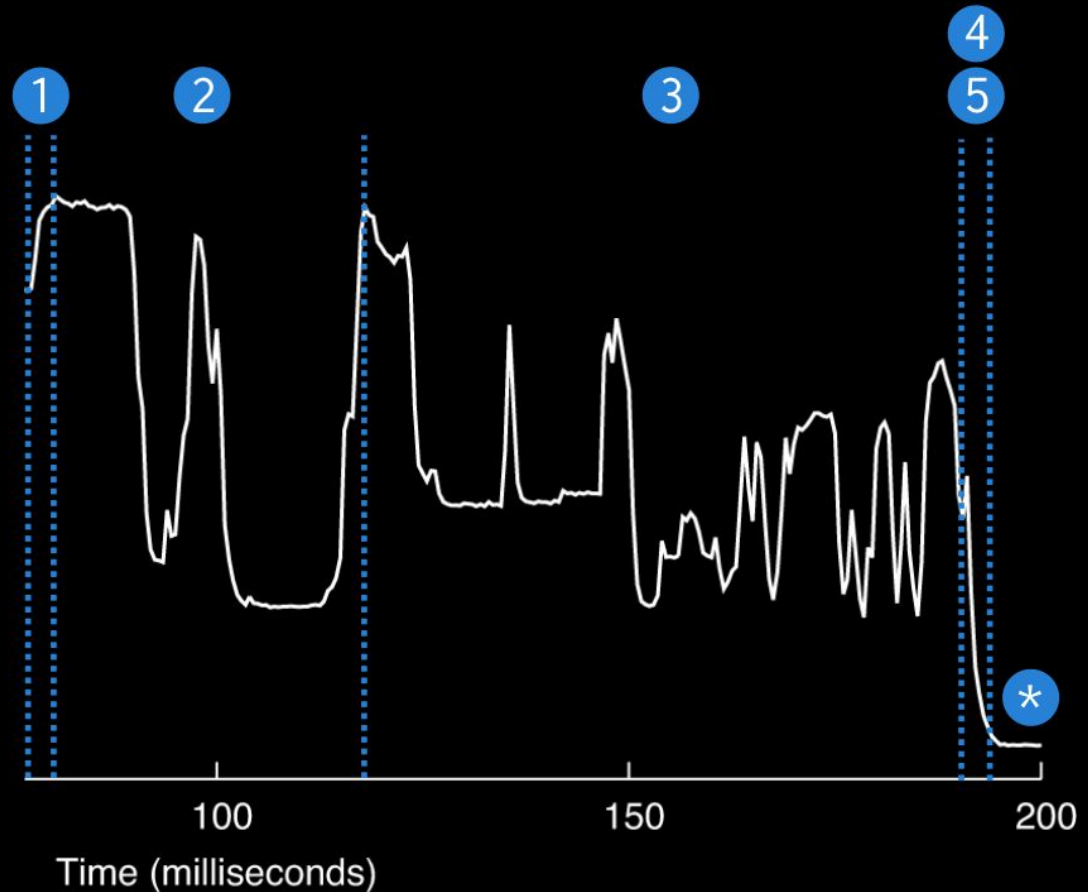


Wake-up and Suspend

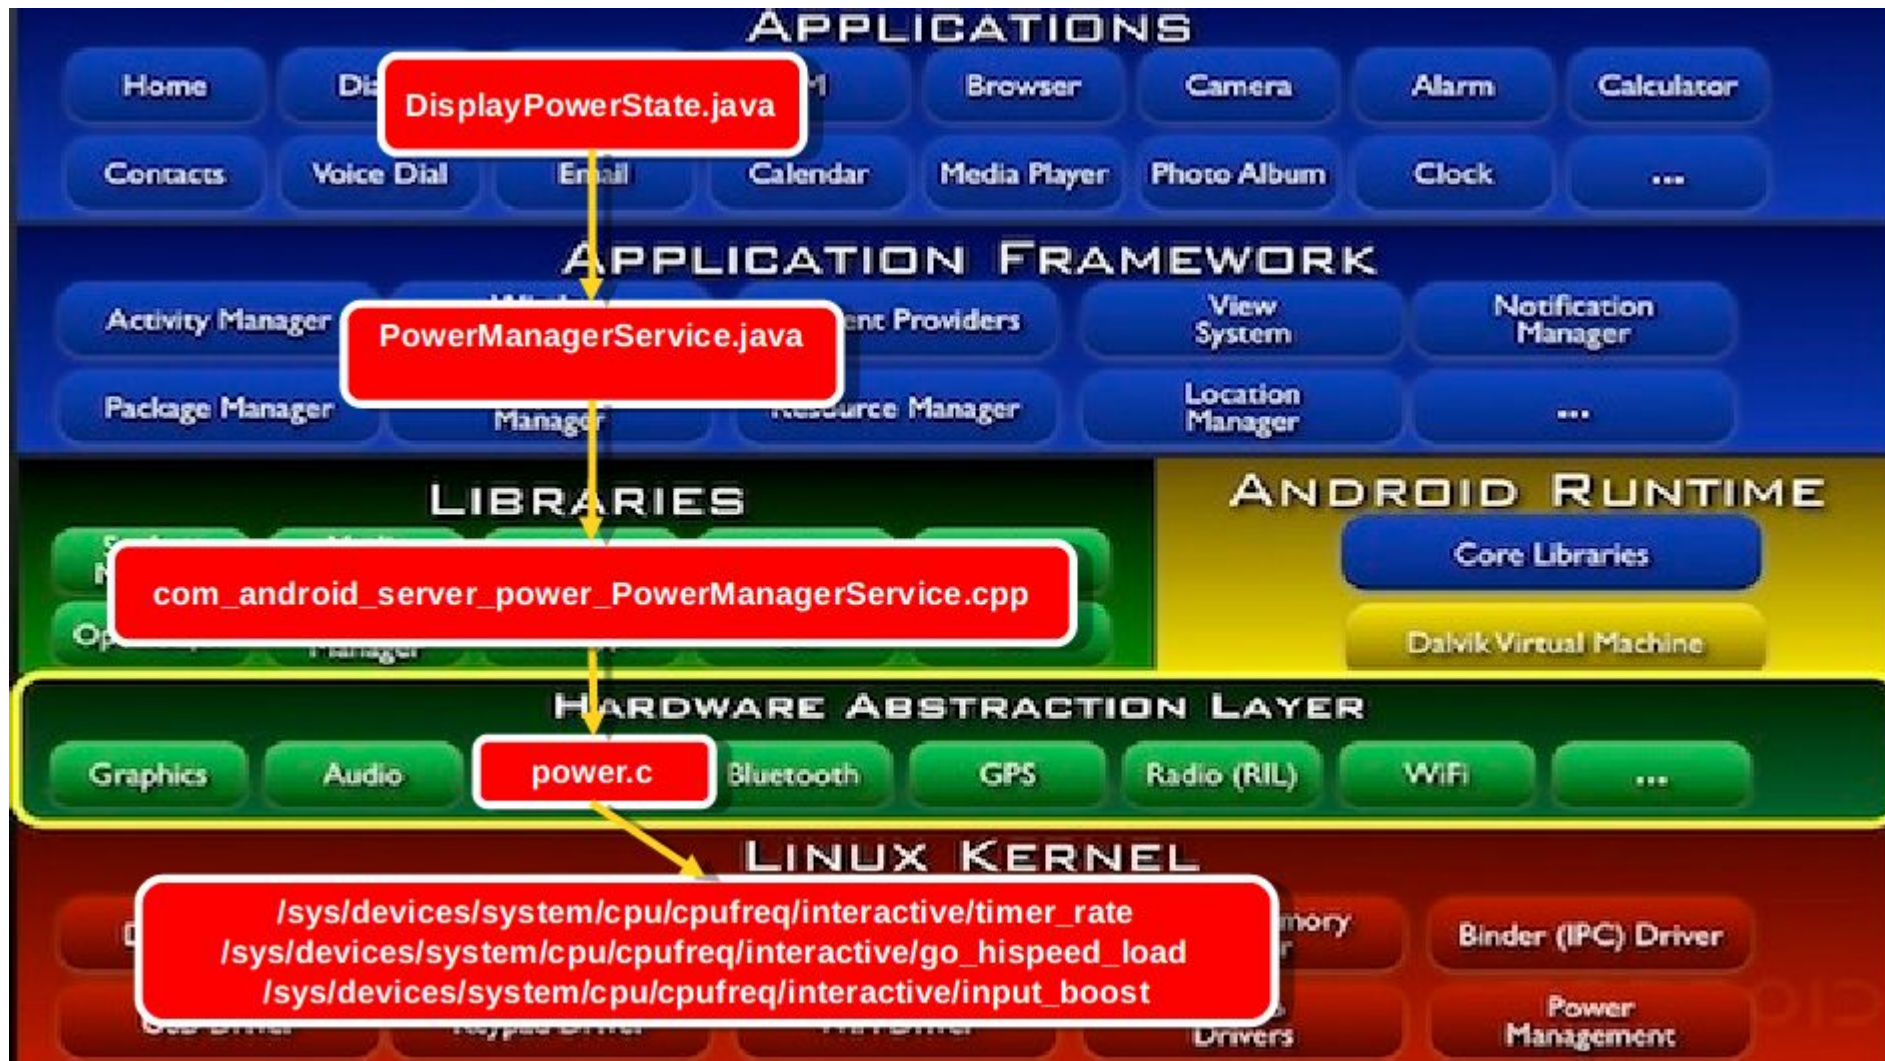


Background Energy Consumption

- 1 Flush filesystem buffers
- 2 Freeze all tasks
- 3 Suspend all devices
- 4 Disable non-boot CPUs
- 5 Set RAM to self-refresh
- * Wait for interrupt



Android power HAL



Android power HAL

- **A way of providing user's experience information.**
- **Android Power HINT -> kernel**
- **Ex) Touch boost, vsync**

Sample implementations of power HAL

power.\$(TARGET_BOARD_PLATFORM).so file can be found (hint: look at the Android.mk's for these as well as the c code):

- hardware/qcom/power
- device/asus/grouper/power
- device/samsung/manta/power
- device/samsung/tuna/power
- device/generic/goldfish/power



Real-time scheduling

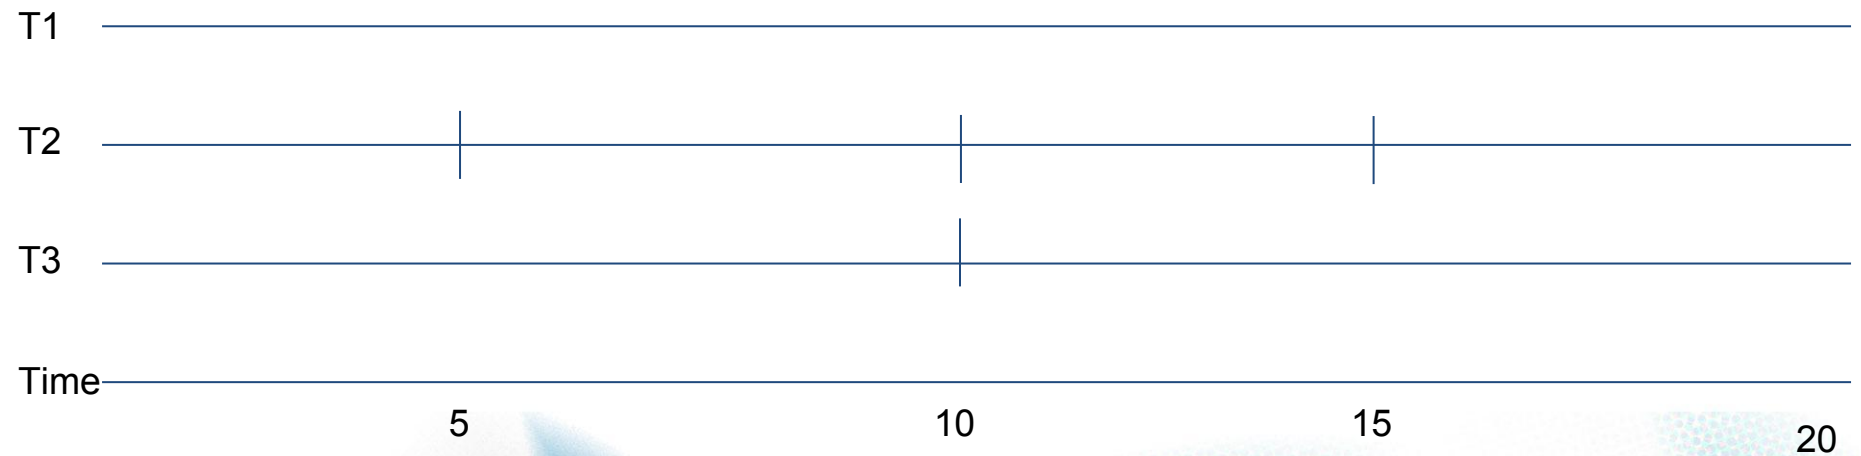
Real-time scheduling

- **Determines the order of real-time task executions**
- **Static-priority scheduling**
 - **Rate Monotonic**
- **Dynamic-priority scheduling**
 - **EDF**

Example : RM

-

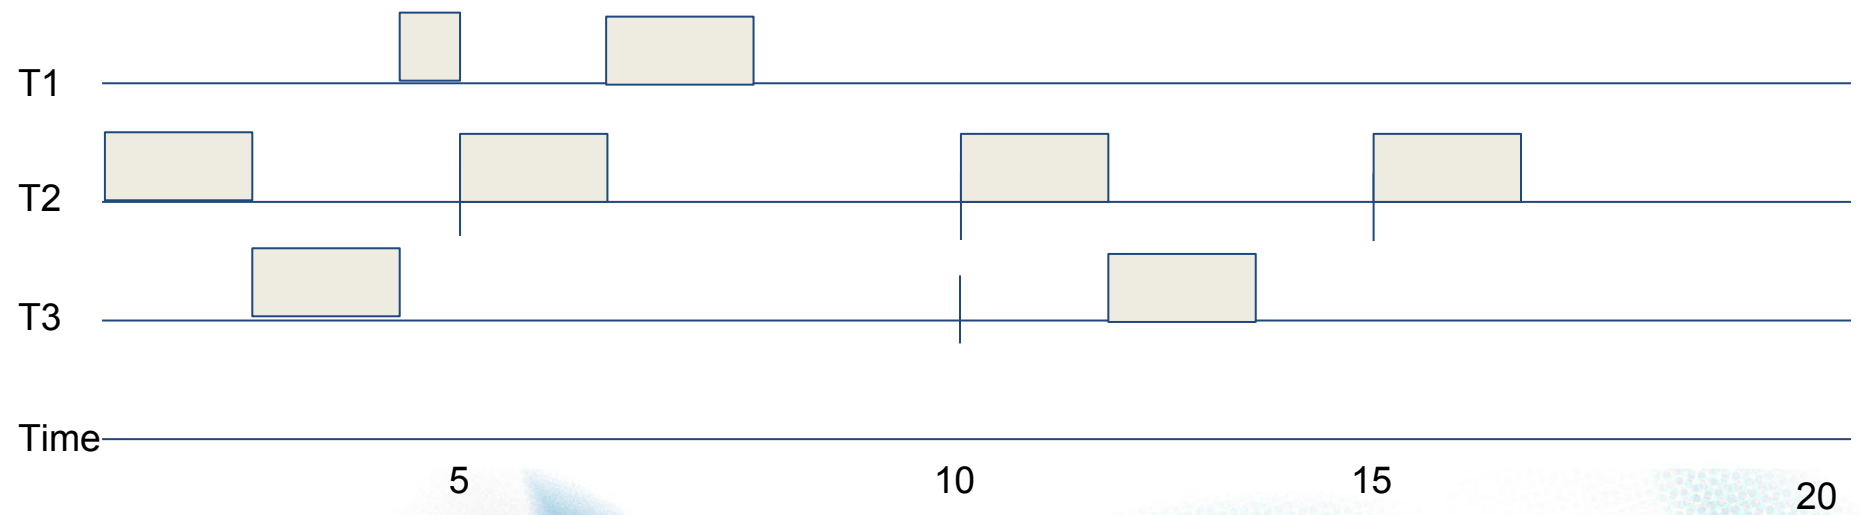
	C	P
T1	3	20
T2	2	5
T3	2	10



Example : RM

-

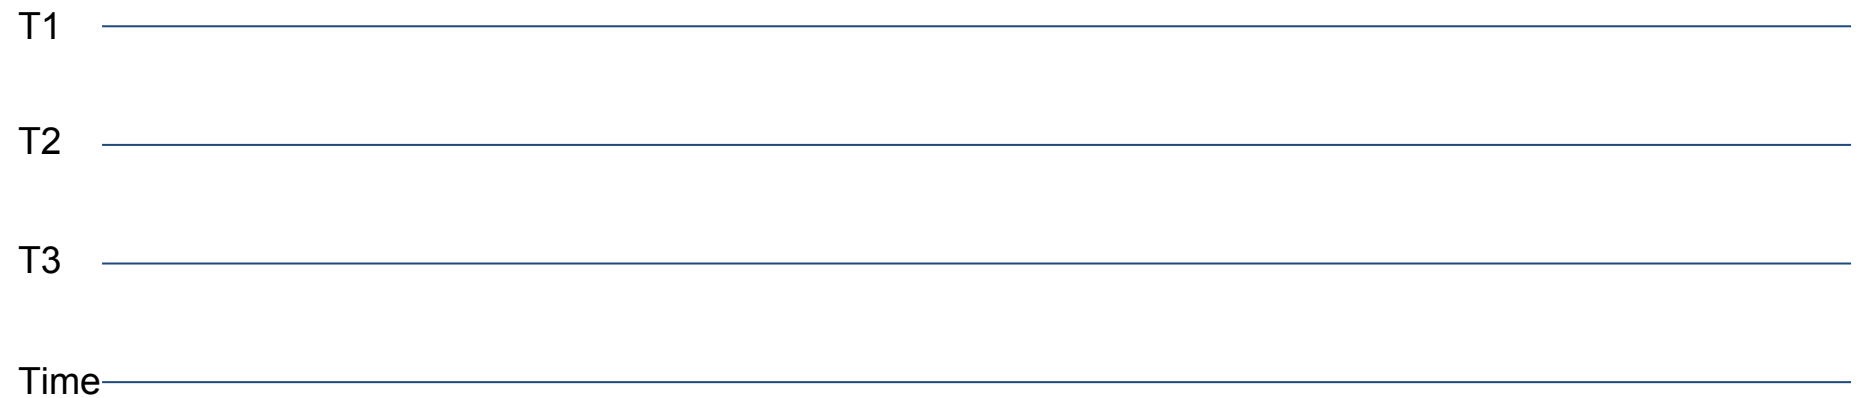
	C	P
T1	3	20
T2	2	5
T3	2	10



Example : EDF

-

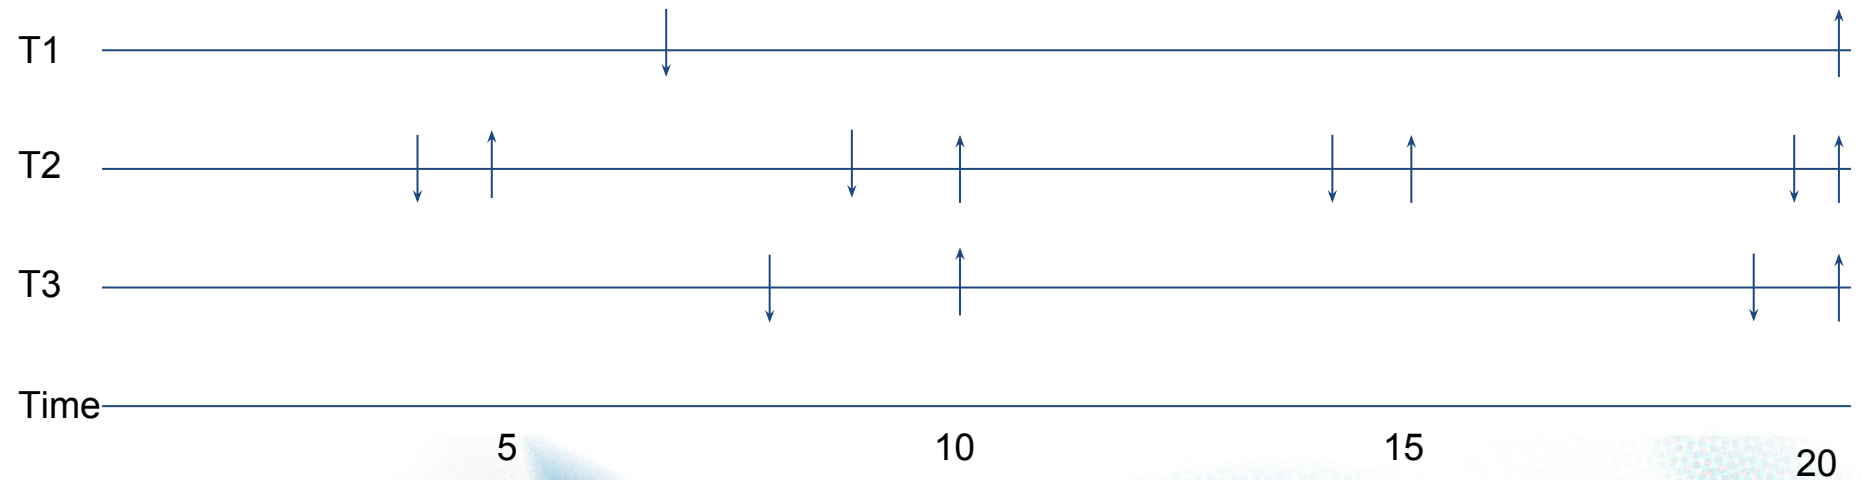
	C	D	P
T1	3	7	20
T2	2	4	5
T3	2	8	10



Example : EDF

-

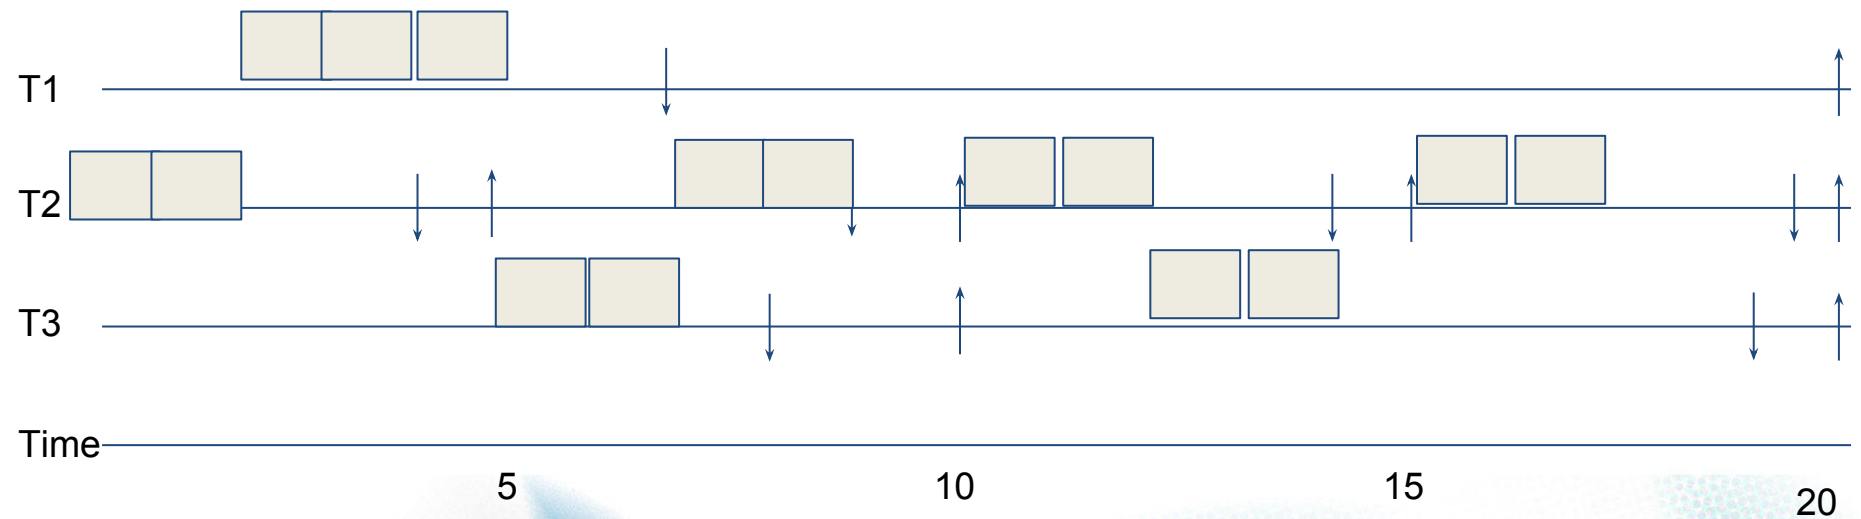
	C	D	P
T1	3	7	20
T2	2	4	5
T3	2	8	10



Example : EDF

-

	C	D	P
T1	3	7	20
T2	2	4	5
T3	2	8	10

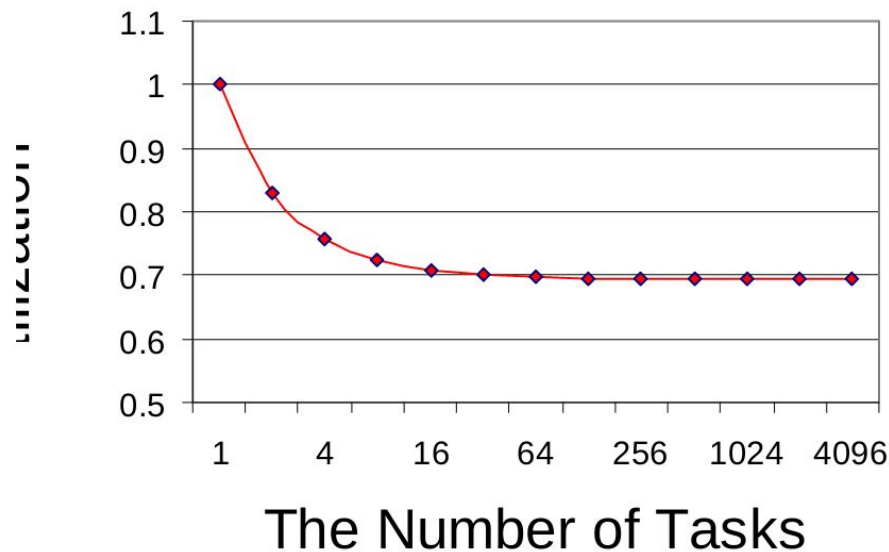


RM Problem – Utilization Bound

- n = number of process

$$U = \sum_{i=1}^n \frac{C_i}{T_i} \leq n(\sqrt[n]{2} - 1) \quad \longrightarrow \quad \lim_{n \rightarrow \infty} n(\sqrt[n]{2} - 1) = \ln 2 \approx 0.693147 \dots$$

RM Utilization Bounds



RM vs. EDF

- **Rate Monotonic**
 - Simpler implementation, even in systems without explicit support for timing constraints (periods, deadlines)
 - Predictability for the highest priority tasks
- **EDF**
 - Full processor utilization
 - Misbehavior during overload conditions
- **For more details: Buttazzo, “Rate monotonic vs. EDF: Judgement Day”, EMSOFT 2003.**

EDF

- **Can EDF be supported in Linux?**
- **Problem**
 - The kernel is not aware of tasks deadline

Can EDF be supported in Linux?

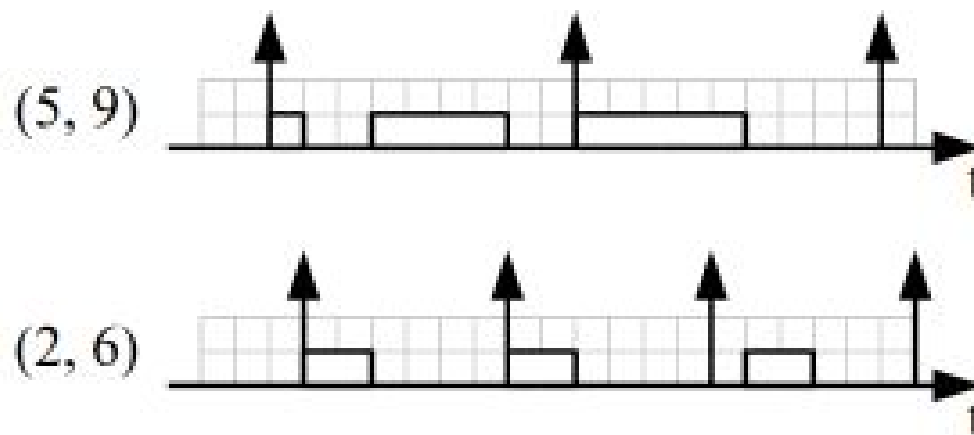
- **Can EDF be supported in Linux?**
- **Problem**
 - The kernel is not aware of tasks deadline
- **Simple Solution**
 - Apps send signal to OS
 - Start()/End()

Linux's sched_attr

```
struct sched_attr {  
    __u32 size;  
  
    __u32 sched_policy;  
    __u64 sched_flags;  
  
    /* SCHED_NORMAL, SCHED_BATCH */  
    __s32 sched_nice;  
  
    /* SCHED_FIFO, SCHED_RR */  
    __u32 sched_priority;  
  
    /* SCHED_DEADLINE (nsec) */  
    __u64 sched_runtime;  
    __u64 sched_deadline;  
    __u64 sched_period;  
};
```

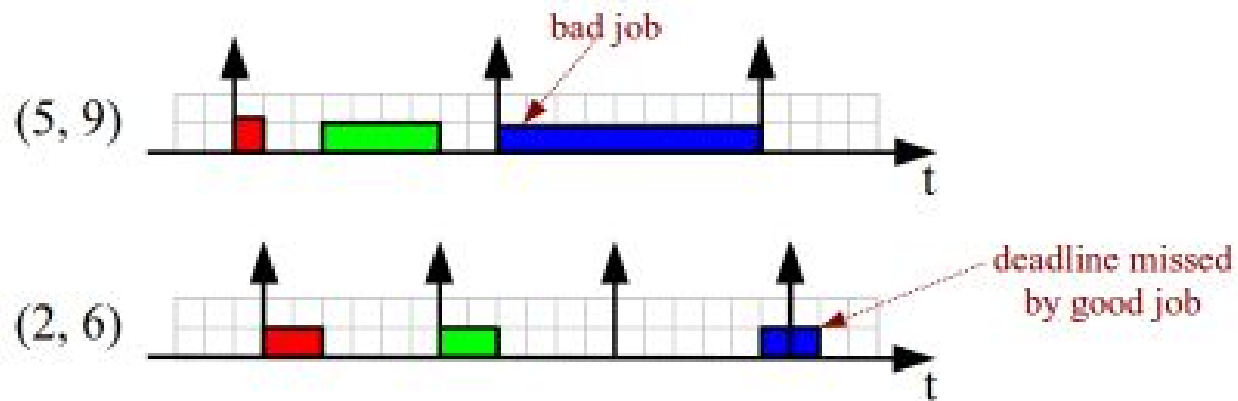
Example

- Two tasks with reservations
- 5ms every 9ms and 2ms every 6ms



~88.9%
Overall
Load

Real world Problem

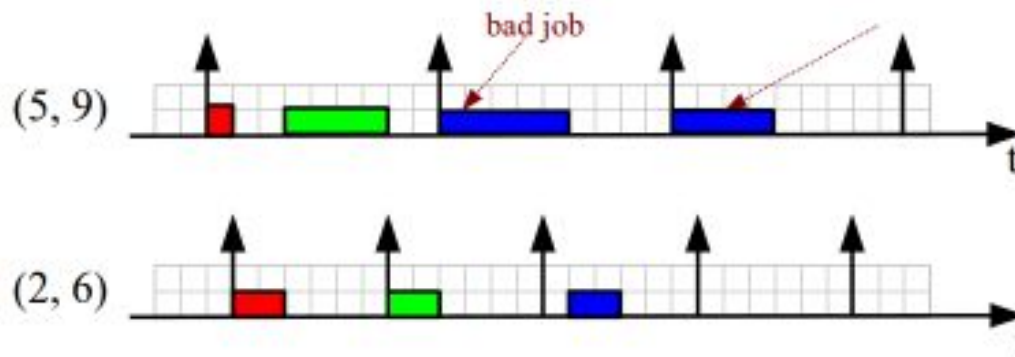


kernel is not aware of tasks deadline

- **Use dynamic scheduling deadlines**
- **Constant Bandwidth Server(CBS)**
 - A resource(CPU) reservation mechanism

Constant Bandwidth Server(CBS)

- Current budget $\rightarrow 0$
Suspended (throttled) till the next activation period
- Resource(CPU) reservation



Constant Bandwidth Server(CBS)

- **Problem**

- properly dimensioning runtime and reservation period
 - **might be too difficult**
- What if tasks occasionally need more bandwidth?
 - e.g., occasional workload fluctuations (network traffic, rendering of particularly heavy frame, etc)

- **Solution**

- Use resource over-allocation.
- Greedy Reclamation of Unused Bandwidth (GRUB)

SCHED_DEADLINE: It's Alive!

- **ARM's presentation.**
 - ELC North America 17, Portland (OR) 02/21/2017
- **Near**
 - Experimenting with Android
 - Reclaiming by demotion towards lower priority class
 - Capacity awareness (for heterogeneous systems)
 - Energy awareness (Energy Aware Scheduling for DEADLINE)
- **And..**
 - Support single CPU affinity
 - Enhanced priority inheritance (M-BWI most probably)
 - Dynamic feedback mechanism (adapt reservation parameters to task' needs)

Reference

- <https://pdos.csail.mit.edu/6.828/2016/schedule.html>
- <http://web.mit.edu/6.033>
- <http://www.rdrop.com/~paulmck/>
- "Is Parallel Programming Hard, And If So, What Can You Do About It?"
- Davidlohr Bueso. 2014. Scalability techniques for practical synchronization primitives. *Commun. ACM* 58

<http://queue.acm.org/detail.cfm?id=2698990>

- "CPUFreq and The Scheduler Revolution in CPU Power Management", Rafael J. Wysocki
- <https://sites.google.com/site/embedwiki/oses/linux/pm/pm-qos>
- <https://intl.aliyun.com/forum/read-916>
- User-level threads : co-routines

<http://www.gamedevforever.com/291>

https://www.youtube.com/watch?v=YYtzQ355_Co

- Scheduler Activations
 - <https://cgi.cse.unsw.edu.au/~cs3231/12s1/lectures/SchedulerActivations.pdf>
- [https://en.wikipedia.org/wiki/FIFO_\(computing_and_electronics\)](https://en.wikipedia.org/wiki/FIFO_(computing_and_electronics))
- <http://jake.dothome.co.kr/>
- <http://www.linuxjournal.com/magazine/completely-fair-scheduler?page=0.0>
- https://www2.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/6_CPU_Scheduling.html
- "Energy Aware Scheduling", Byungchul Park, LG Electronic
- "Update on big.LITTLE scheduling experiments", ARM
- "EAS Update" 2015 september ARM
- "EAS Overview and Integration Guide", ARM TR
- "Drowsy Power Management", Matthew Lentz, SOSP 2015
- <https://www.slideshare.net/nanik/learning-aosp-android-hardware-abstraction-layer-hal>
- <https://www.youtube.com/watch?v=oTGQXqD3CNI>
- <https://www.youtube.com/watch?v=P80NcKUKpuo>
- <https://lwn.net/Articles/398470/>
- "SCHED_DEADLINE: It's Alive!", ARM, 2017